

Dependencies and Performance

Computer Architecture
ECE 6913

Brandon Reagen

What we'll cover today

- 1) Review:
 - 1) Control flow example
 - 2) Stack
- 2) Dependencies and pipelining
- 3) Understanding performance
- 4) Lab 2

MIPS Instructions: I-Type

opcode	rs	rt	immediate
6 bits	5 bits	5 bits	16 bits

addi rs, rt, imm // $R[rt] \leftarrow R[rs] + \{\text{SignExtend, imm}\}; \text{MSB of imm is extended to 32 bits}$

ori rs, rt, imm // $R[rt] \leftarrow R[rs] | \{\text{ZeroExtend, imm}\}; \text{bit-wise Boolean OR operation}$

beq rs, rt, imm // if{ $R[rs] == R[rt]$ } branch to $\text{PC} + 4 + \text{BranchAddress}$; (“PC relative”)
// Else go to PC+4
// BranchAddress = {SignExtend, imm, oo}

lw rs, rt, imm // $R[rt] \leftarrow \text{Mem}[\{\text{SignExtend, imm}\} + R[rs]]$ (“Displaced/based”)

MIPS Instructions: J-Type

opcode	address
6 bits	26 bits

j address // JumpAddress = { (PC+4)[31:28], address, 2'bo }
 // (Pseudodirect)
jal address // PC ← JumpAddress
 // R[31] ← PC+4; PC ← JumpAddress;

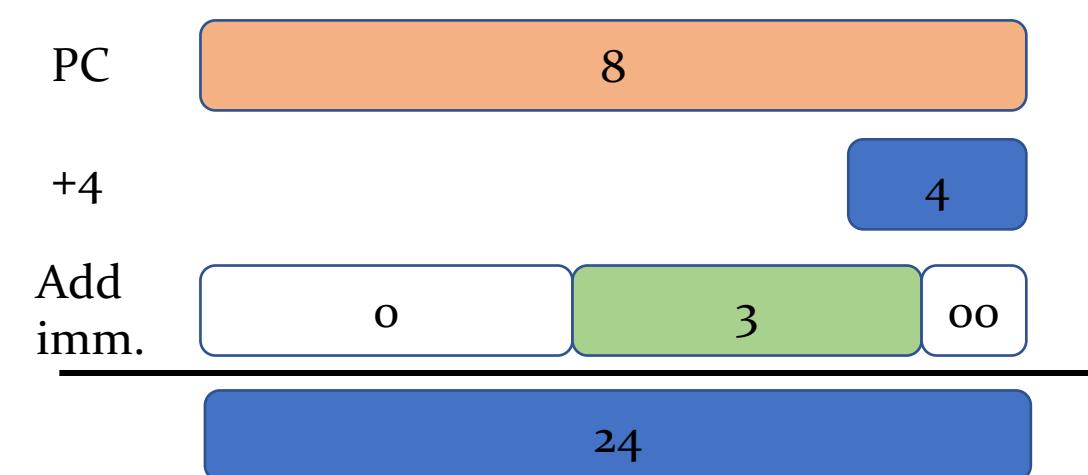
Branch instructions verses jump?
Why store PC+4?

First quiz
for (j=0; j<10; j++)
{
 b = b + j;
}

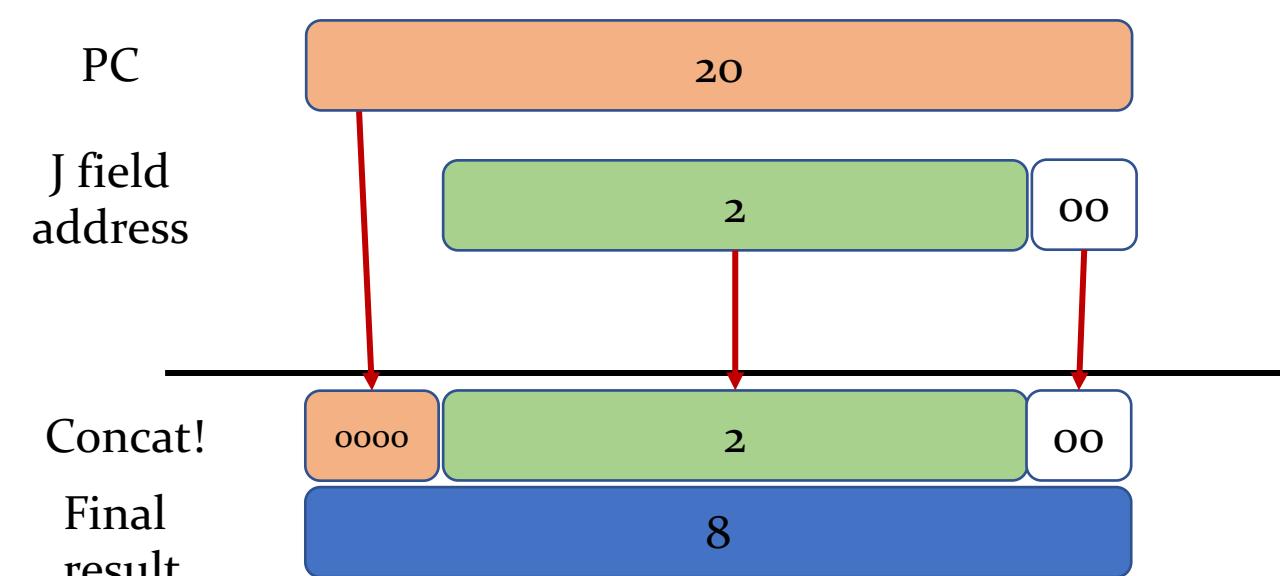
Address	Instruction	Comment
0		
4		
8		
12		
16		
20		
24	...	Done with loop

First quiz
for ($j=0; j<10; j++$)
{
 $b = b + j;$
}

Address	Instruction	Comment
0	Addi r5, ro, 0	$J = 0$
4	Addi r1, ro, 10	$R1 = 10$; exit condition
8	Beq r5, r1, 3	If j is 10, go to 24
12	Add r6, r6, r5	$b = b + j$
16	Addi r5, r5, 1	Update j , $j = j+1$
20	J 2	Goto 8
24	...	Done with loop



$$\text{Beq: } \text{PC} = 8 + 4 + (3 \ll 2) = 24$$



$$\text{Jump: } \text{PC} = \text{oooo} || 2 || \text{oo} = 8$$



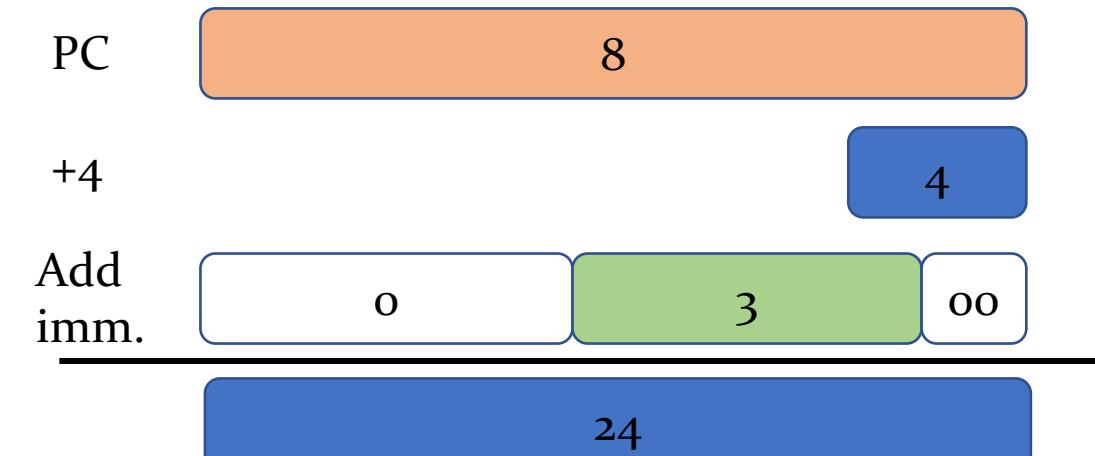
NYU

TANDON SCHOOL
OF ENGINEERING

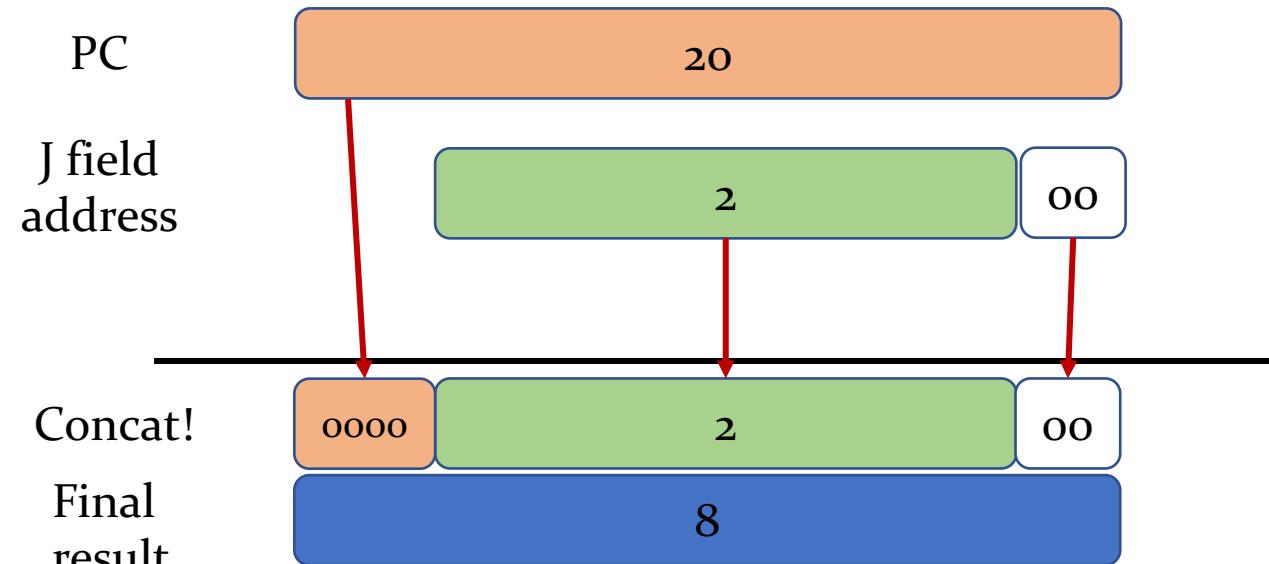
First quiz

Follow-up:

- 1) Why didn't we add 4 to the jump instruction?
- 2) What if I want to go further than 28 bits??
- 3) Do I seriously have to do this arithmetic???



$$\text{Beq: } \text{PC} = 8 + 4 + (3 \ll 2) = 24$$



$$\text{Jump: } \text{PC} = \text{oooo} || 2 || \text{oo} = 8$$



NYU

TANDON SCHOOL
OF ENGINEERING

Contrived example

Initialize
[li \$s0, 7
li \$s1, 3

Some compute
[sub \$s2, \$s0, \$s1

Jump prep
[move \$ao, \$s0
move \$a1, \$s1

Jump!
[jal compute_f1

Get return val
[move \$s3, \$vo

Compute..
[slt \$s4, \$s2, \$s3

Q: what happens if a function
calls another function..?

Procedure prep:
save all used s* regs

Free to compute!

Return prep:
save return value
restore s* regs
restore sp!

Return control!
(Jump)

compute_f1:

```
[ addi $sp, $sp, -12
sw $s0, 8($sp)
sw $s1, 4($sp)
sw $s2, 0($sp)
add $s0, $ao, $a1
addi $s1, $a1, 1
sub $s2, $ao, $s1
mult $s2, $s1
mflo $vo, $lo
lw $s0, 8($sp)
lw $s1, 4($sp)
lw $s2, 0($sp)
addi $sp, $sp, 12
jr $ra
```

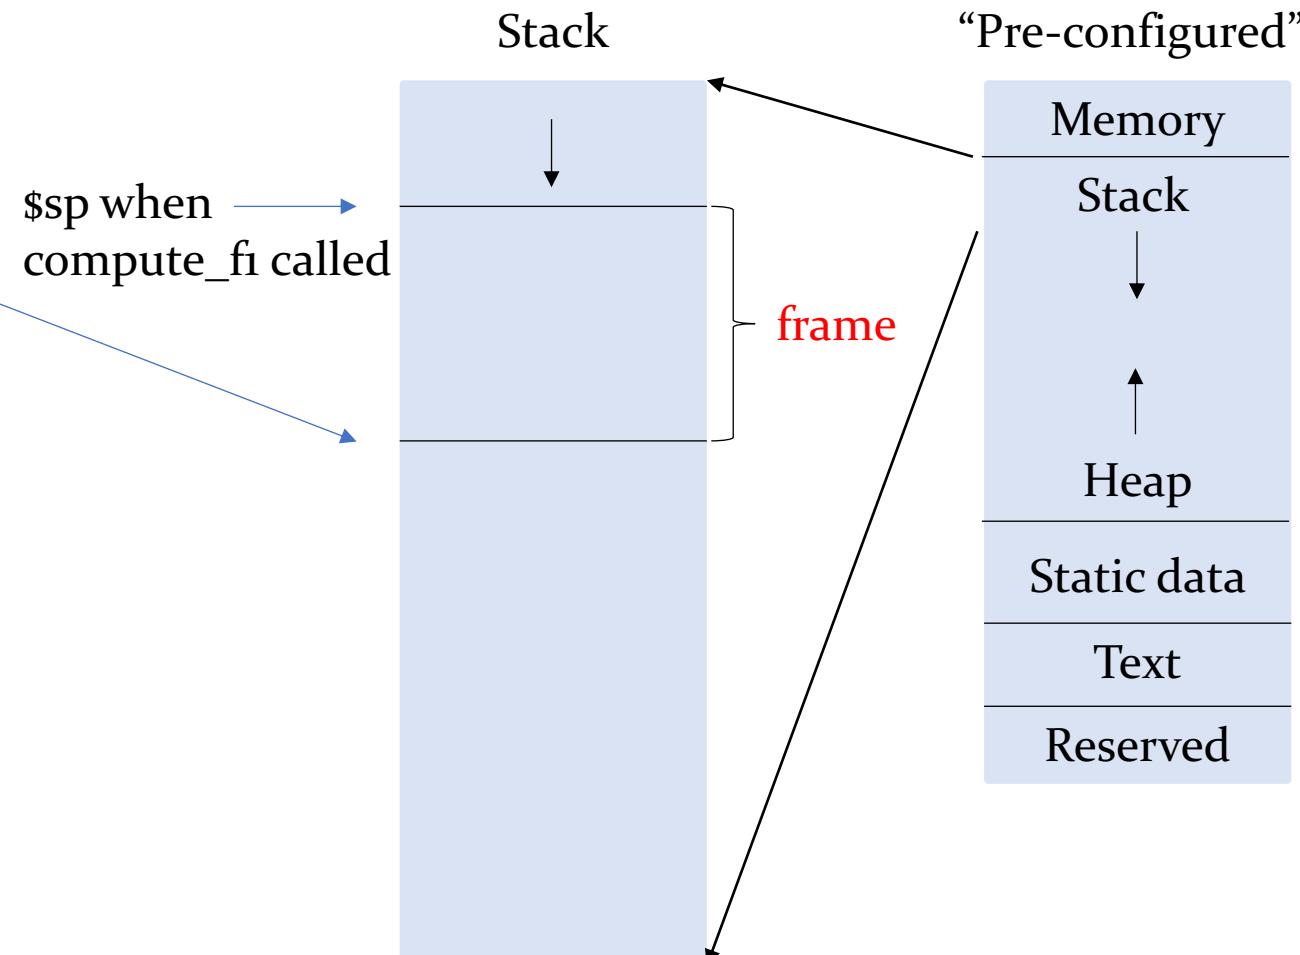
Visualizing the stack

compute_f1:

```
Setup      addi $sp, $sp, -12  
           sw $s0, 8($sp)  
           sw $s1, 4($sp)  
           sw $s2, 0($sp)
```

Teardown

```
.....  
       lw $s0, 8($sp)  
       lw $s1, 4($sp)  
       lw $s2, 0($sp)  
       addi $sp, $sp, 12  
       jr $ra
```



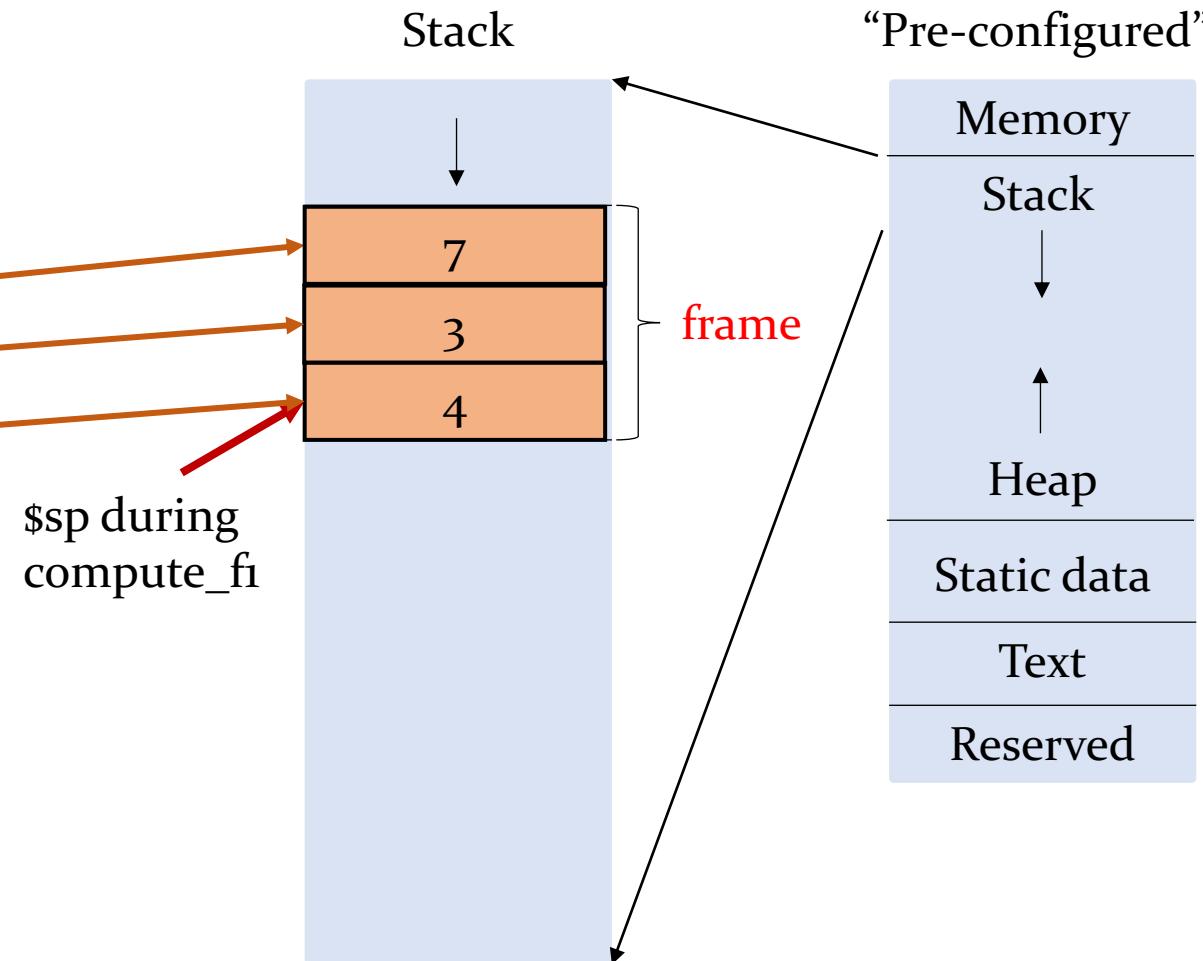
Visualizing the stack

compute_f1:

```
Setup      addi $sp, $sp, -12  
           sw $s0, 8($sp)  
           sw $s1, 4($sp)  
           sw $s2, 0($sp)
```

.....

```
Teardown   lw $s0, 8($sp)  
           lw $s1, 4($sp)  
           lw $s2, 0($sp)  
           addi $sp, $sp, 12  
           jr $ra
```

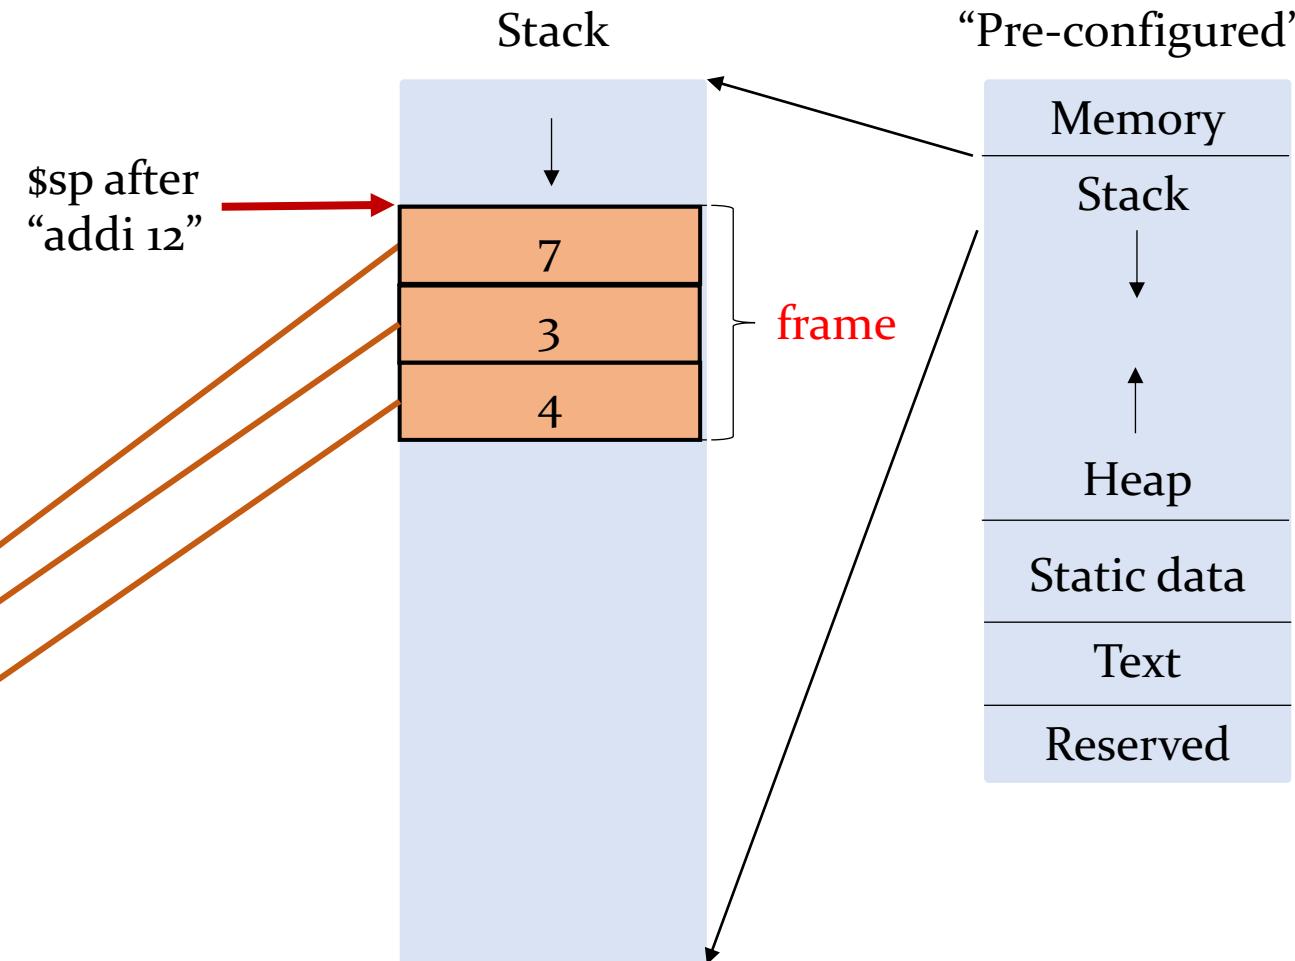


Visualizing the stack

compute_f1:

```
Setup: [ addi $sp, $sp, -12  
        sw $s0, 8($sp)  
        sw $s1, 4($sp)  
        sw $s2, 0($sp)
```


Teardown: [.....
 lw \$s0, 8(\$sp)
 lw \$s1, 4(\$sp)
 lw \$s2, 0(\$sp)
 addi \$sp, \$sp, 12
 jr \$ra



Q: what happens if a function calls another function..?

Instruction Set Architecture (ISA)

What is the architectural state of the machine?

- 1) Registers
- 2) Memory
- 3) Program counter

Are pipeline registers part of the architecture or micro architecture?

Why?

What about hazards?

Why is the distinction so important?

MIPS Summary

32b RISC architecture

- Fixed instruction length
- What if your simulators supported variable length?
- 32b registers and memory addresses
- Byte addressable memory.. How much?

Load-store architecture

- Only special instructions access memory
- Think about your simulators and the control signals
- Now, what if any instruction could access memory?

32 general purpose registers

- We will assume the simplified calling convention
- Program counter for finding next instruction
- Hi/lo registers for multiplication
- What if we had 128?

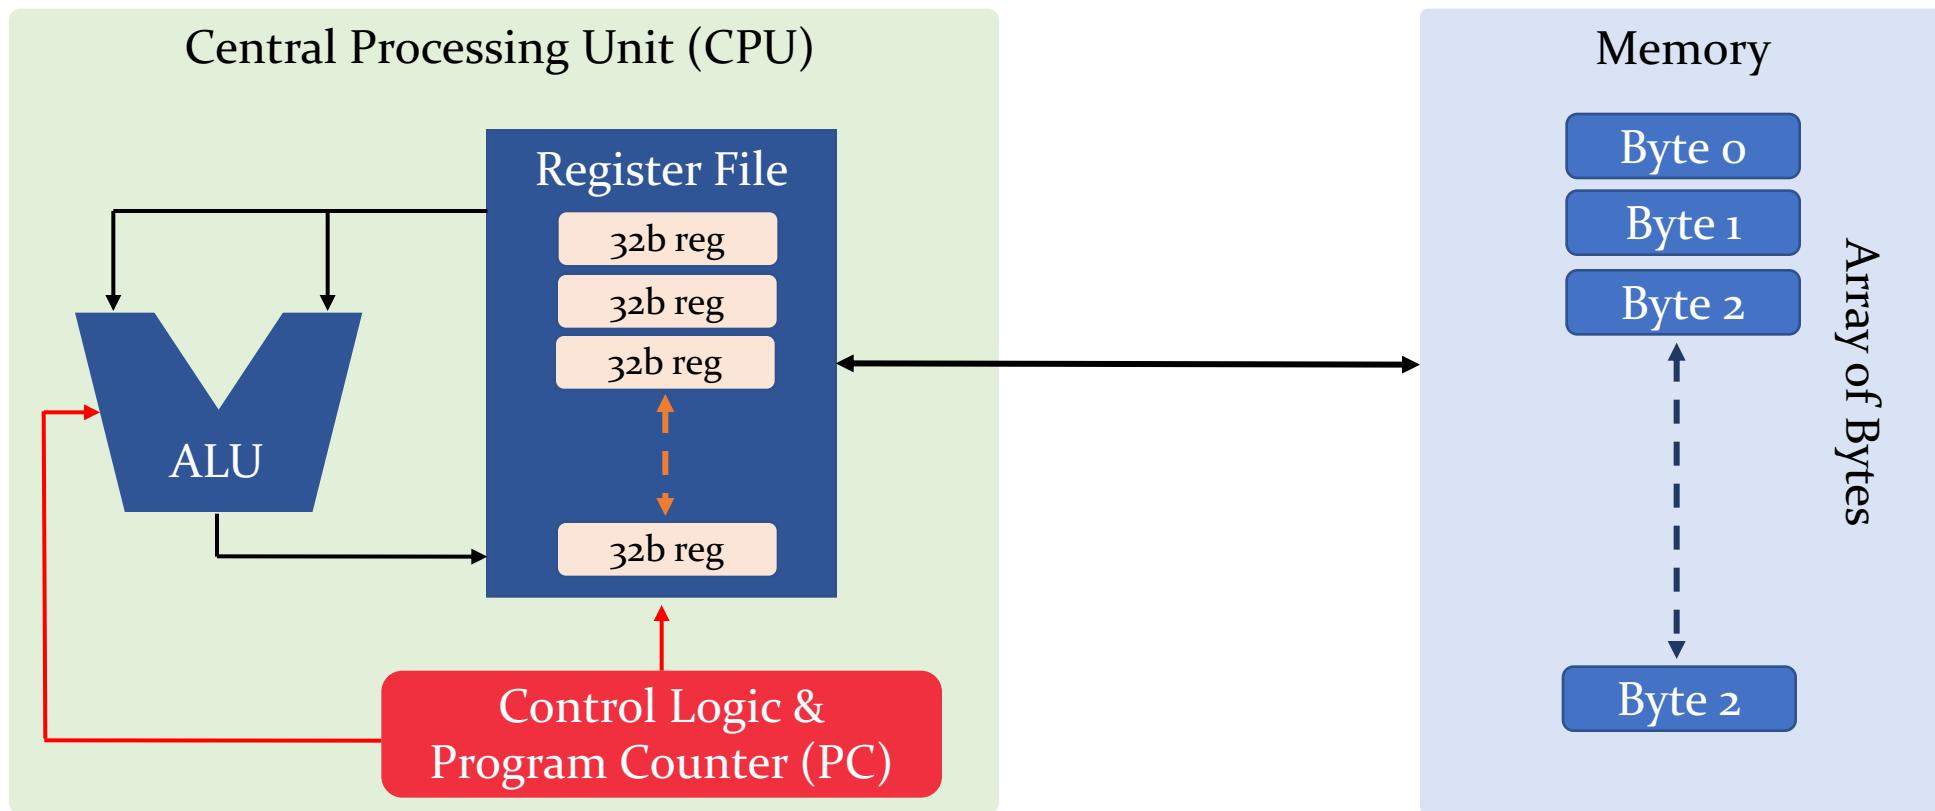
Design principle: smaller is faster.



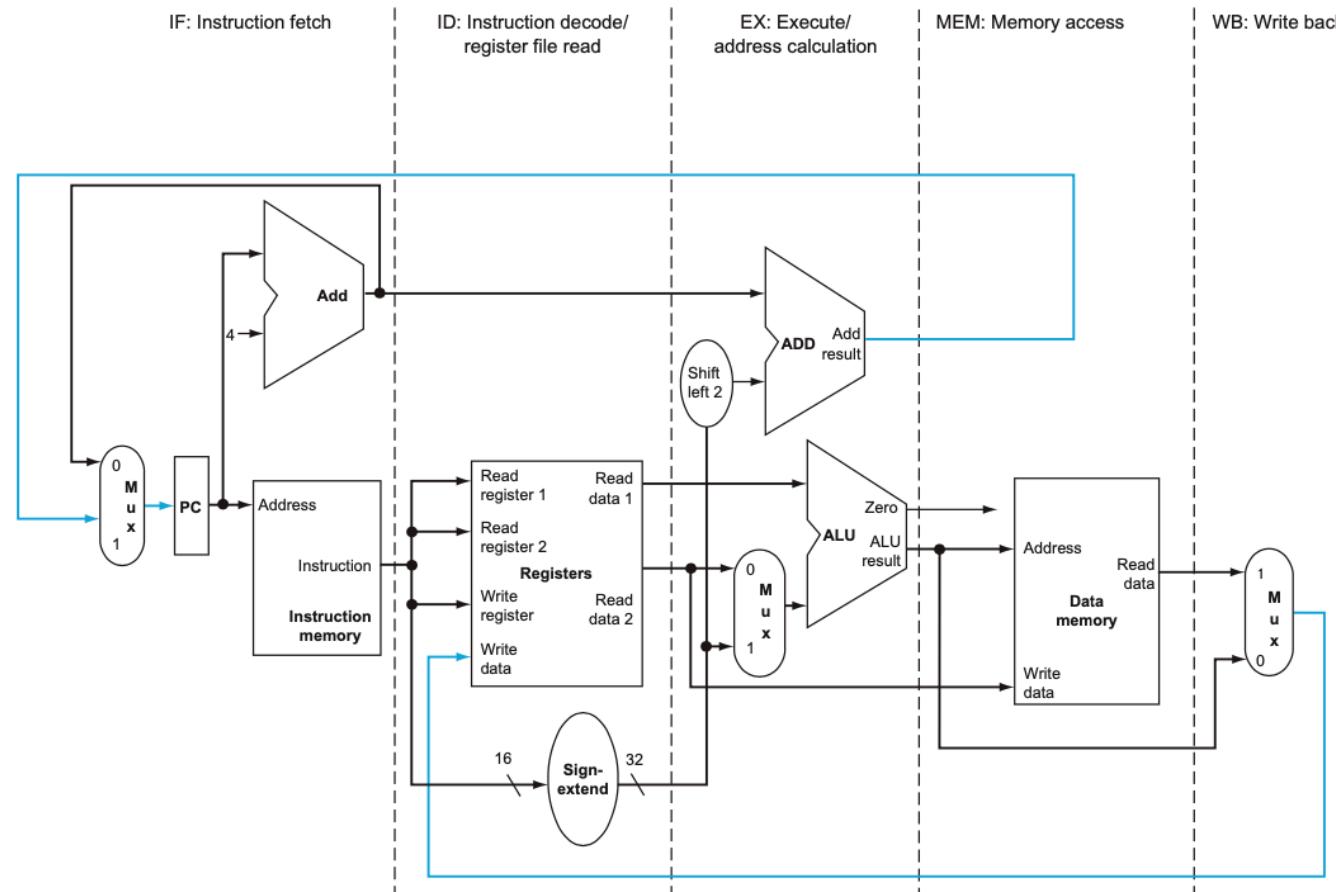
NYU

TANDON SCHOOL
OF ENGINEERING

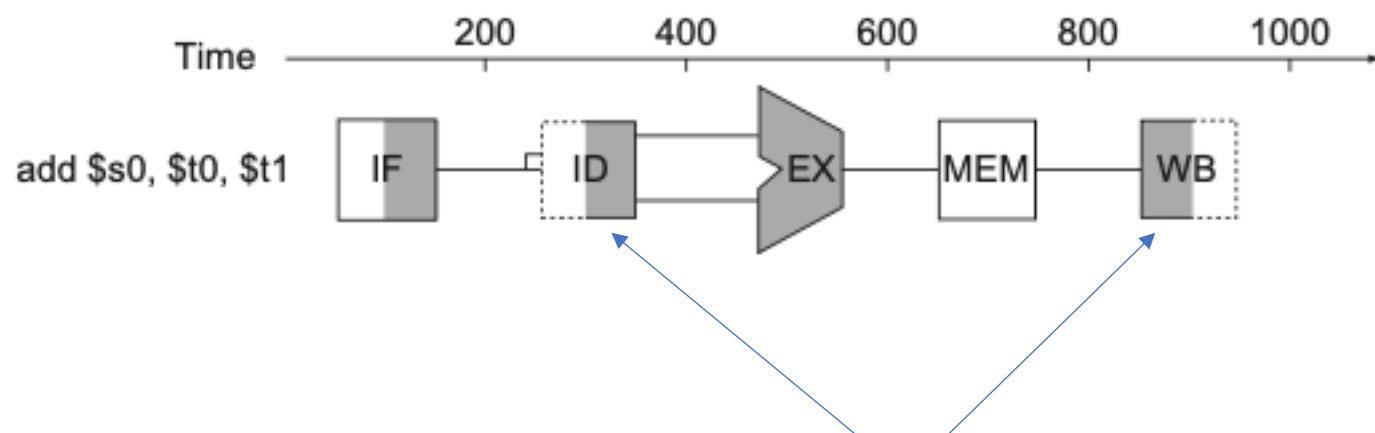
Architectural (ISA) view of a processor



Basic (datapath) pipeline machine



Basic pipeline diagram

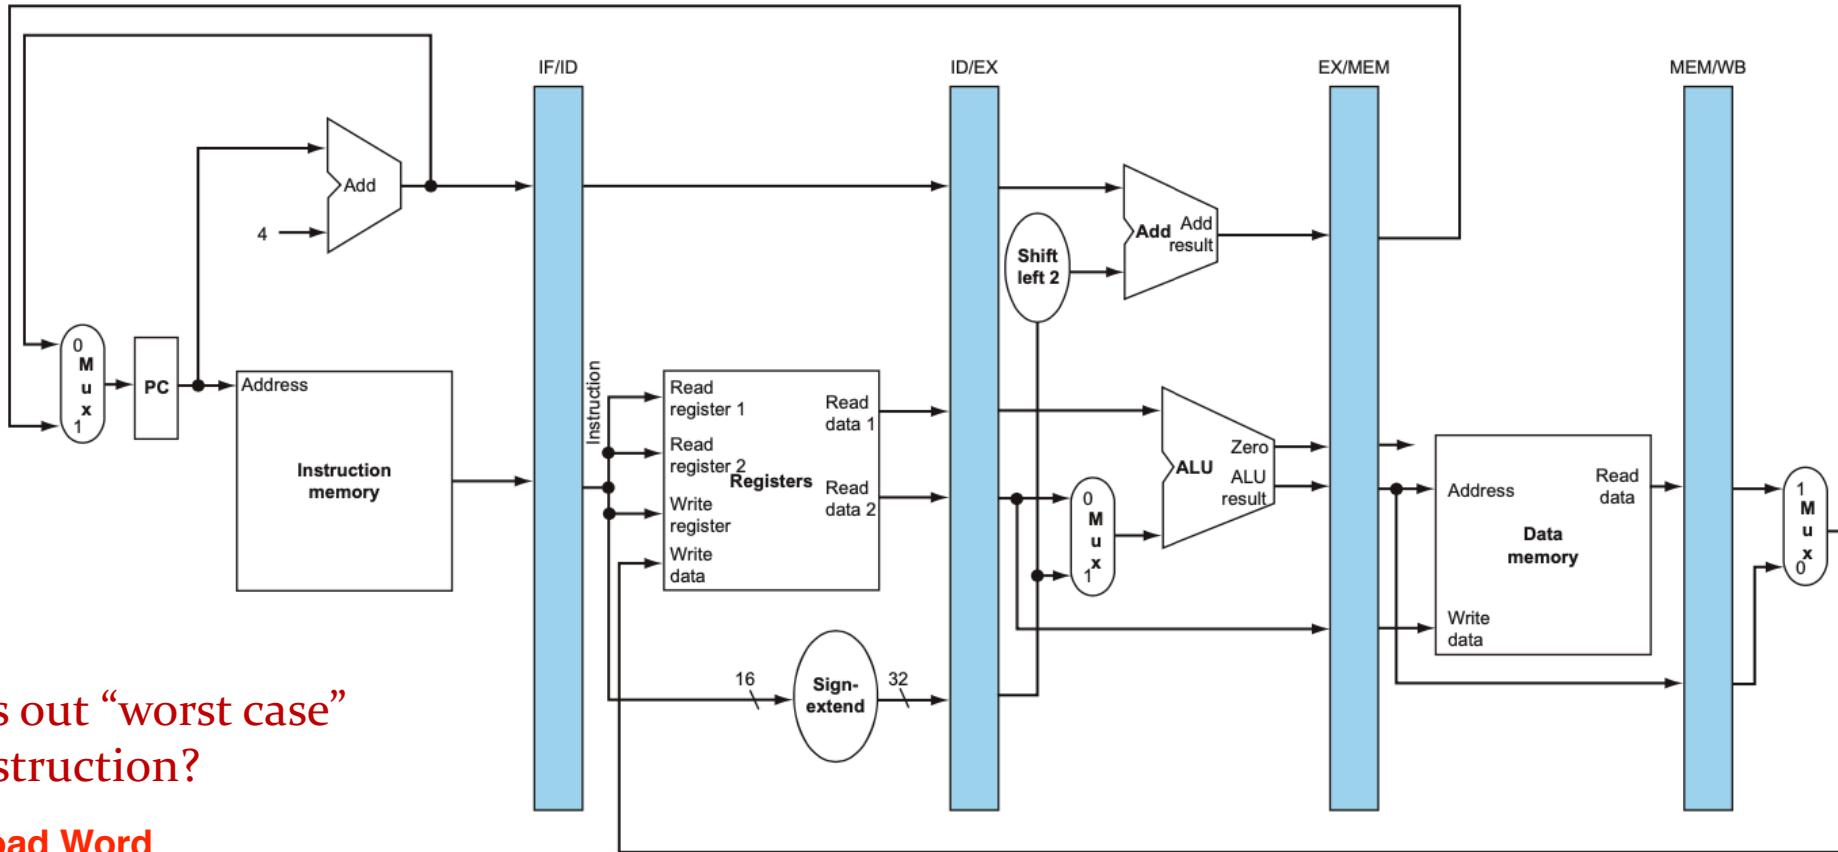


Architectural state written in first half
and read in second half of clock cycle!

More detailed pipeline diagram
Shaded sections imply use
Note: WB and ID both using RF

Question: what's the problem here
and how do we avoid it?

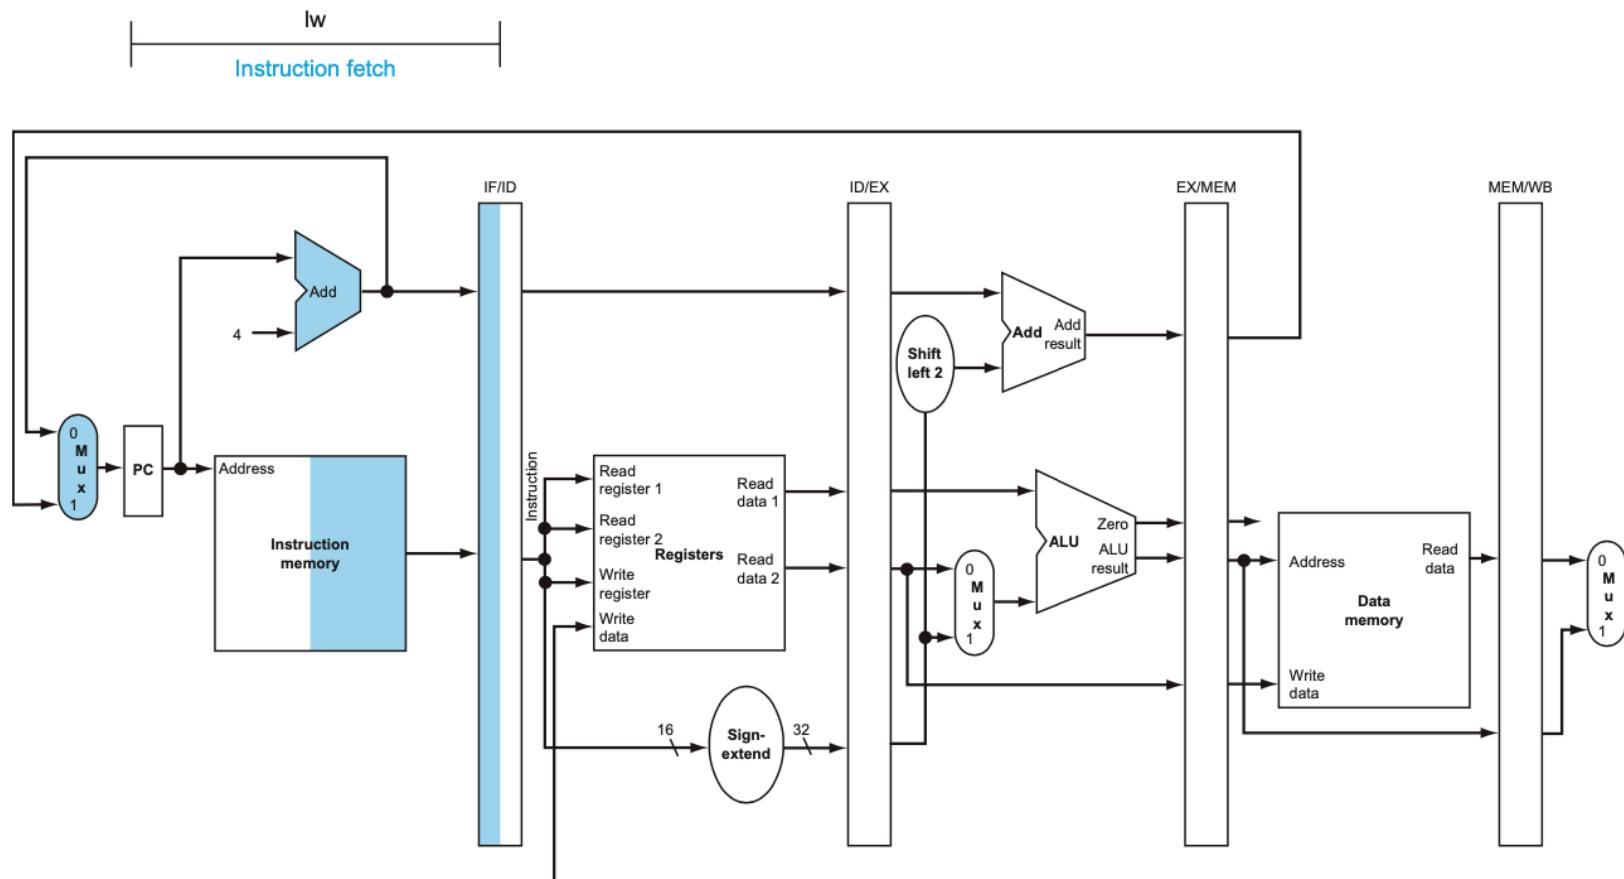
Pipelined datapath



NYU

TANDON SCHOOL
OF ENGINEERING

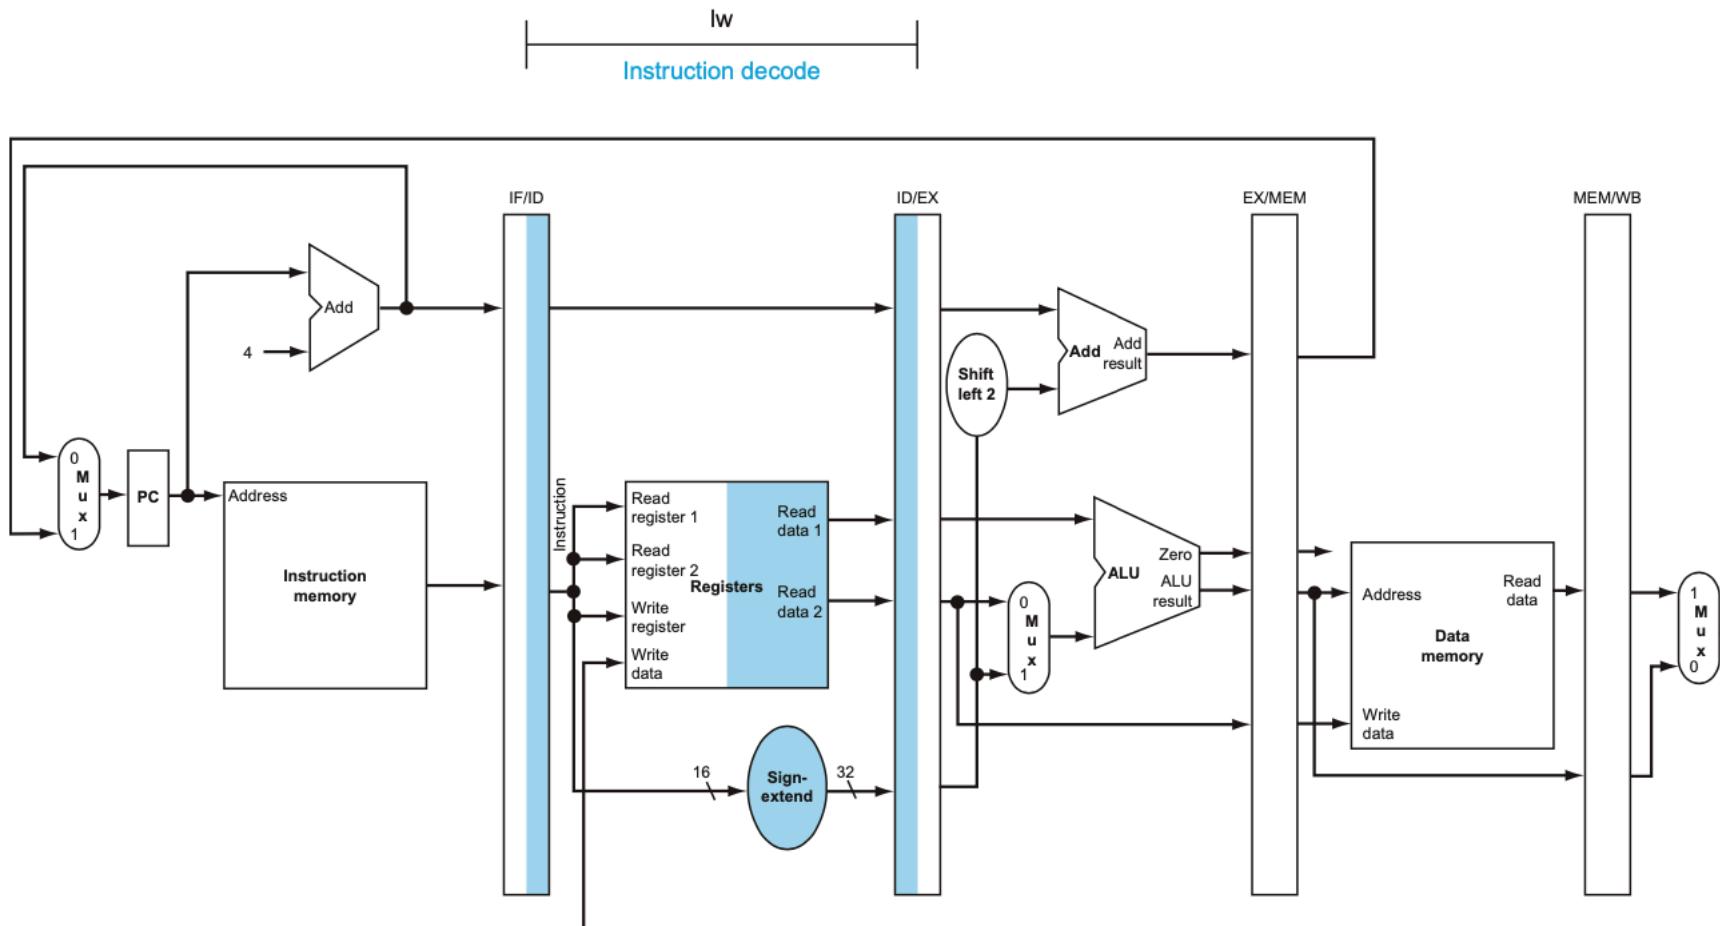
Pipeline: IF



NYU

TANDON SCHOOL
OF ENGINEERING

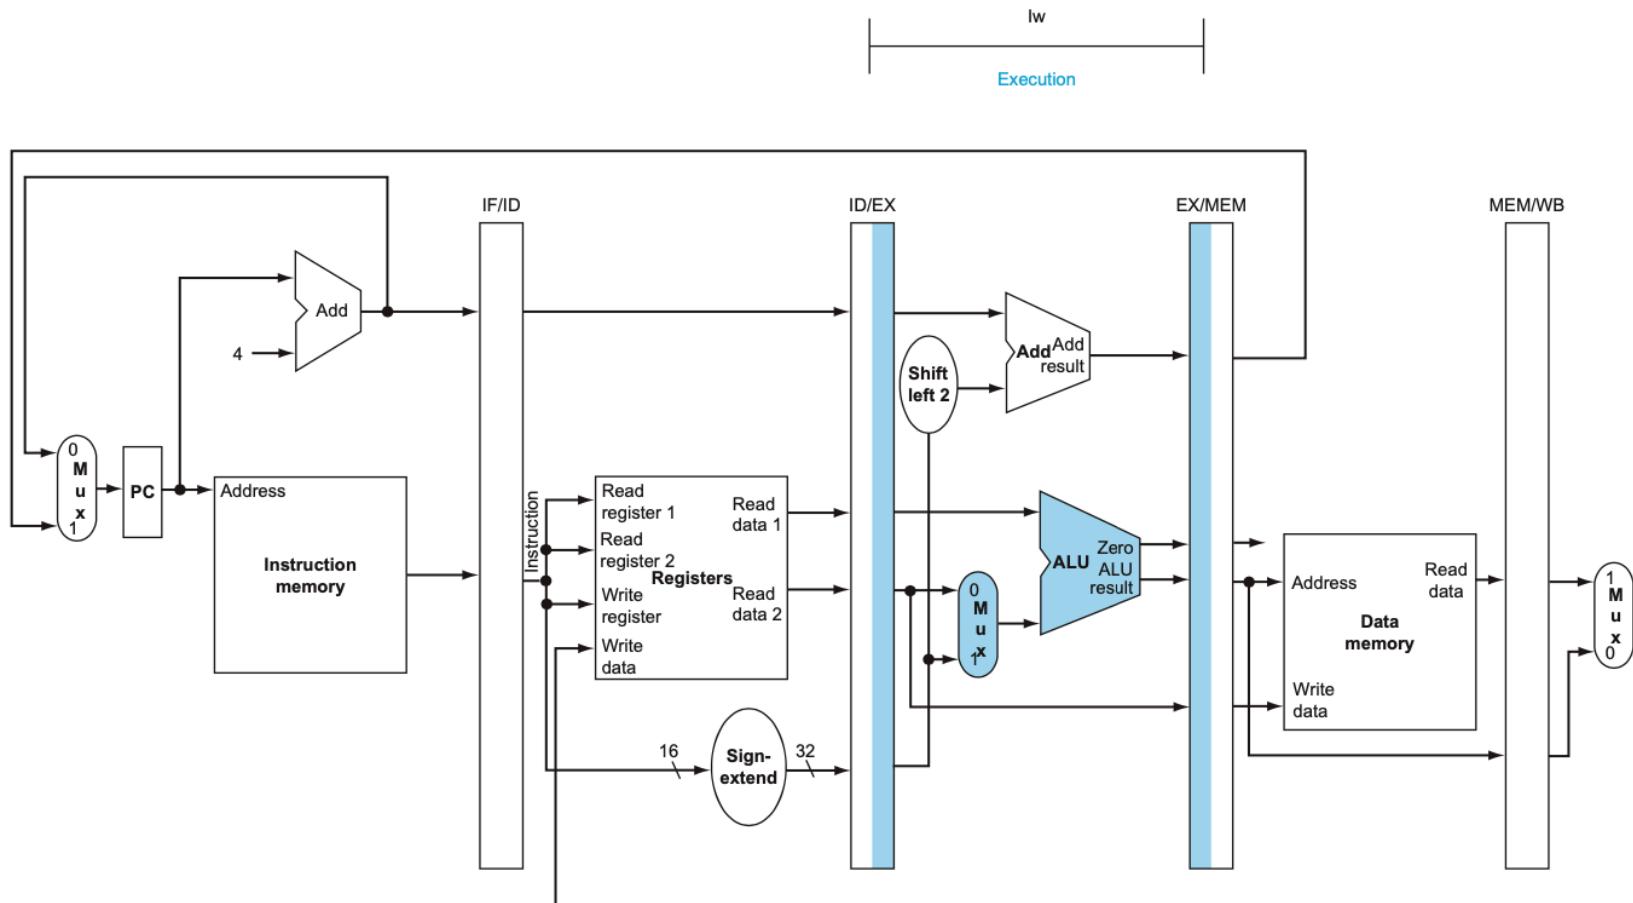
Pipeline: ID



NYU

TANDON SCHOOL
OF ENGINEERING

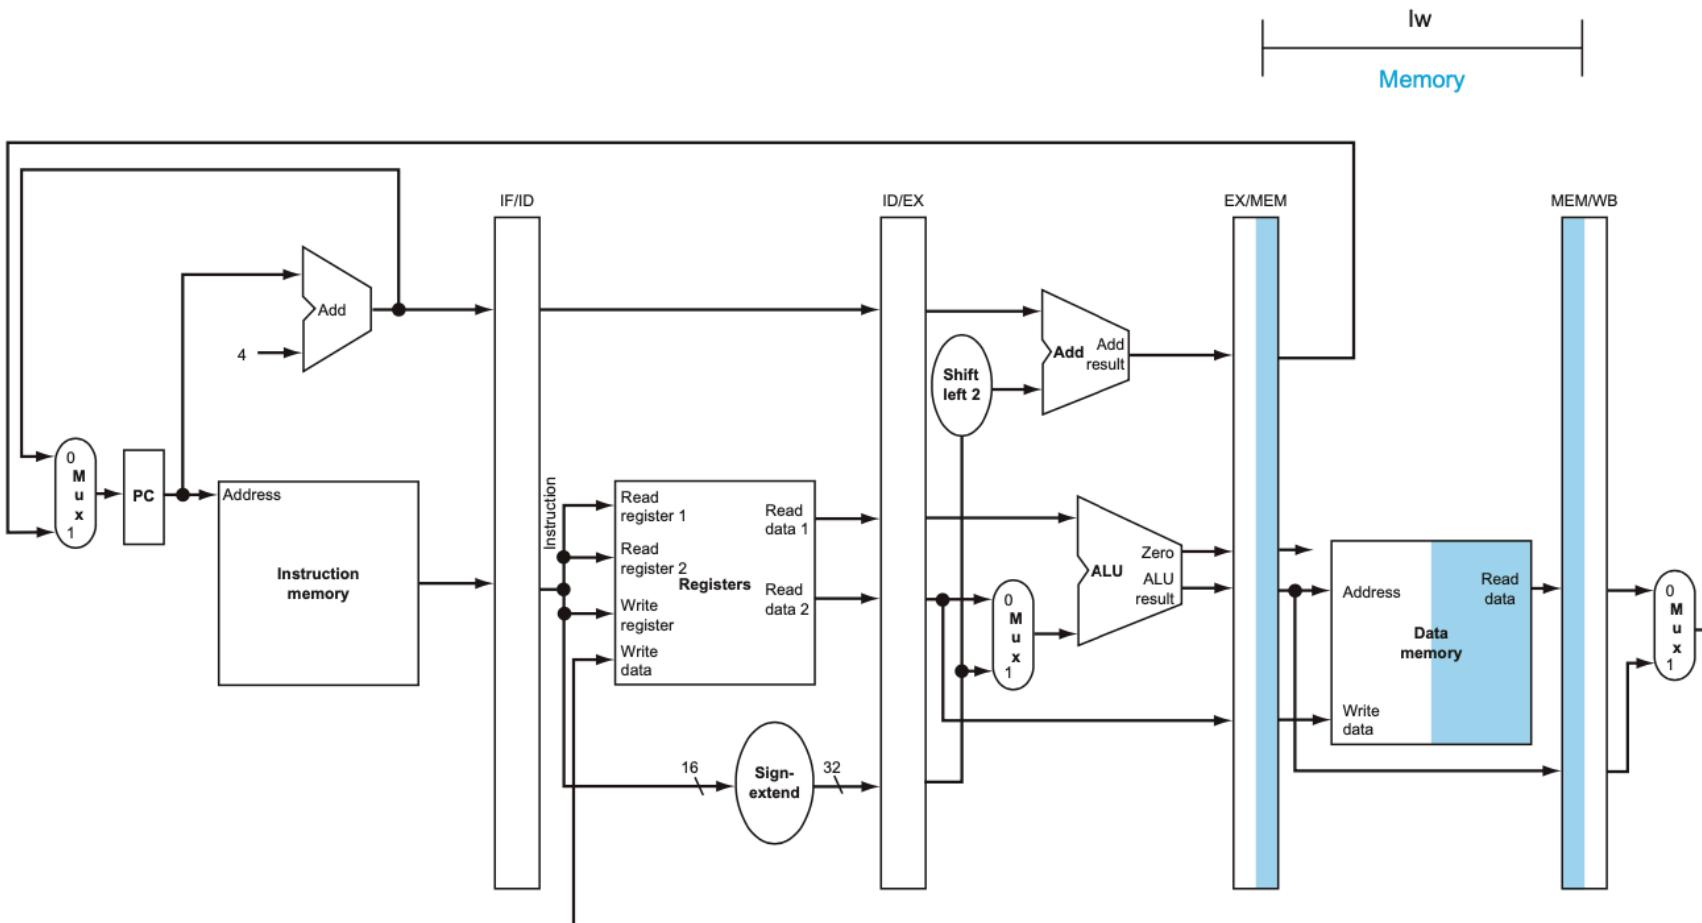
Pipeline: EX



NYU

TANDON SCHOOL
OF ENGINEERING

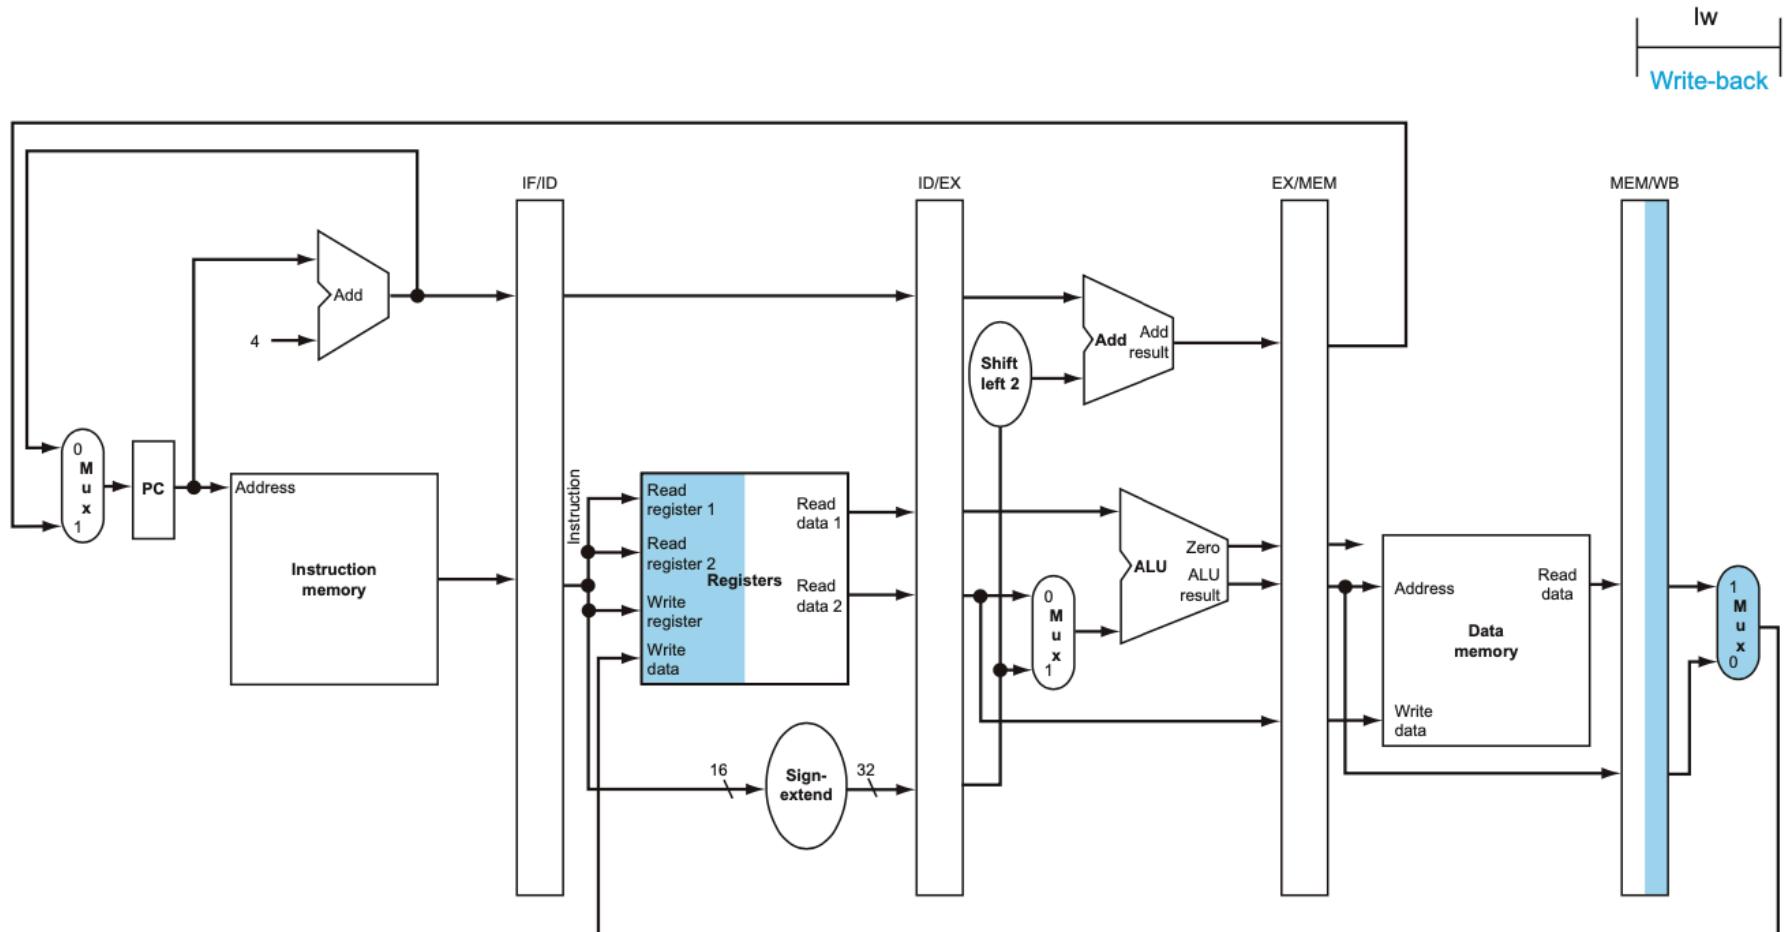
Pipeline: Mem



NYU

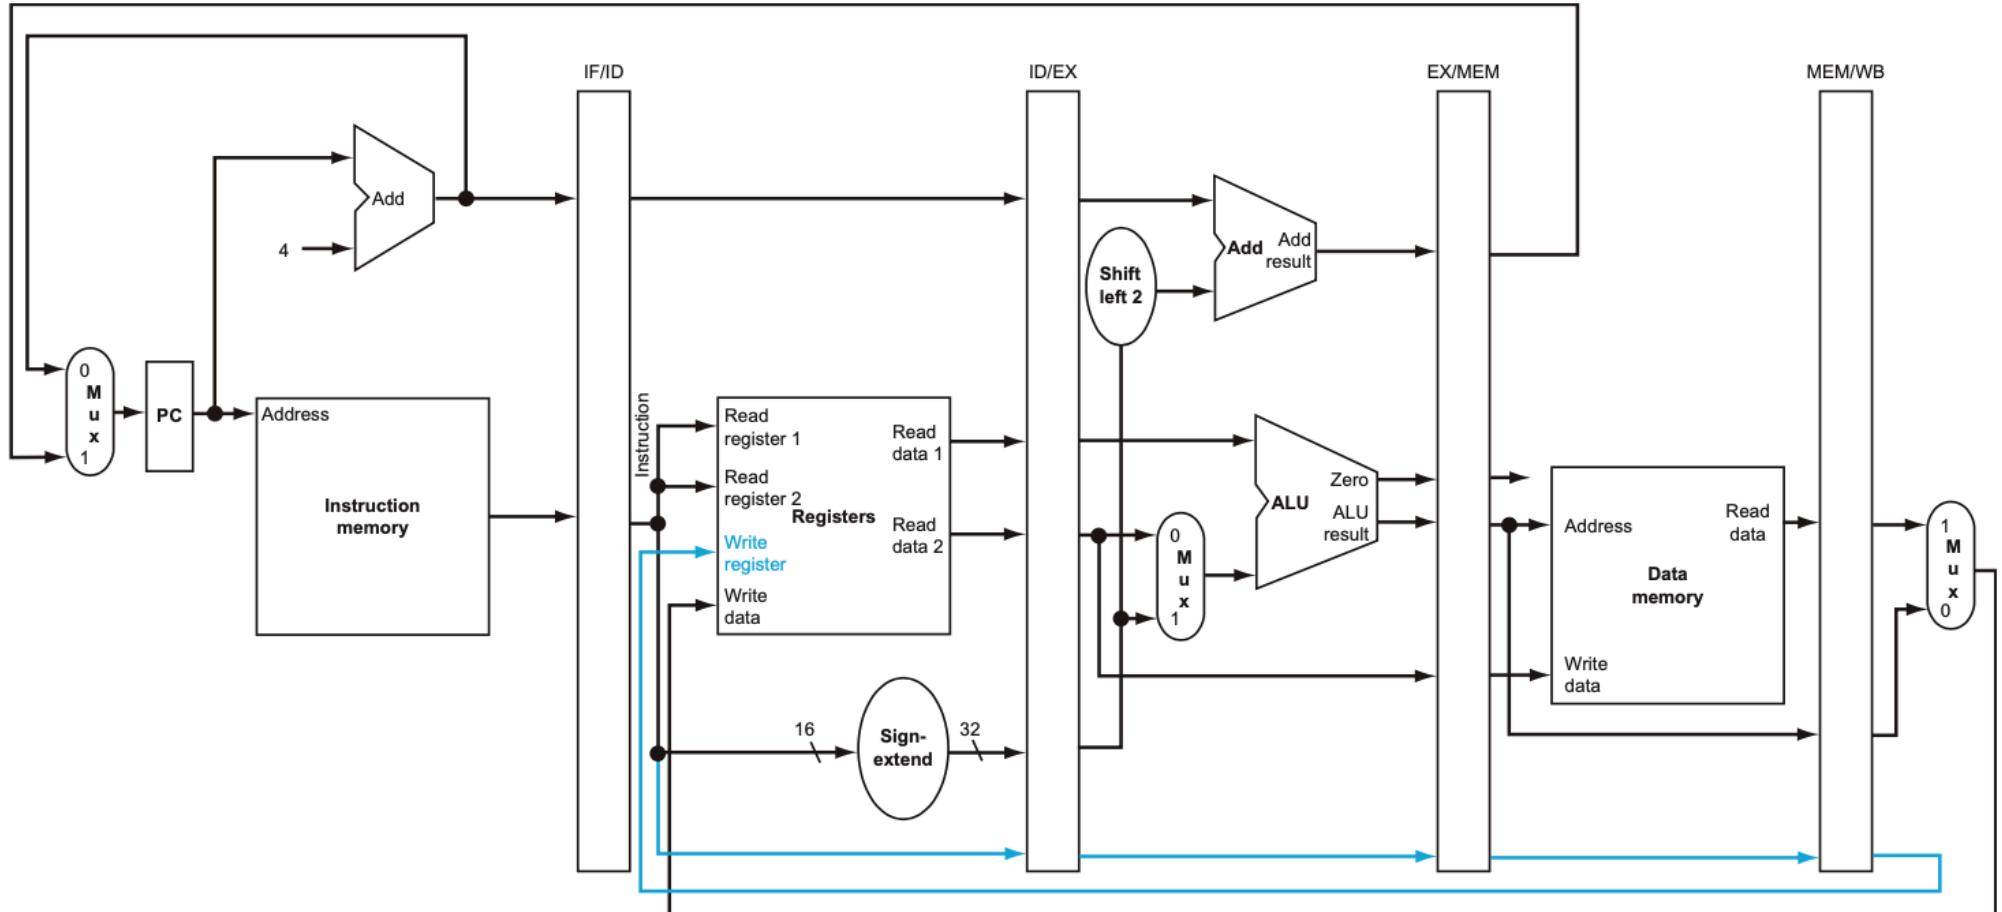
TANDON SCHOOL
OF ENGINEERING

Pipeline: WB



What did we forget?

Where to writeback!



NYU

TANDON SCHOOL
OF ENGINEERING

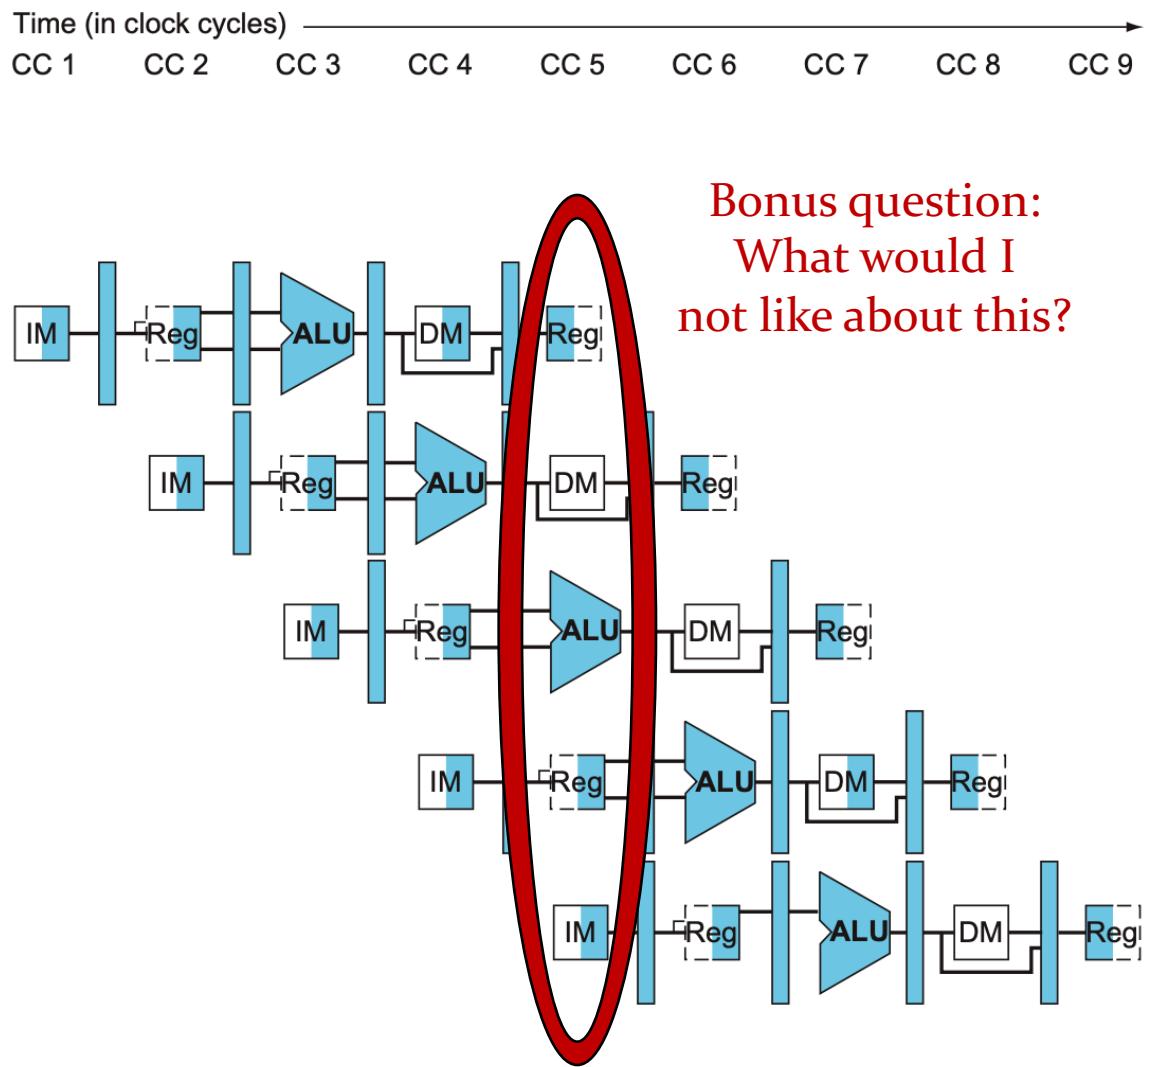
Pipeline diagram

lw	\$10, 20(\$1)
sub	\$11, \$2, \$3
add	\$12, \$3, \$4
lw	\$13, 24(\$1)
add	\$14, \$5, \$6

What's great about this program?

Program execution order (in instructions)

- lw \$10, 20(\$1)
- sub \$11, \$2, \$3
- add \$12, \$3, \$4
- lw \$13, 24(\$1)
- add \$14, \$5, \$6



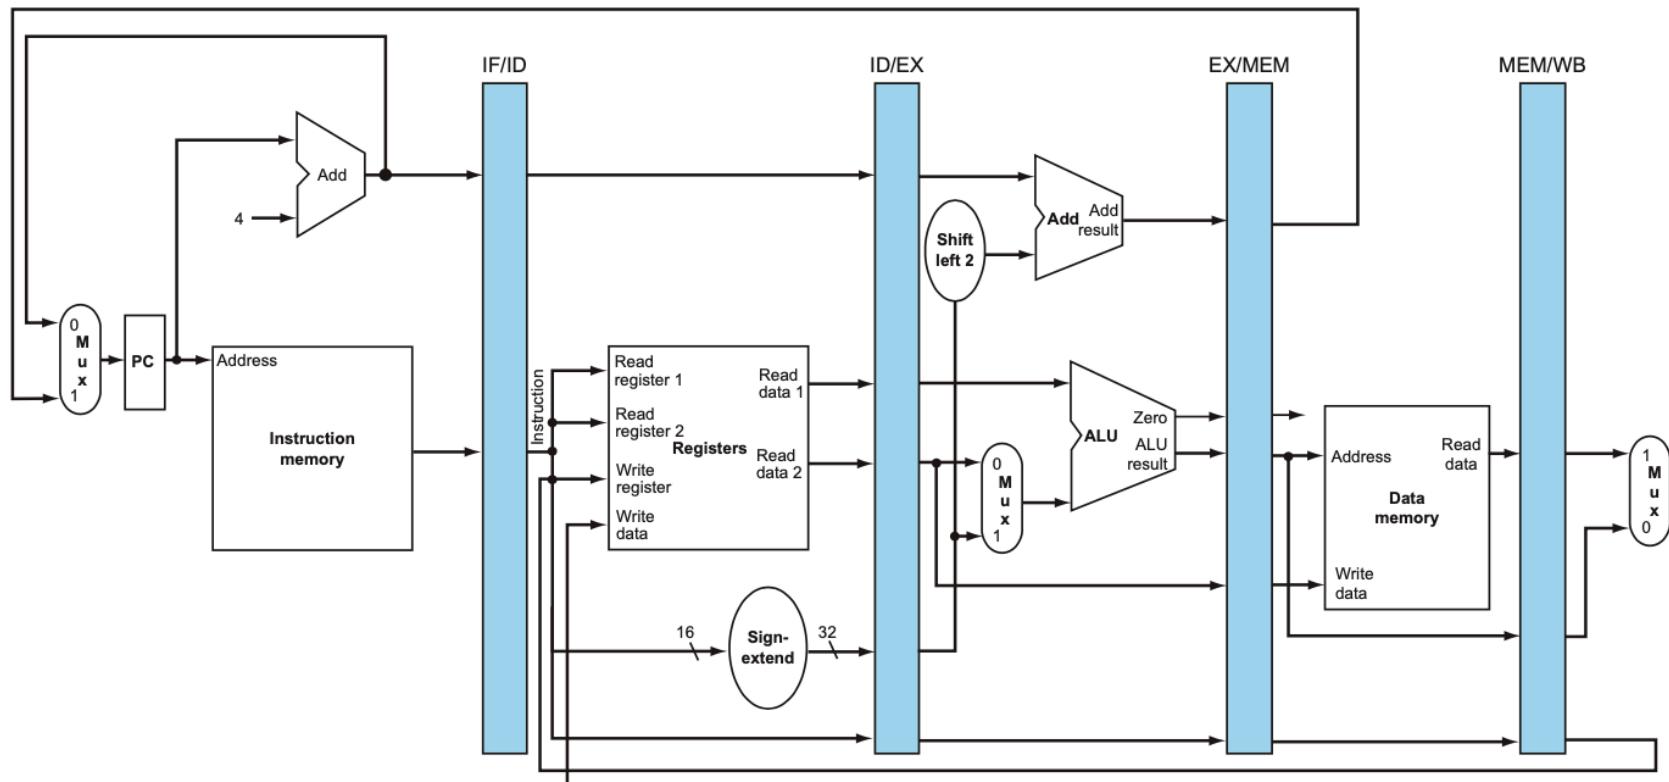
Bonus question:
What would I
not like about this?



NYU

TANDON SCHOOL
OF ENGINEERING

Visualizing cycle (cc) 5

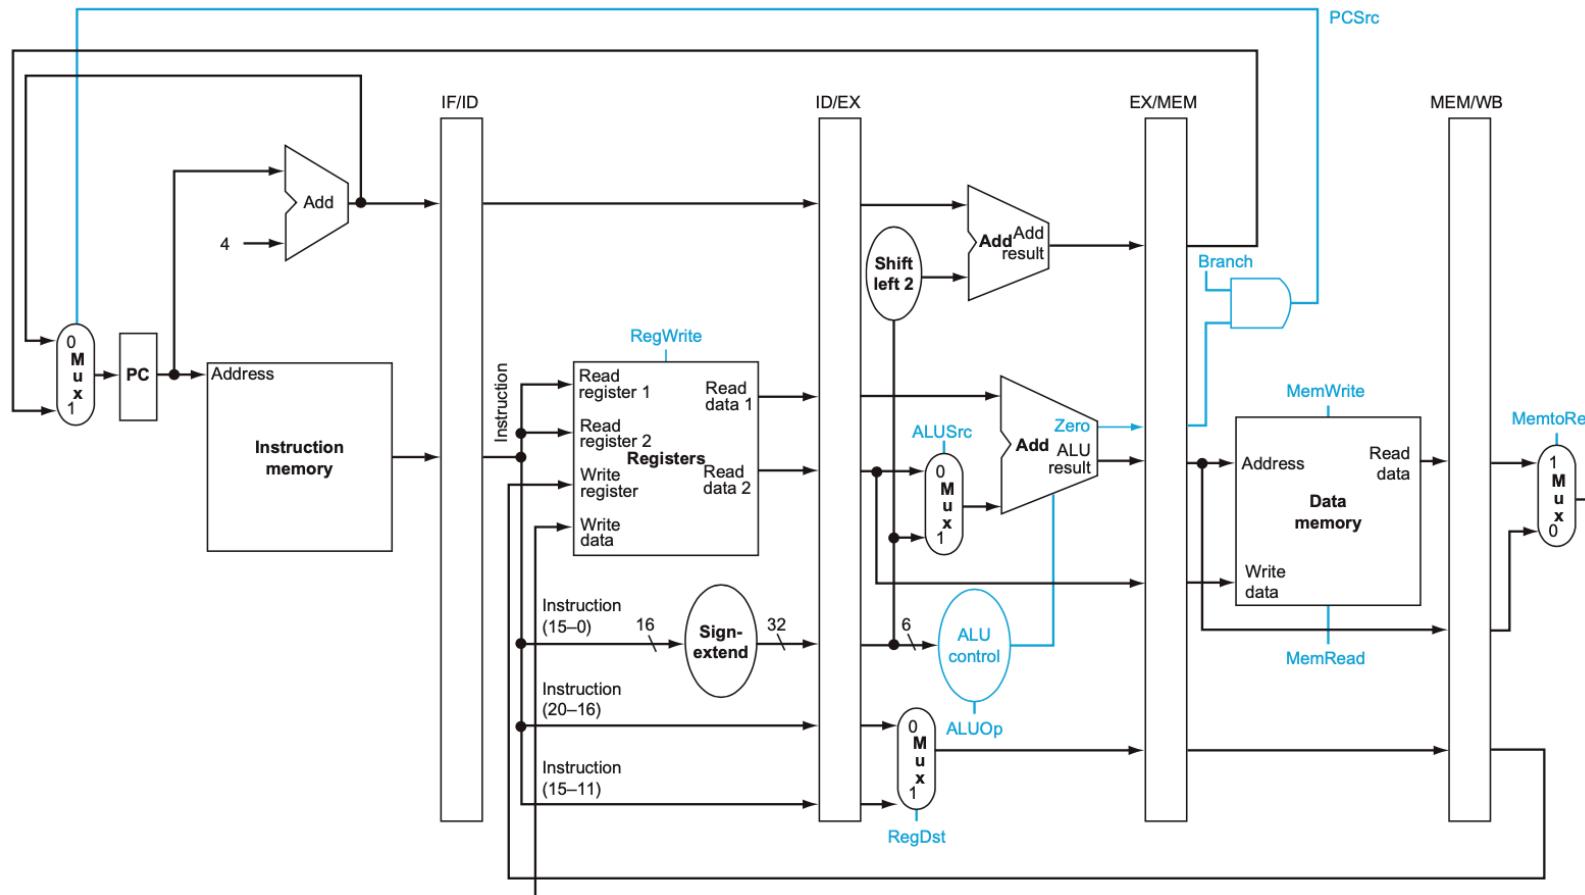


Why is pipelining so powerful?

What happens to:

- Instruction latency?
- Instruction throughput?
- Overall performance?
- Hardware utilization?
- Chip area?
- Chip power?
- Chip energy?

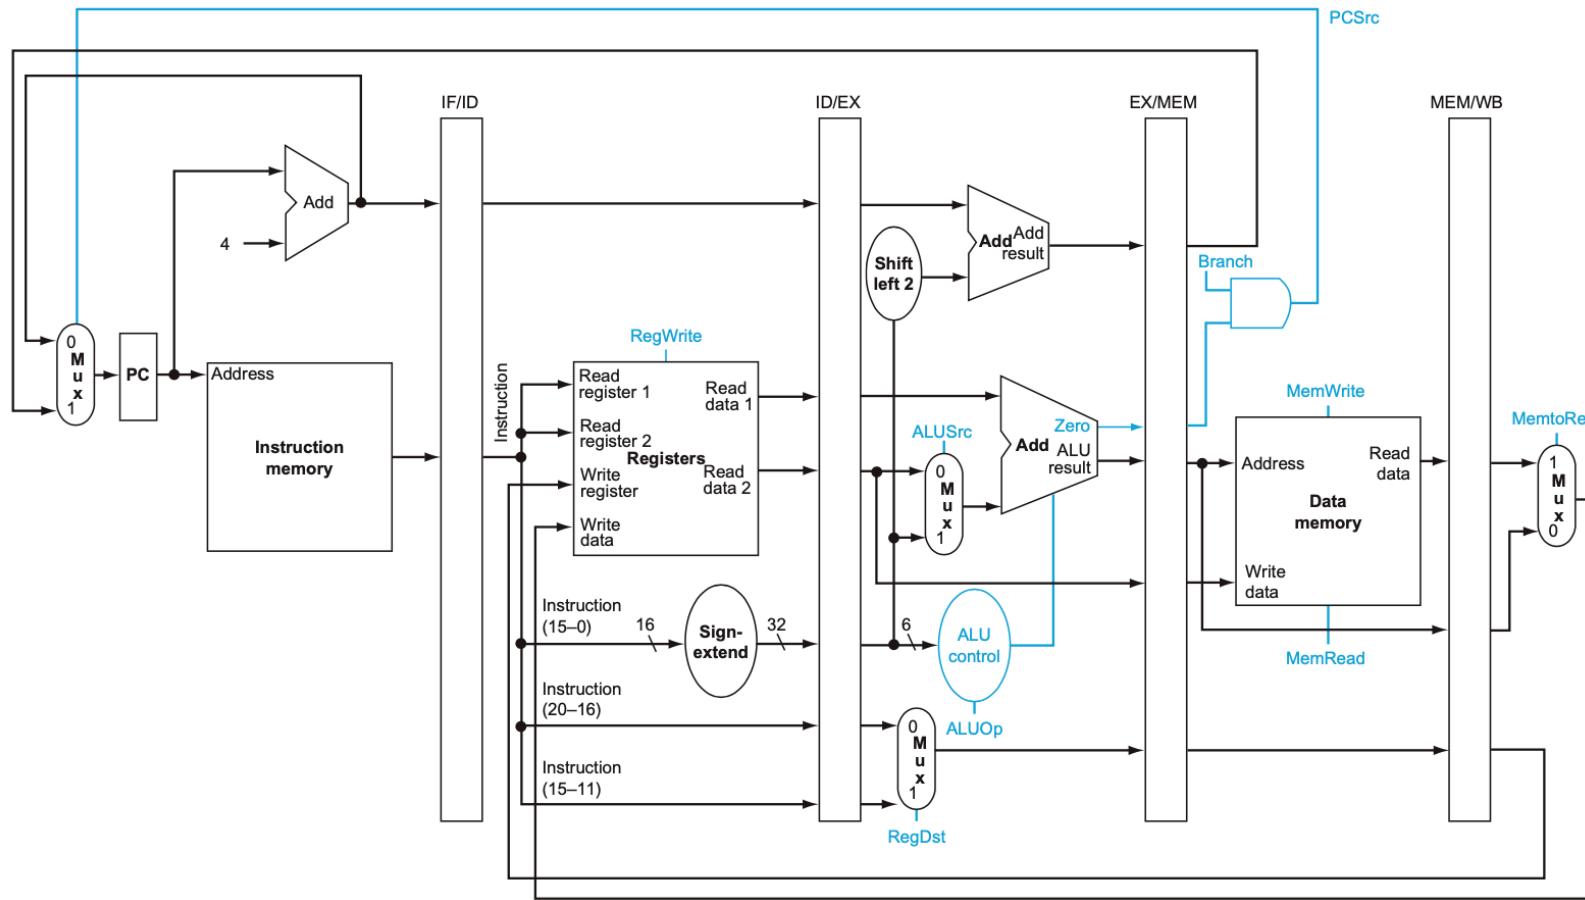
Pipelining control: when are signals needed?



IF stage:
None.
Always read PC memory.

ID stage:
None.
Same thing every cycle.

Pipelining control: when are signals needed?



EXE stage:

- 1) regDest
- 2) ALUop
- 3) ALUsrc

regDst: where to write value

ALUop: what ALU should do

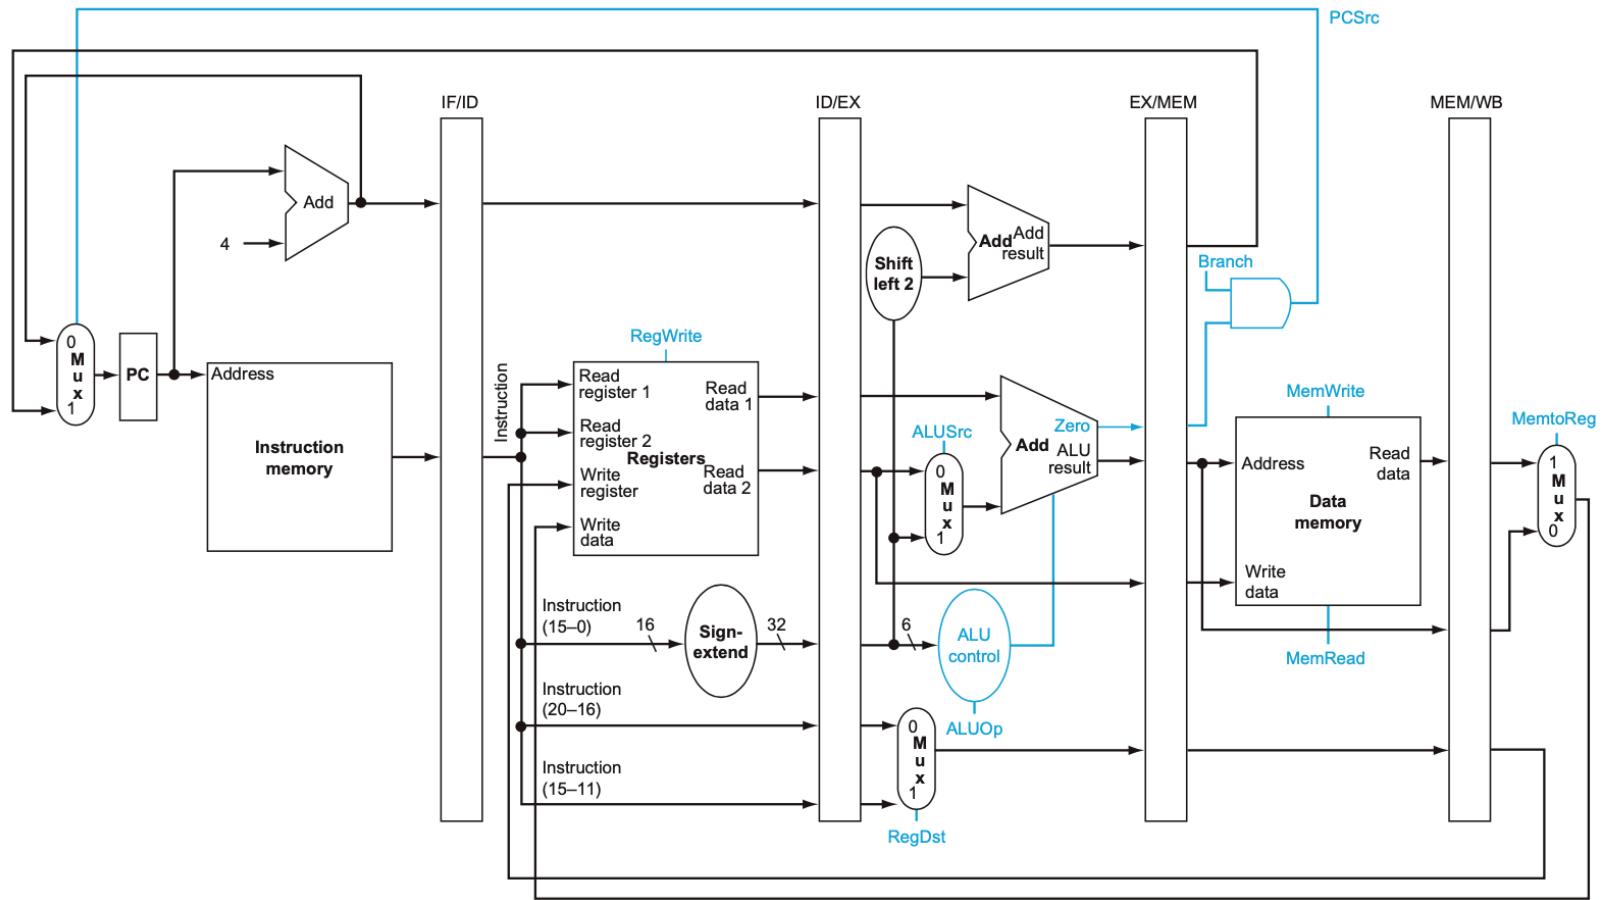
ALUsrc: reg2 or immediate value



NYU

TANDON SCHOOL
OF ENGINEERING

Pipelining control: when are signals needed?



MEM stage:

- 1) Branch
- 2) MemRd
- 3) MemWr

Branch: is branch instruction

MemRd: are we reading memory

MemWr: are we writing memory

Generates PCSrc signal to select next instruction.

What's going on here?

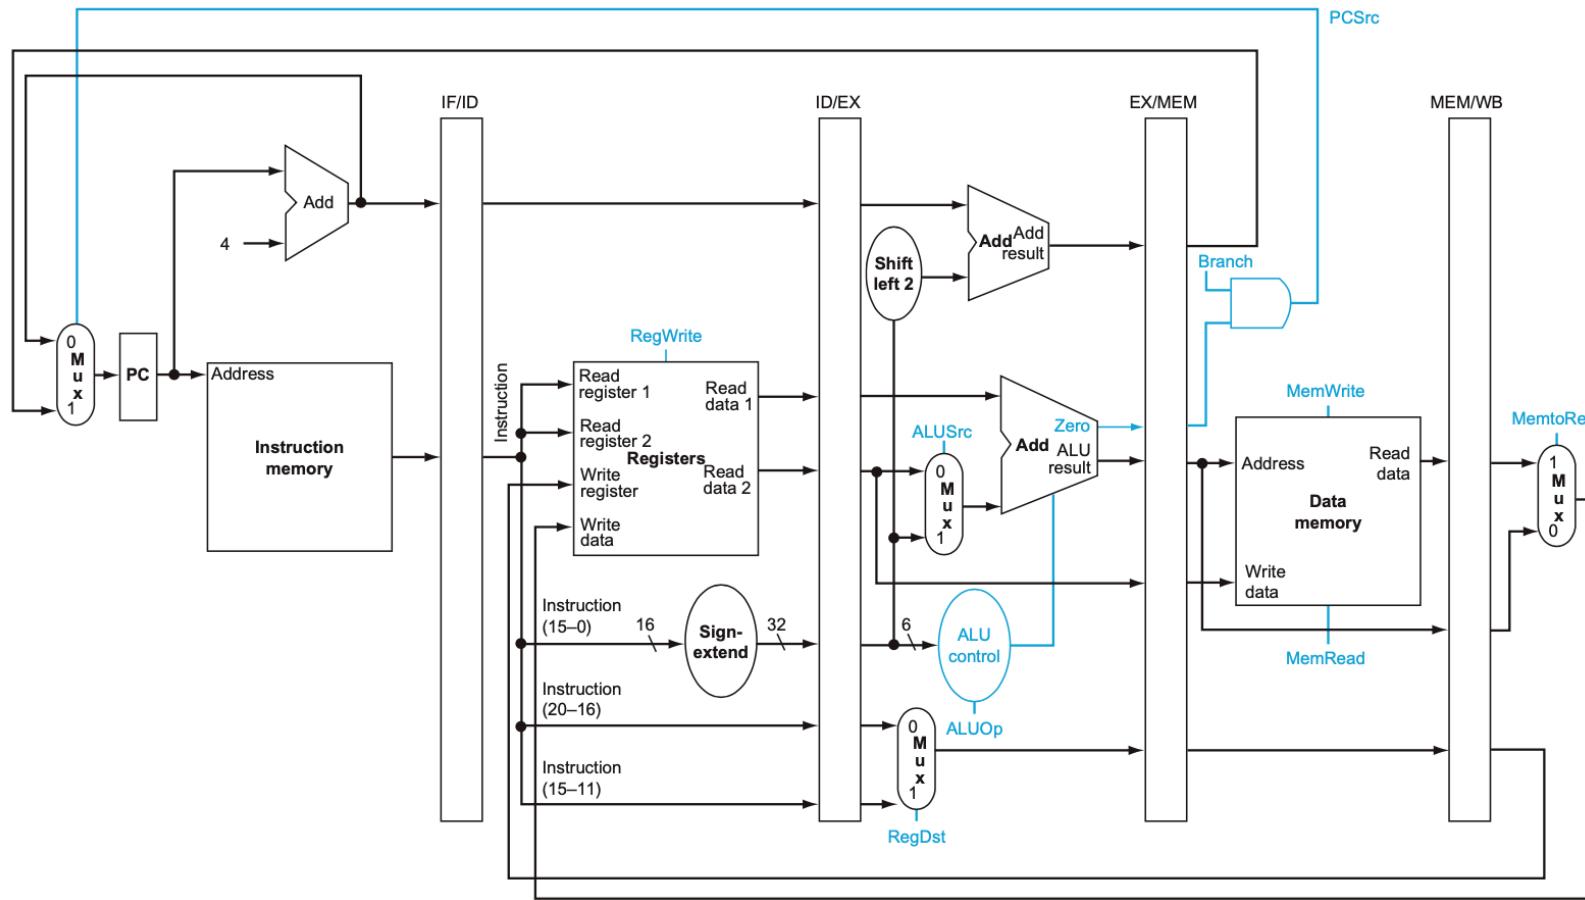
We'll talk about
branches in depth
in another class.



NYU

TANDON SCHOOL
OF ENGINEERING

Pipelining control: when are signals needed?



WB stage:
1) MemtoReg
2) RegWrite

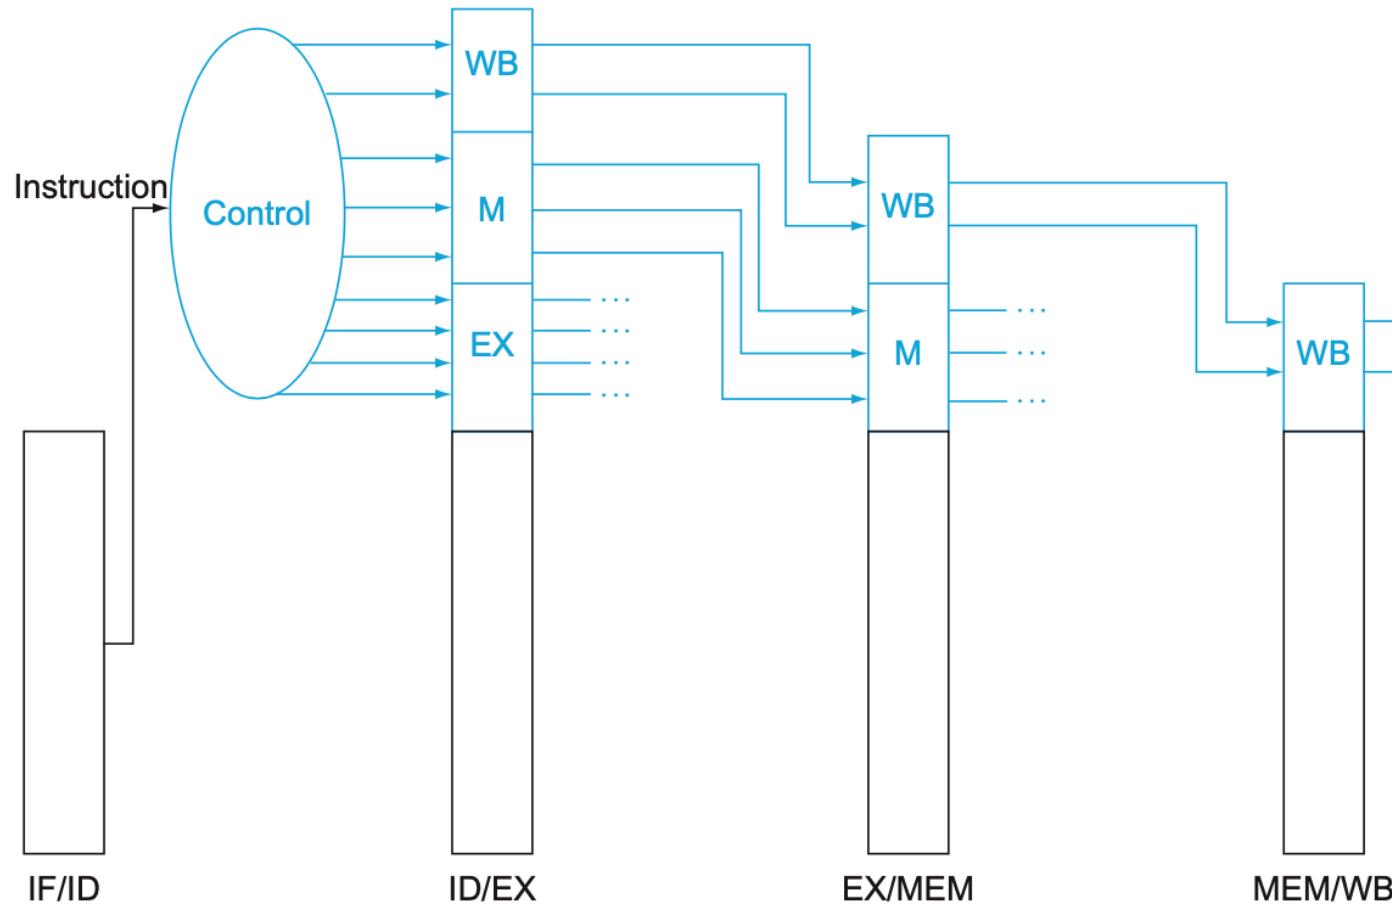
MemtoReg: which value to wr. back
RegWrite: enable RF writing



NYU

TANDON SCHOOL
OF ENGINEERING

Computing “Control lines” in ID stage

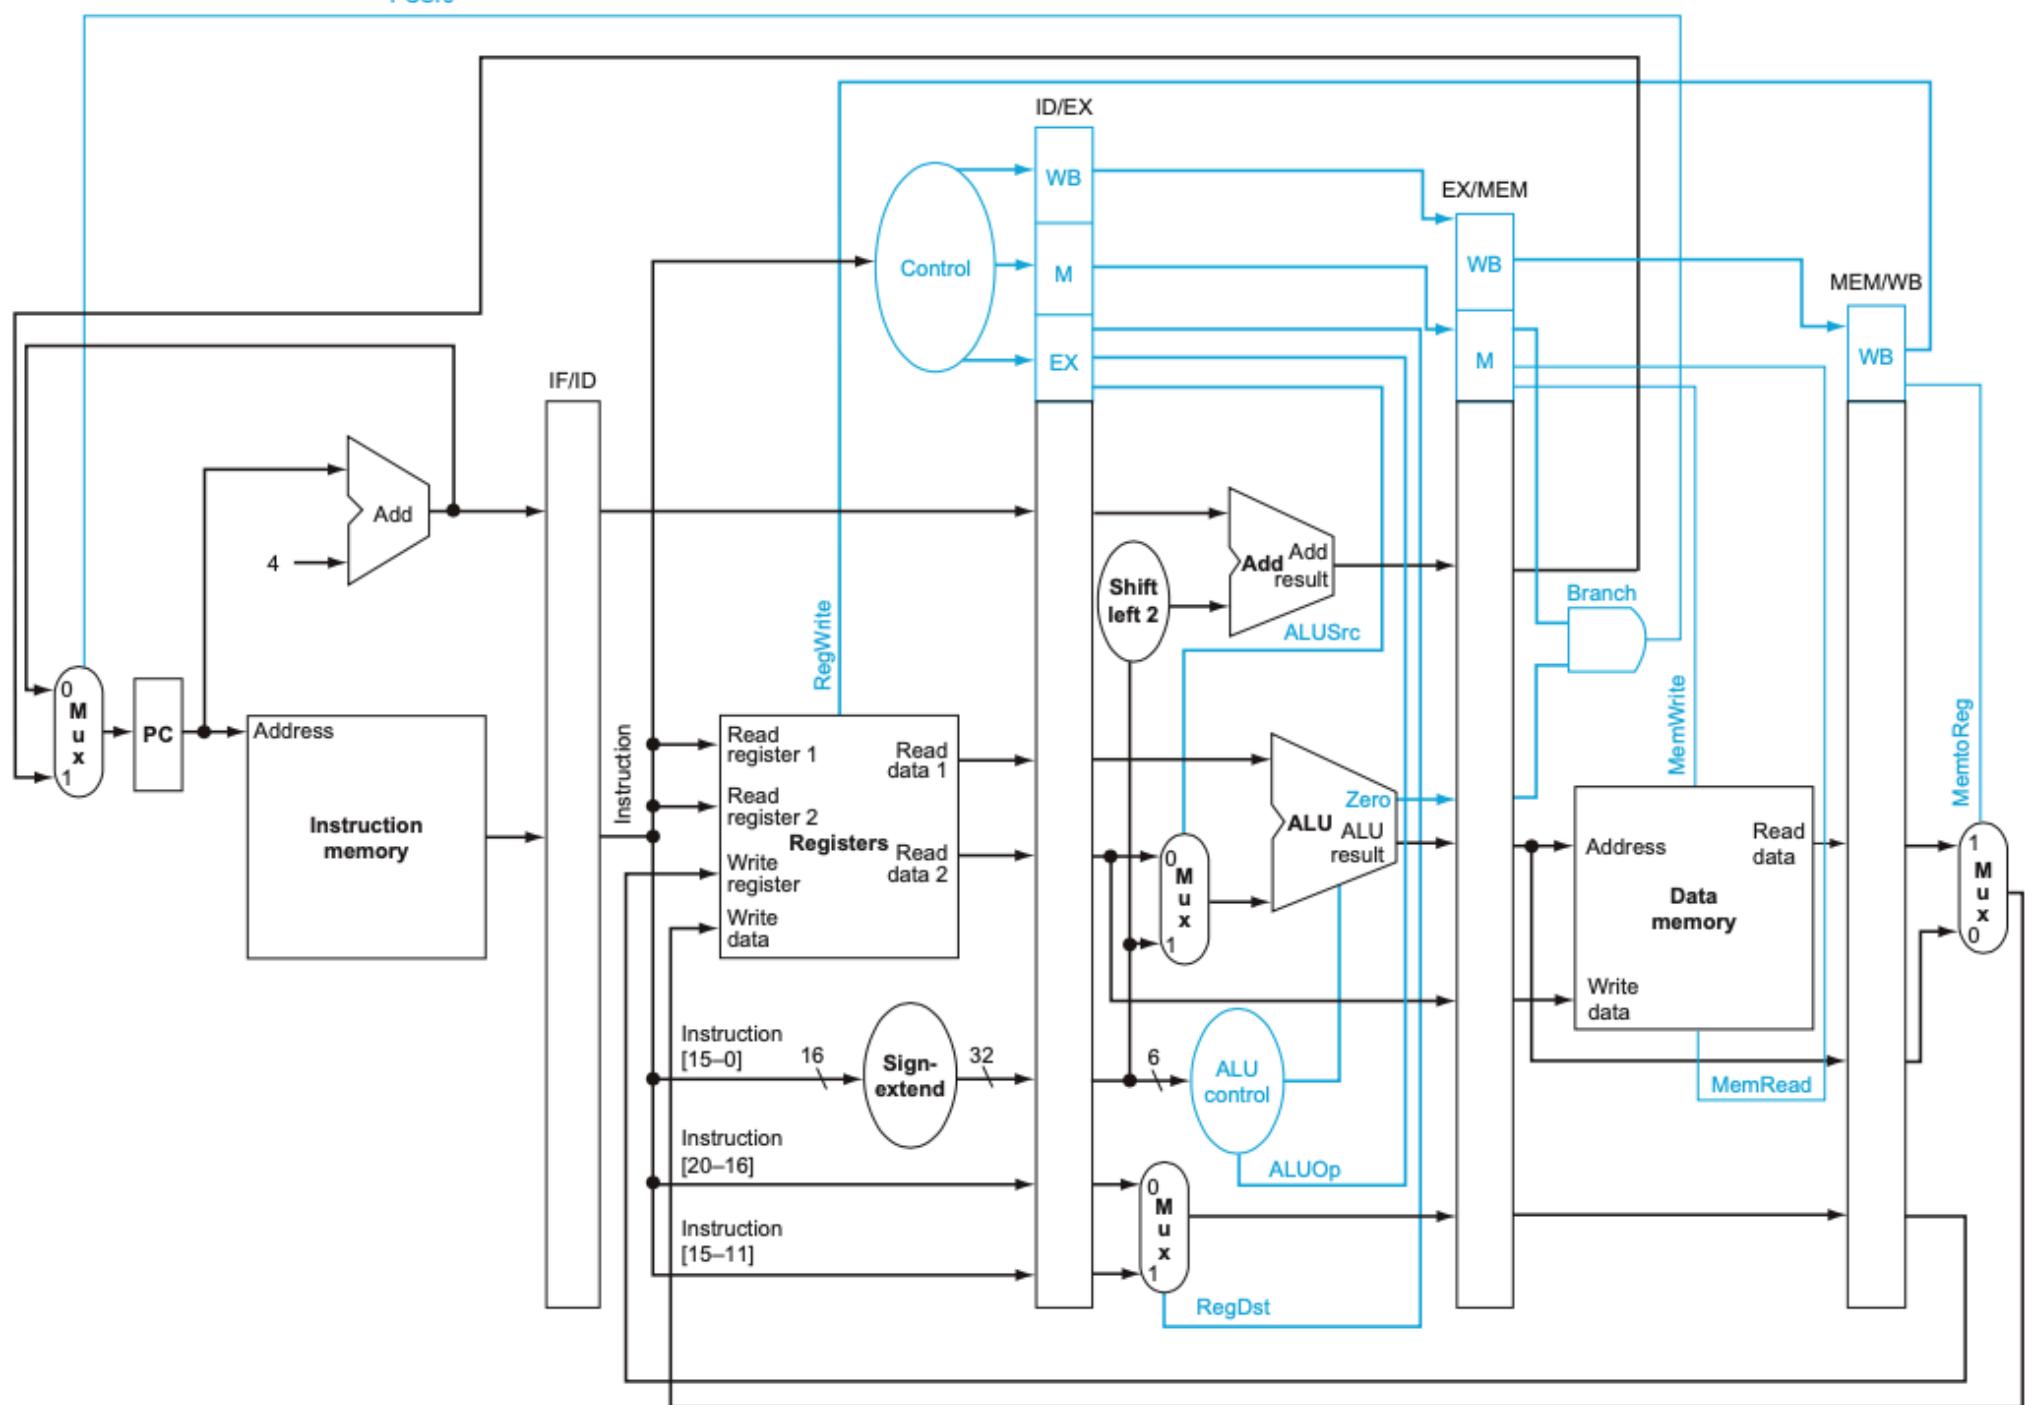


Control signals only needed in EX, MEM, WB stages

All signals depend on instr.,
which we know in decode (ID).
Actually at IF, why not there??

So we compute all control and
string it through the pipeline registers!

Our pipeline registers got bigger,
But that's fine b/c it's not
defined in the ISA.



NYU

Hazards!

Hazard (formal) := “when the next instruction cannot execute
in the following clock cycle”

Hazard (me) := “you need something you don’t have”

There are three of them..

Think about what goes into a program and how they execute

- 1) Structural : hardware resources
- 2) Data : need result from prior computation
- 3) Control : need to know where to go next



NYU

TANDON SCHOOL
OF ENGINEERING

Hazards!

Hazard (formal) := “when the next instruction cannot execute in the following clock cycle”

Hazard (me) := “you need something you don’t have”

There are three of them..

Think about what goes into a program and how they execute

- 1) Structural : hardware resources
- 2) Data : need result from prior computation
- 3) Control : need to know where to go next



NYU

TANDON SCHOOL
OF ENGINEERING

Hazards!

Hazard (formal) := “when the next instruction cannot execute in the following clock cycle”

Hazard (me) := “you need something you don’t have”

There are three of them..

Think about what goes into a program and how they execute

- 1) Structural : hardware resources : SOLVED!
- 2) Data : need result from prior computation
- 3) Control : need to know where to go next => Branch prediction (later)



NYU

TANDON SCHOOL
OF ENGINEERING

Program with dependencies

Sub \$r2, \$r1, \$r3

And \$r12, \$r2, \$r5

Or \$r13, \$r6, \$r2

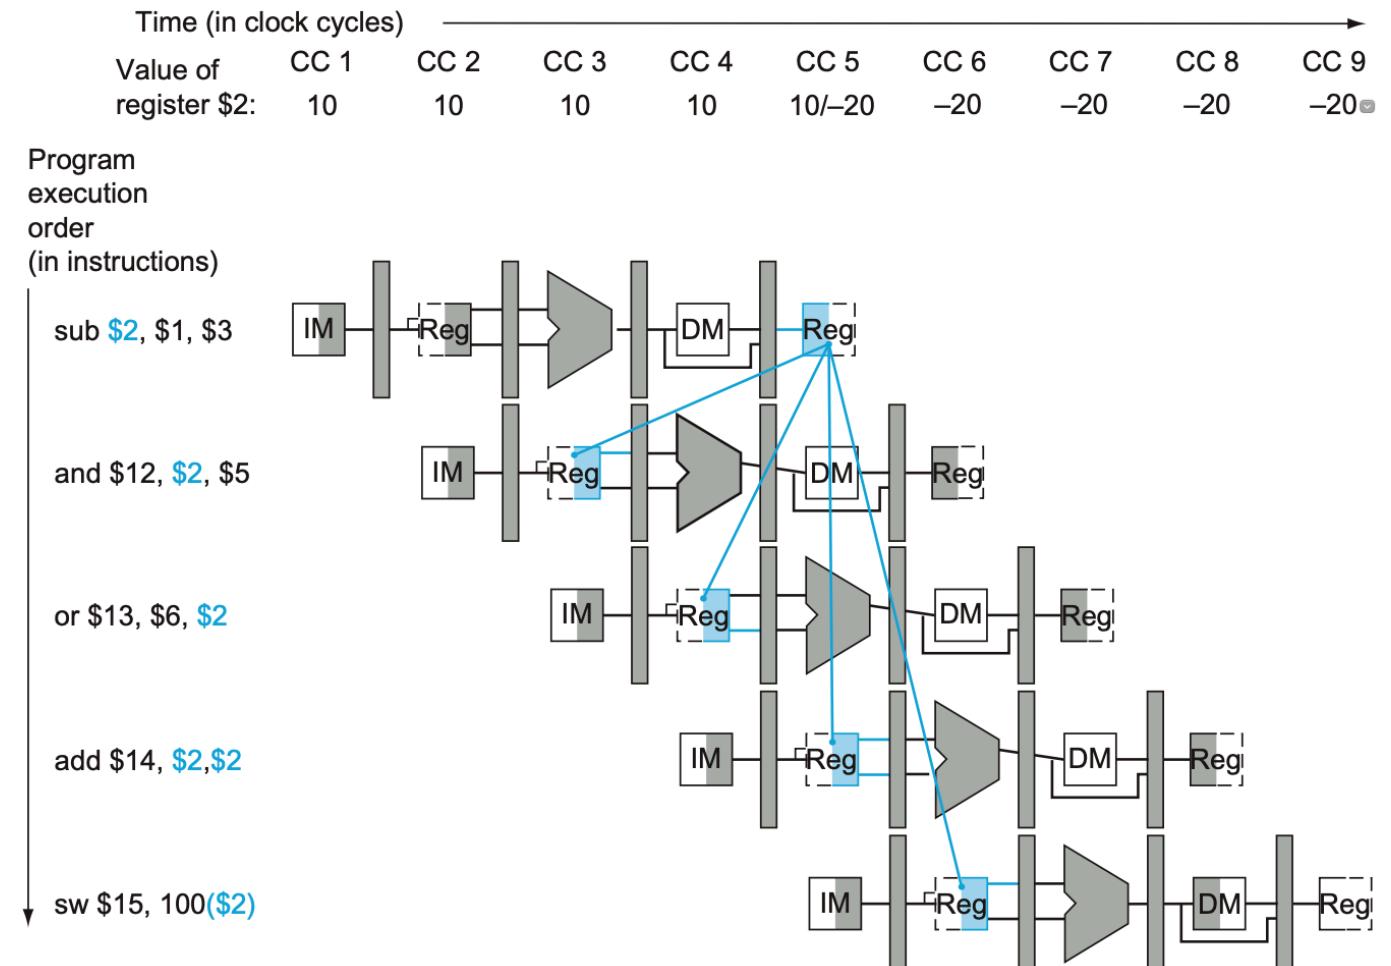
Add \$r14, \$r2, \$r2

Sw \$r15, 100(\$r2)

Take a minute to
Highlight all RAW
Dependencies.

Program with dependencies

Sub \$r2, \$r1, \$r3
And \$r12, \$r2, \$r5
Or \$r13, \$r6, \$r2
Add \$r14, \$r2, \$r2
Sw \$r15, 100(\$r2)



Notation for describing detecting hazards

Notation: <StagePrev> / <StageNext> . <PipelineReg>

1a. EX/MEM.RegisterRd = ID/EX.RegisterRs

1b. EX/MEM.RegisterRd = ID/EX.RegisterRt

2a. MEM/WB.RegisterRd = ID/EX.RegisterRs

2b. MEM/WB.RegisterRd = ID/EX.RegisterRt

Ex-Ex hazards

Mem-Ex hazards



NYU

TANDON SCHOOL
OF ENGINEERING

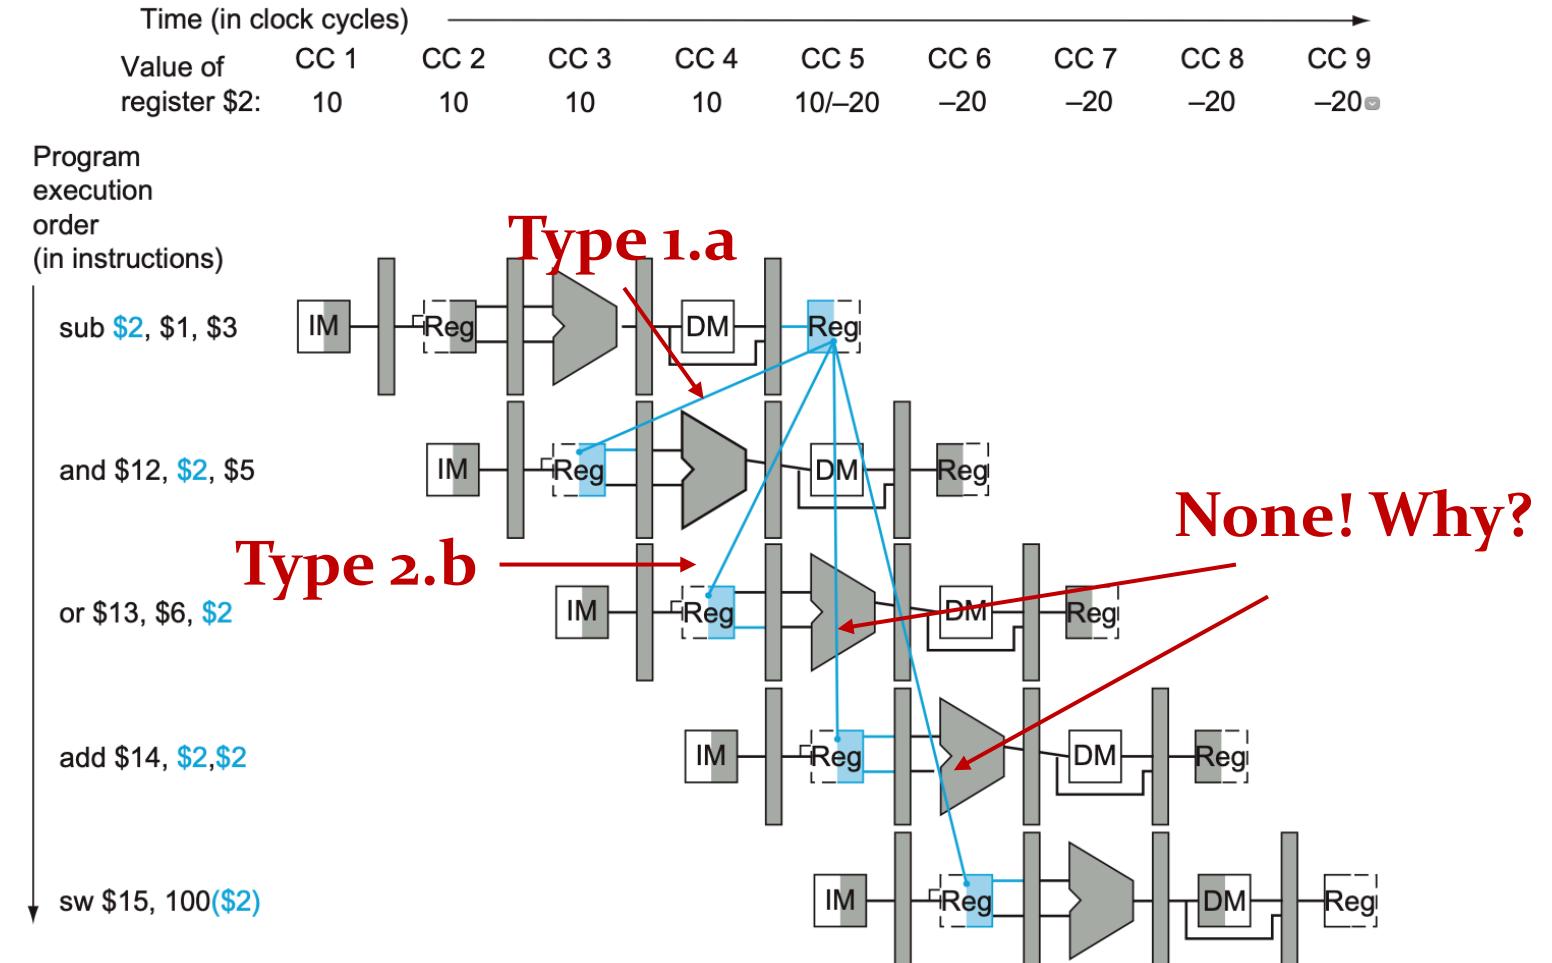
Program with dependencies

1a. EX/MEM.RegisterRd =
ID/EX.RegisterRs

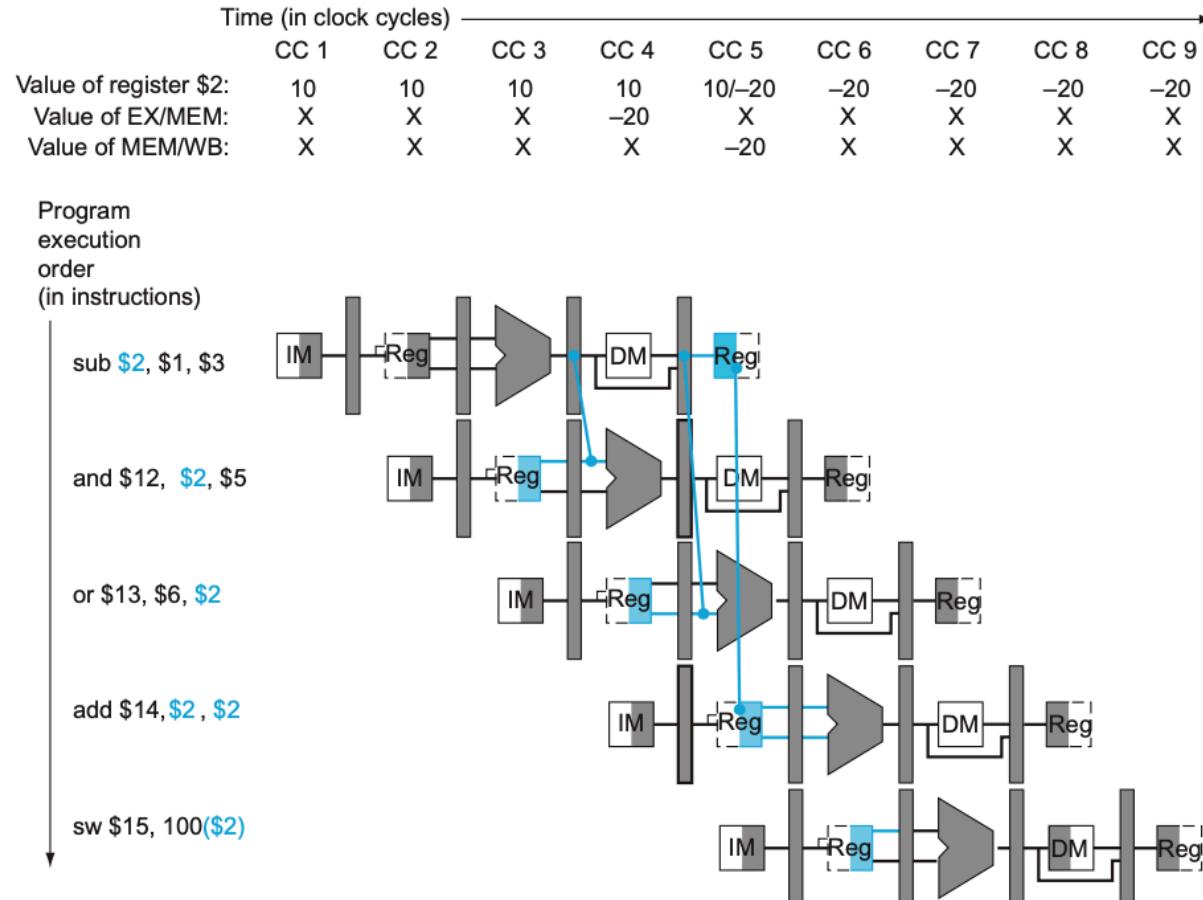
1b. EX/MEM.RegisterRd =
ID/EX.RegisterRt

2a. MEM/WB.RegisterRd =
ID/EX.RegisterRs

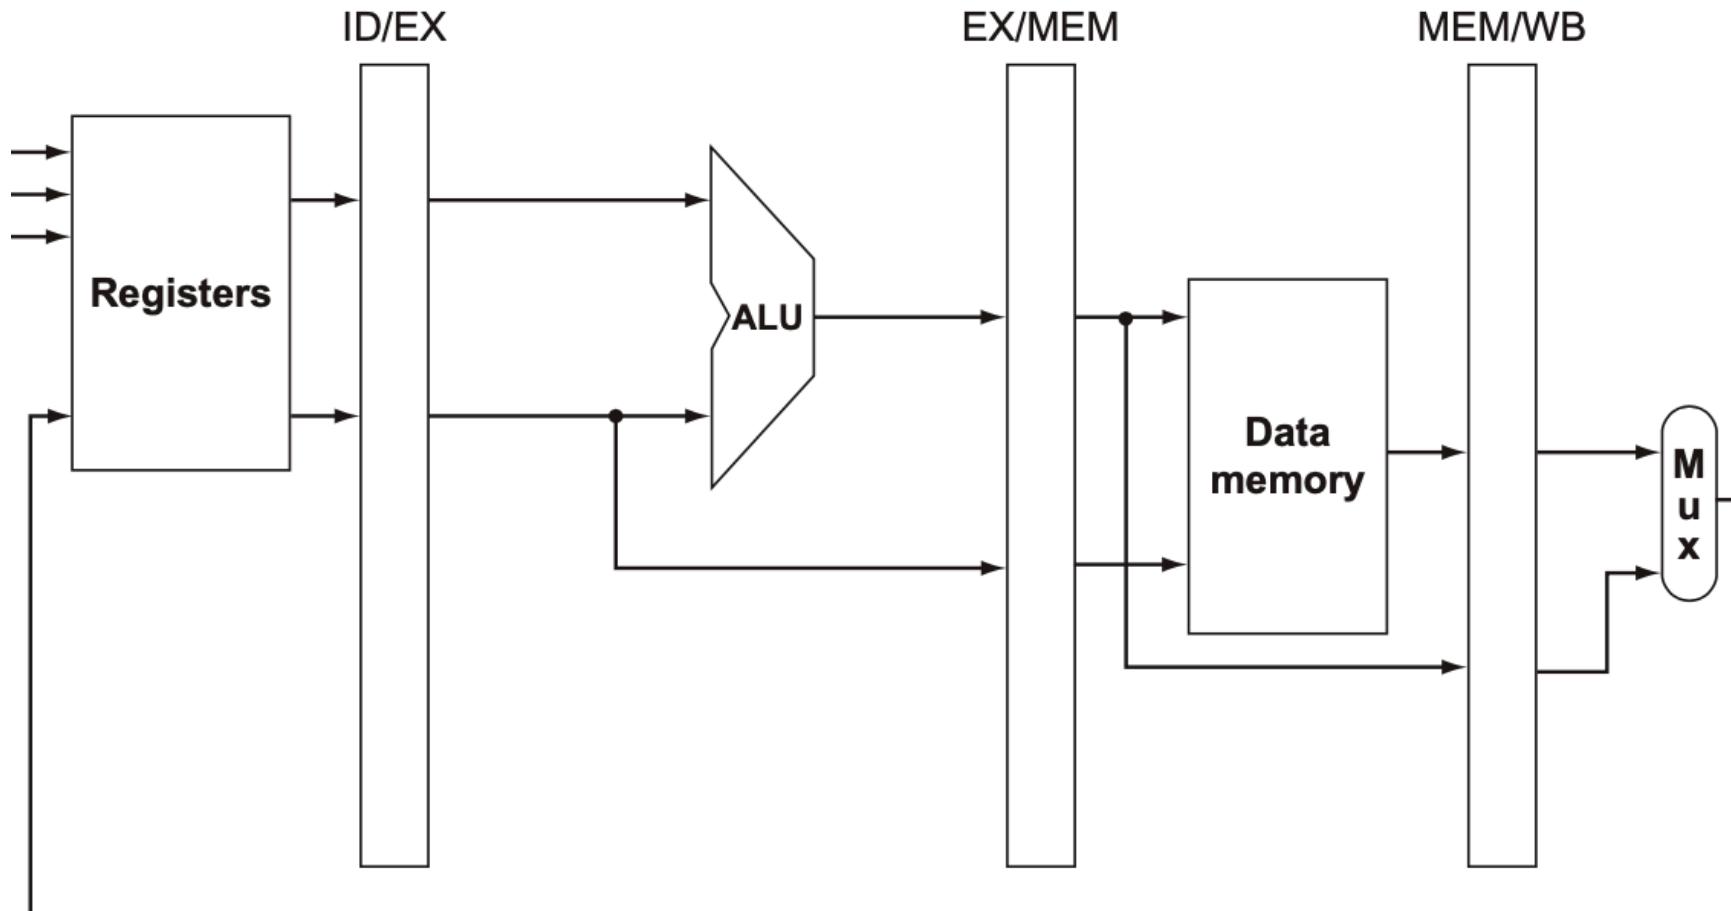
2b. MEM/WB.RegisterRd =
ID/EX.RegisterRt



Solve register hazards with forwarding



Implementing forwarding in the pipeline



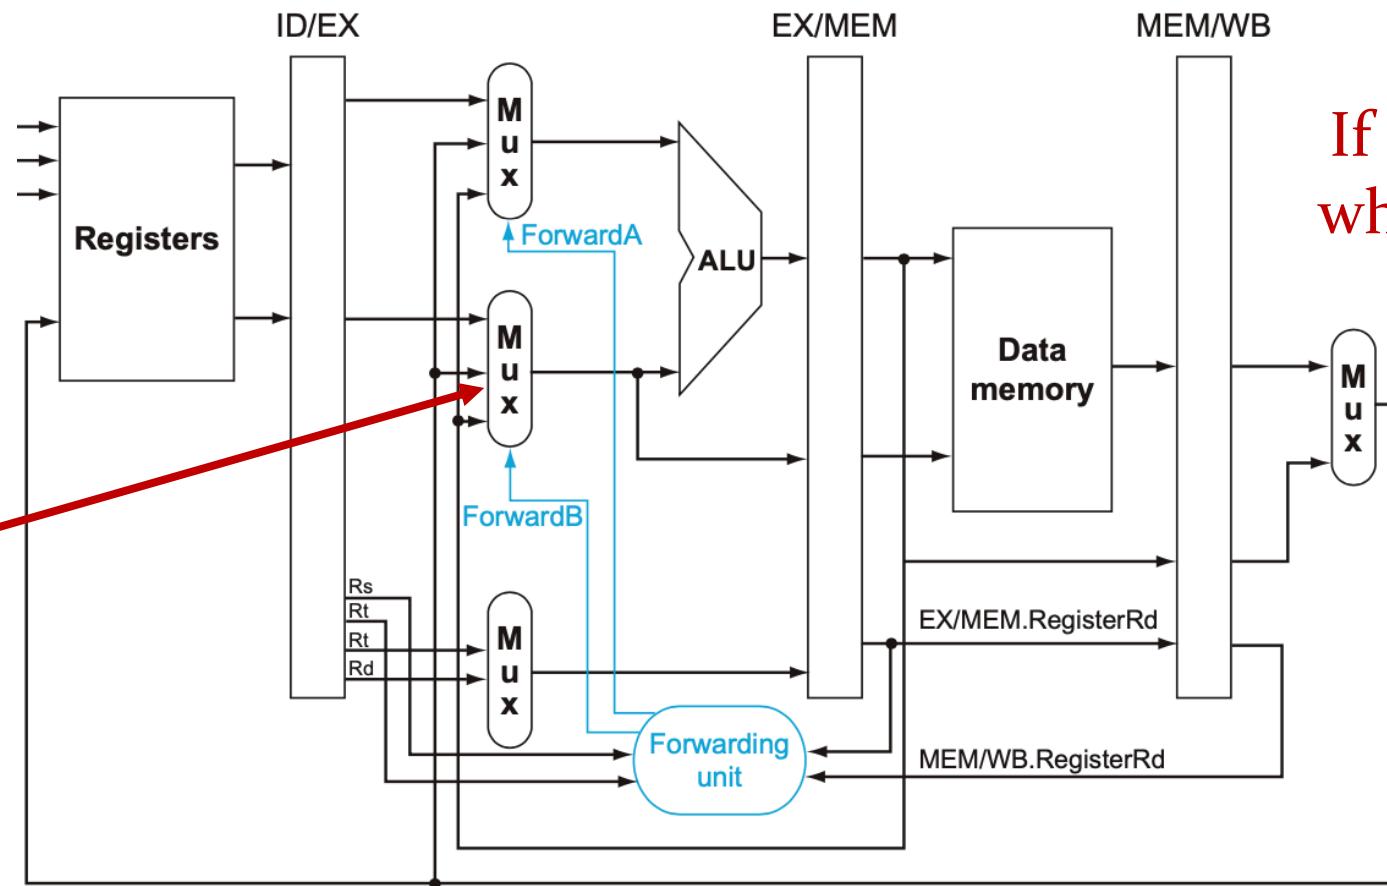
NYU

TANDON SCHOOL
OF ENGINEERING

Forward datapath logic added

3 sources:

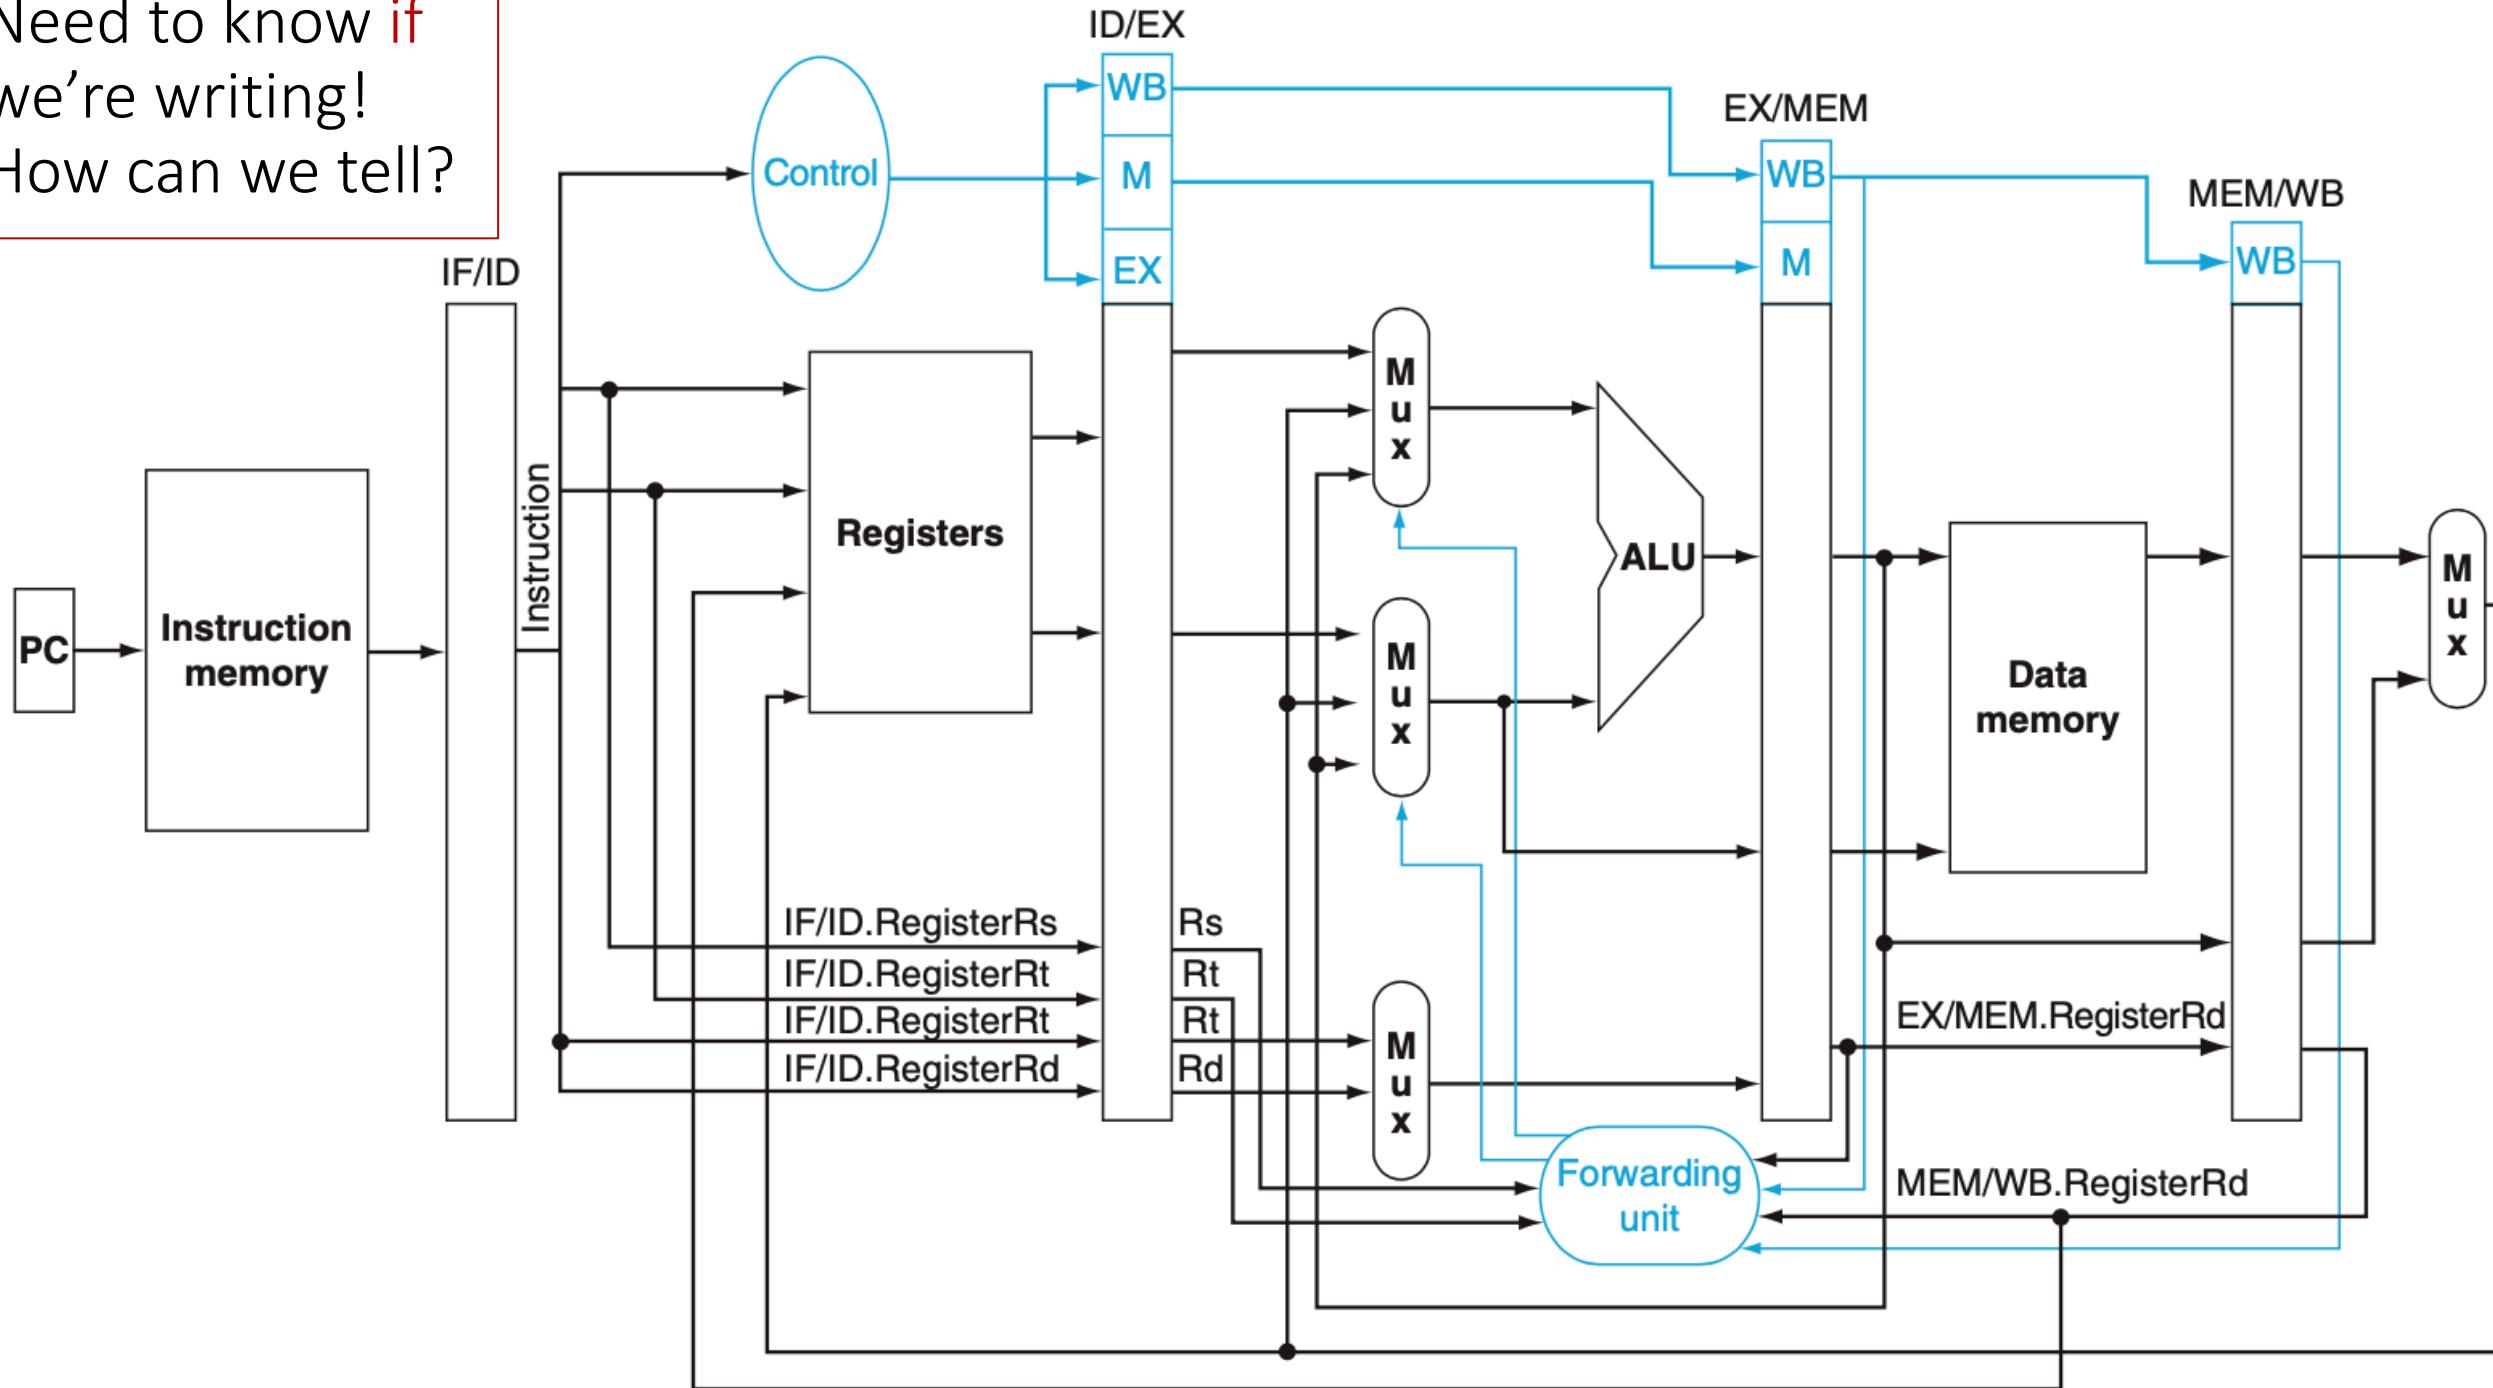
- ID/EX (from RF)
- EX/MEM (from ALU)
- MEM/WB (from Mem)



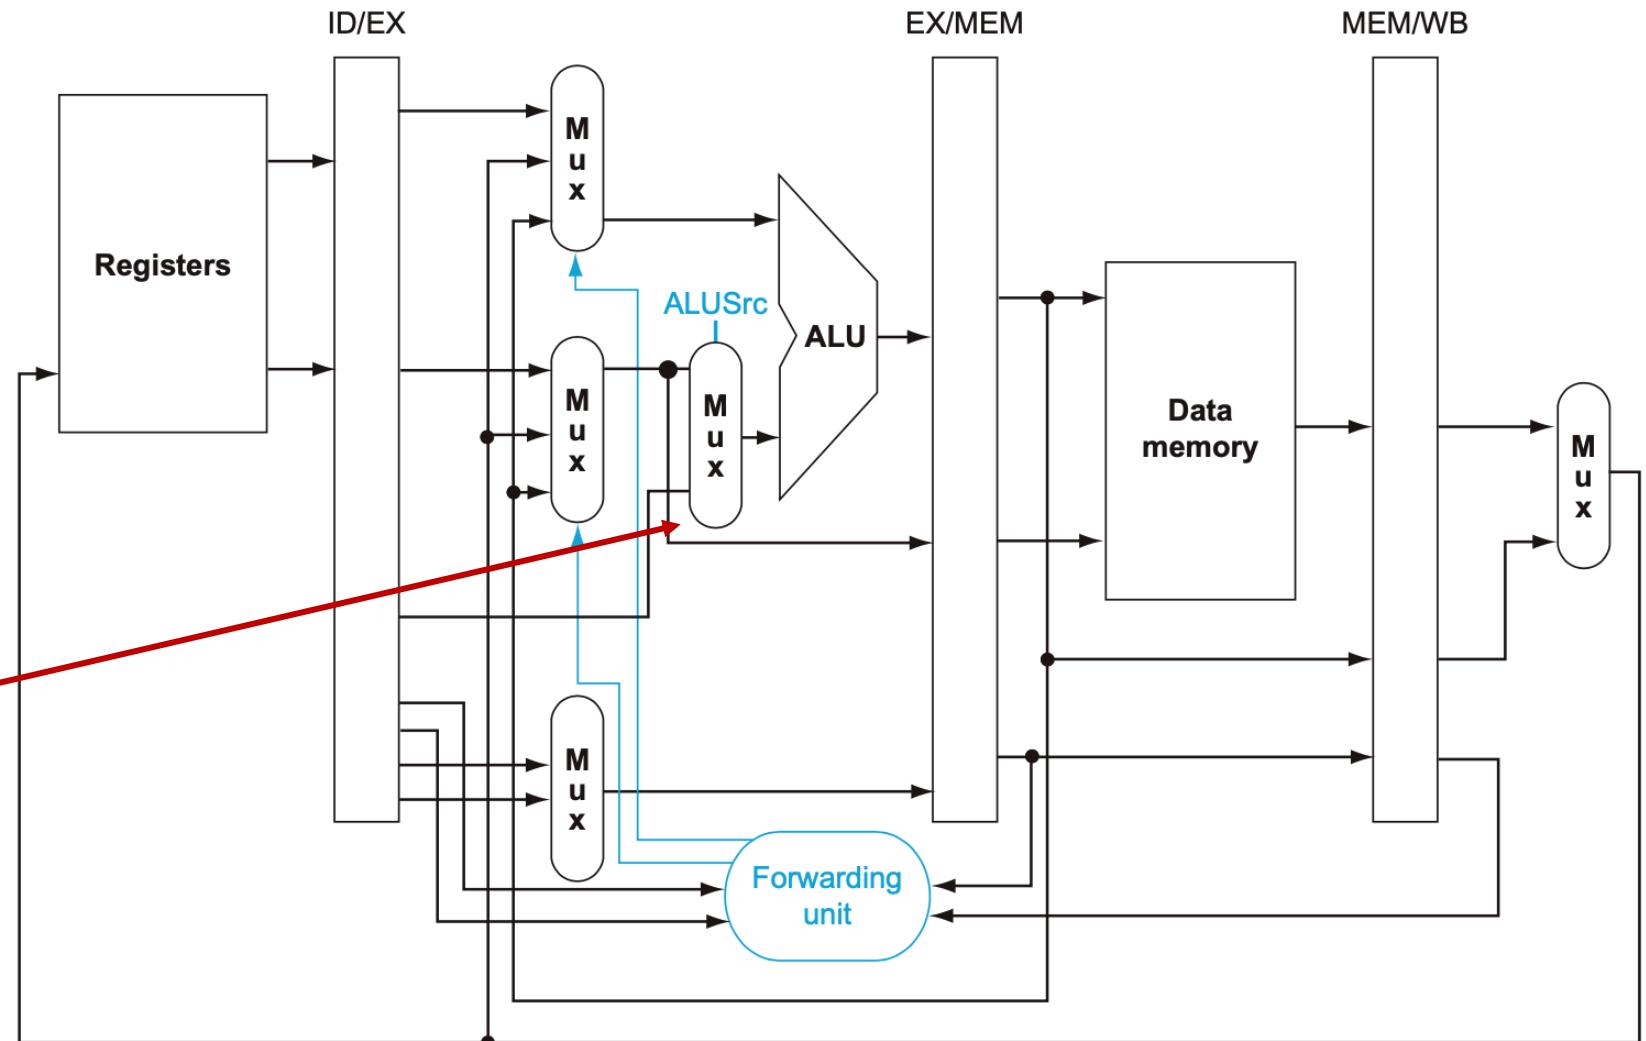
If both cases are true
which should we use?

Q: What are we
missing here?
(Take a min)

Need to know if
we're writing!
How can we tell?



For completeness..



Add another MUX to choose
between Immediate and register



NYU

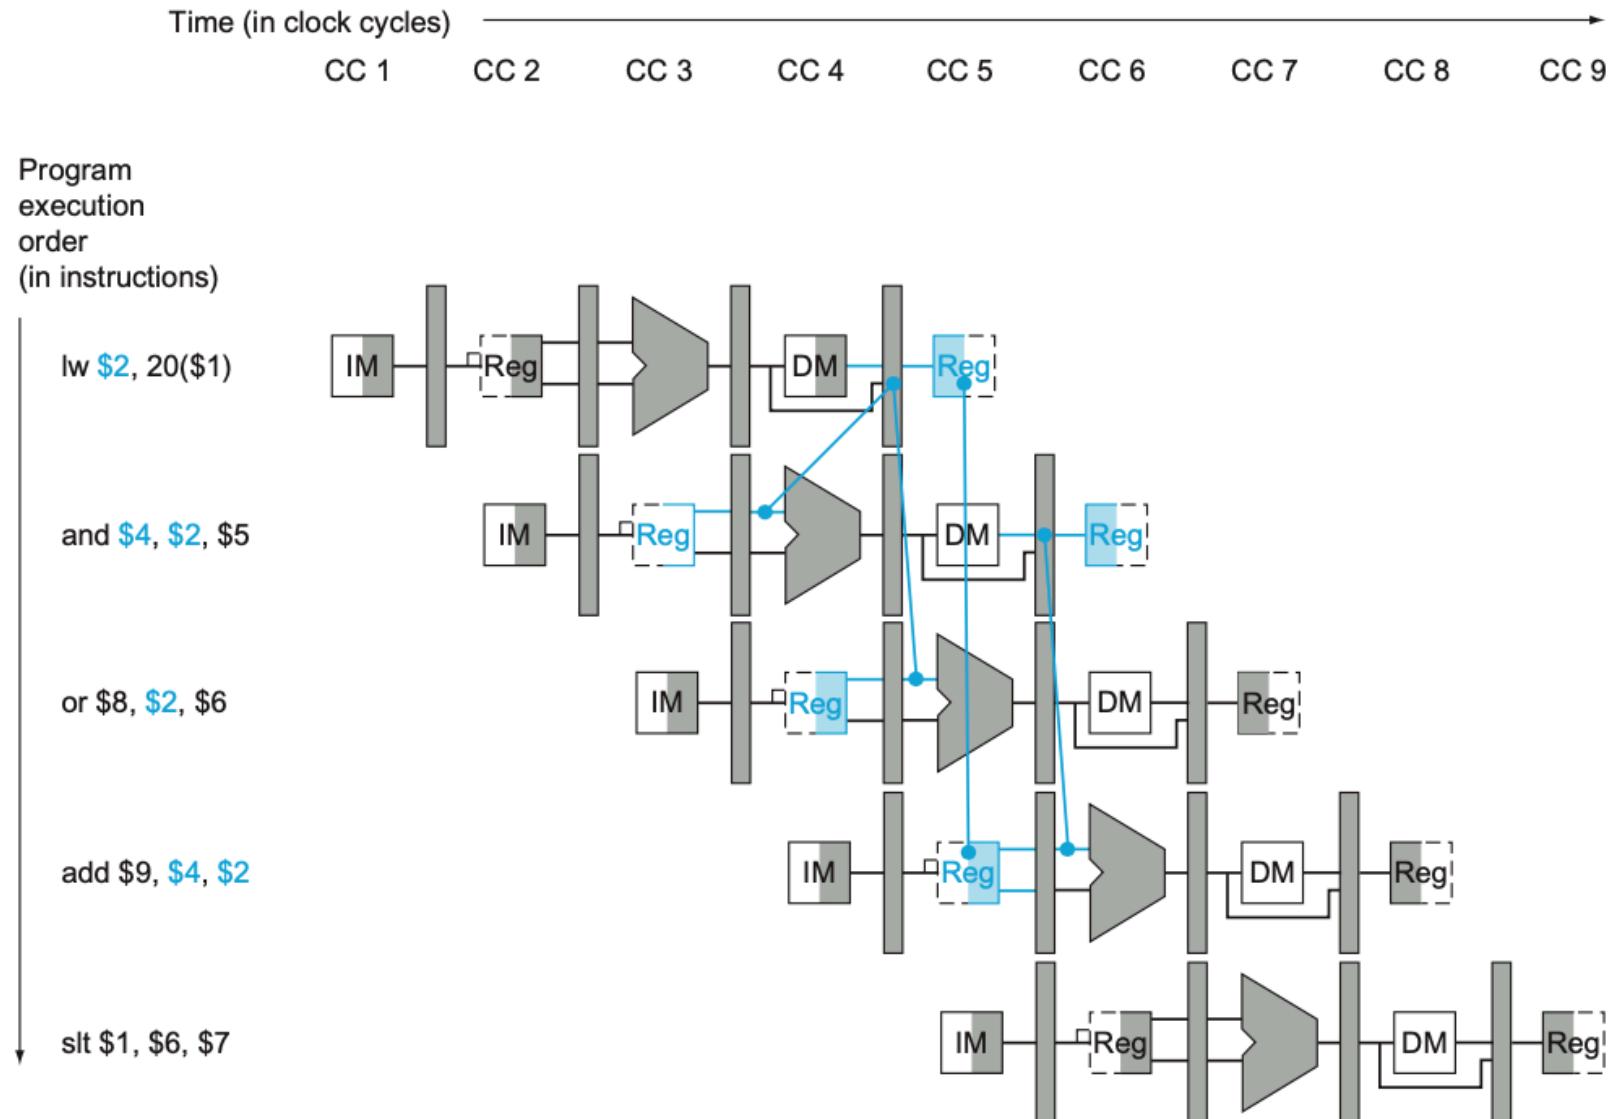
TANDON SCHOOL
OF ENGINEERING

When does forwarding fail?

Can't go backward in time.

With a LW instruction,
we don't have the value
until end of MEM stage.

If we have a dependence
between lw and the next
instruction, we must stall.



When do we have a hazard?

What stage do we detect the hazard in?
Decode!

The previous instruction was a load

```
if (ID/EX.MemRead and  
((ID/EX.RegisterRt = IF/ID.RegisterRs) or  
(ID/EX.RegisterRt = IF/ID.RegisterRt)))  
stall the pipeline
```

The previous instruction Destination matches either Source registers.

Dealing with data (lw -> add) hazards

When that condition is met, we have to **stall** one cycle

First, we disable updating the PC and ID/RF pipeline registers
“Freeze” the front end to stop progressing until hazard dealt with
If we didn’t stall PC, what would happen?

Second, insert a NOP into the ID/EX pipeline register
Can zero everything out, or just the control signals

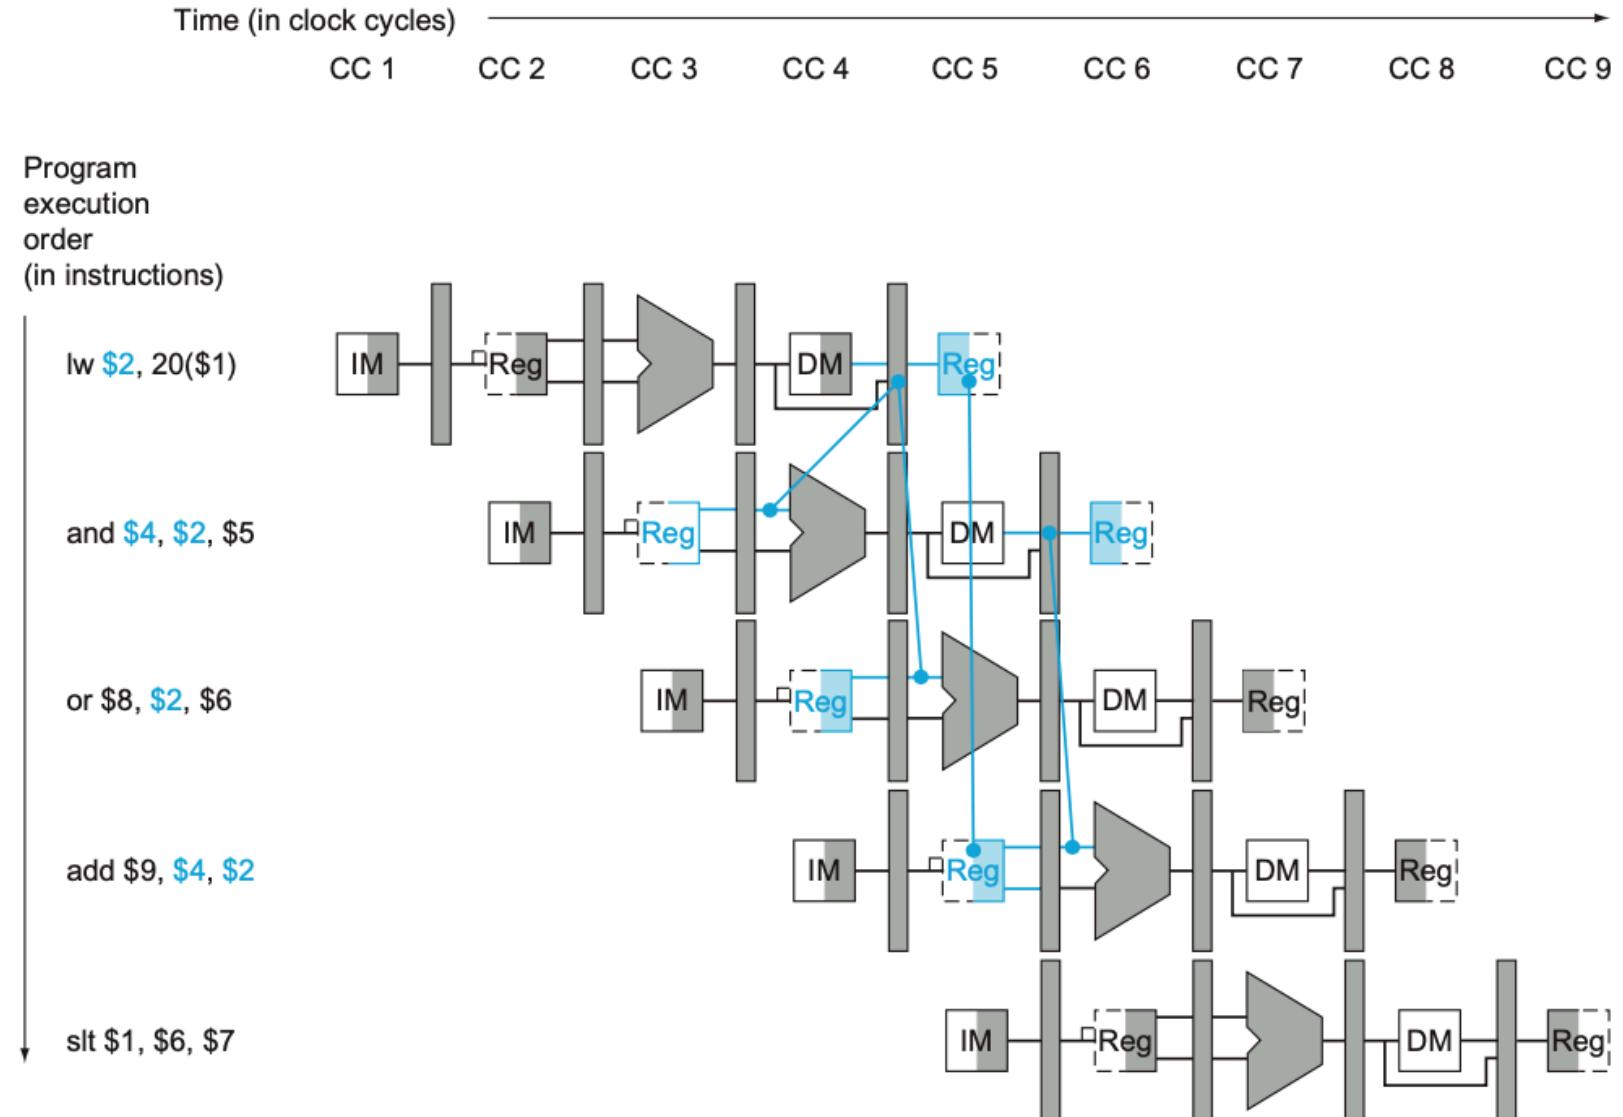
Why?

When does forwarding fail?

Can't go backward in time.

With a LW instruction,
we don't have the value
until end of MEM stage.

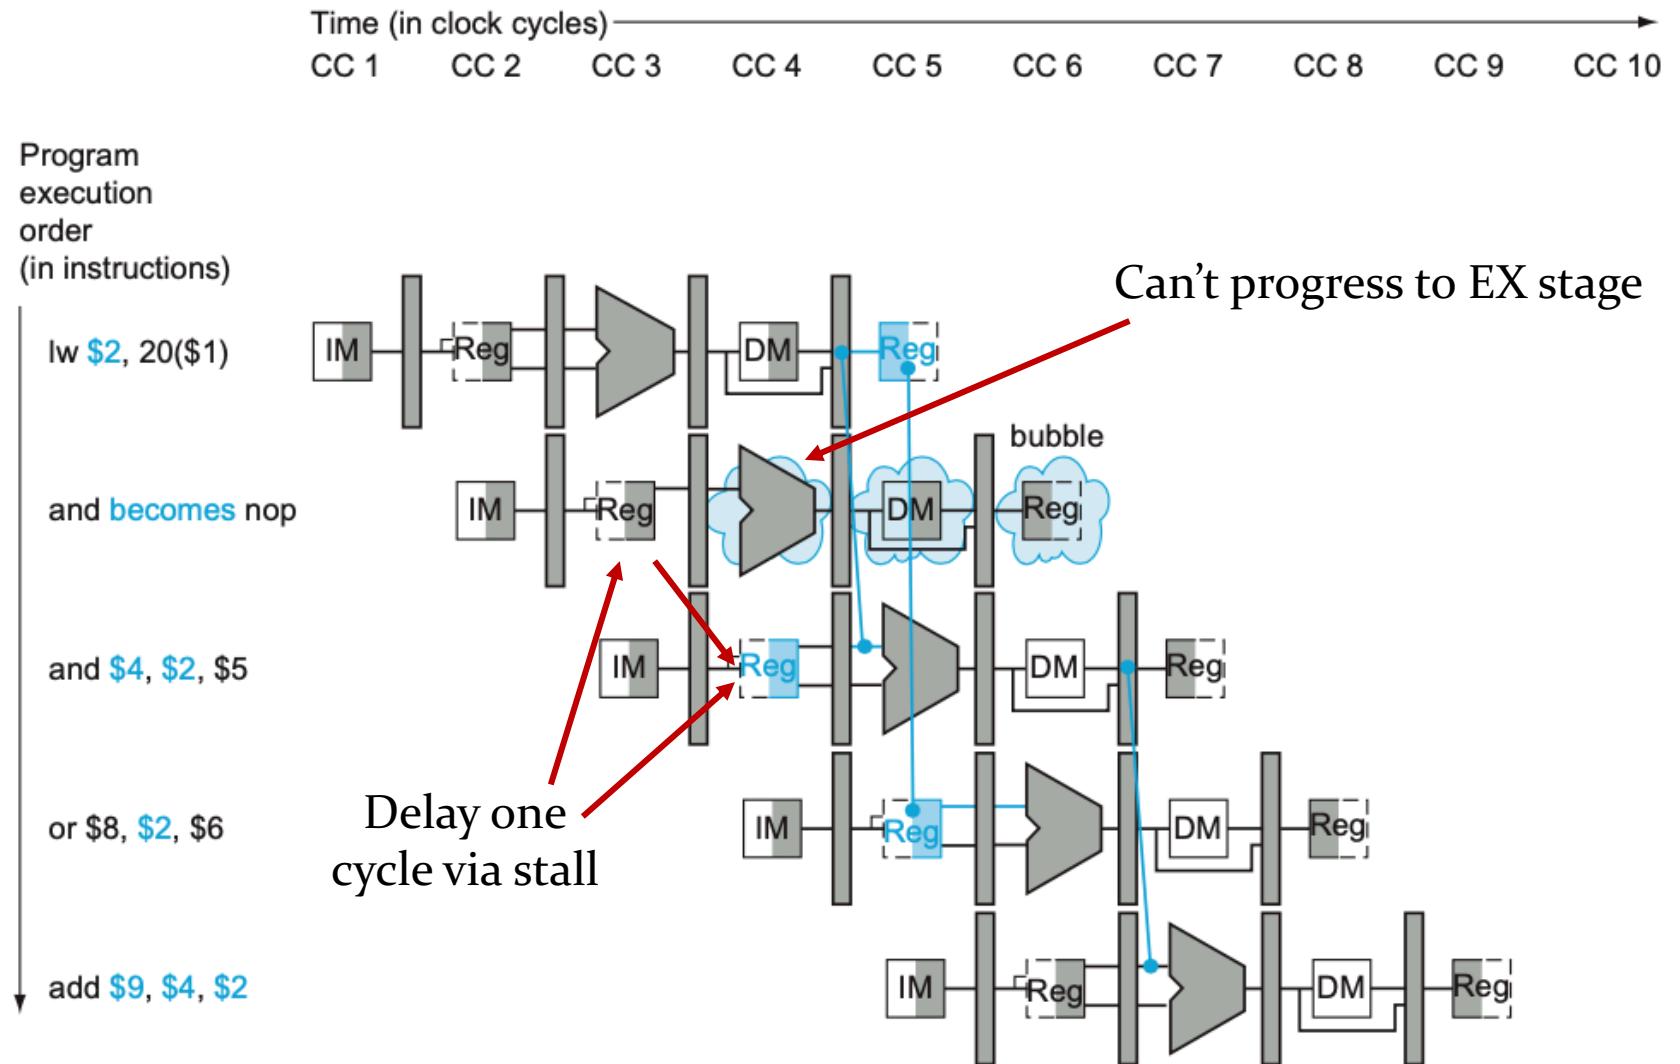
If we have a dependence
between lw and the next
instruction, we must stall.



NYU

TANDON SCHOOL
OF ENGINEERING

Stall+NOP creates a bubble in the pipeline



NYU

TANDON SCHOOL
OF ENGINEERING

QUIZ: Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1					
2					
3					
4					
5					
6					
7					
8					
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1					
2					
3					
4					
5					
6					
7					
8					
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2					
3					
4					
5					
6					
7					
8					
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3					
4					
5					
6					
7					
8					
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)		
4					
5					
6					
7					
8					
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)		
4	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	Lw \$r1, o(\$r2)	
5					
6					
7					
8					
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)		
4	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	Lw \$r1, o(\$r2)	
5	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!	Lw \$r1, o(\$r2)
6					
7					
8					
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)		
4	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	Lw \$r1, o(\$r2)	
5	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!	Lw \$r1, o(\$r2)
6	Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!
7					
8					
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)		
4	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	Lw \$r1, o(\$r2)	
5	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!	Lw \$r1, o(\$r2)
6	Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!
7		Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!
8					
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)		
4	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	Lw \$r1, o(\$r2)	
5	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!	Lw \$r1, o(\$r2)
6	Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!
7		Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!
8			Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3
9					
10					



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)		
4	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	Lw \$r1, o(\$r2)	
5	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!	Lw \$r1, o(\$r2)
6	Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!
7		Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!
8			Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3
9				Mul \$r3, \$r3	Add \$r3, \$r1, \$r2
10					



Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)		
4	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	Lw \$r1, o(\$r2)	
5	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!	Lw \$r1, o(\$r2)
6	Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	BUBBLE!
7		Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!
8			Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3
9				Mul \$r3, \$r3	Add \$r3, \$r1, \$r2
10					Mul \$r3, \$r3

Instructions: 4
Cycles: 10



NYU

TANDON SCHOOL
OF ENGINEERING

Example no MEM fwd

Program

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All dependencies

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

All hazards

Lw	\$r1, o(\$r2)
Add	\$r2, \$r1, \$r3
Add	\$r3, \$r1, \$r2
Mul	\$r3, \$r3

Cycle	IF	ID	EX	MEM	WB
1	Lw \$r1, o(\$r2)				
2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)			
3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	Lw \$r1, o(\$r2)		
4	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	Lw \$r1, o(\$r2)	
5	Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!	Lw \$r1, o(\$r2)
6		Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3	BUBBLE!
7			Mul \$r3, \$r3	Add \$r3, \$r1, \$r2	Add \$r2, \$r1, \$r3
8				Mul \$r3, \$r3	Add \$r3, \$r1, \$r2
9					Mul \$r3, \$r3

Instructions: 4
Cycles: 9

Performance metrics!

As we'll see now, comparing computer performance isn't as trivial as it sounds..

Revisit some definitions

Execution/Response time (**Latency**)

- Elapsed time between the start and completion of an event
- How long did your project take?
- Where have we seen this in class?

Throughput (**Bandwidth**)

- Total amount of work done within some amount of time
- How many lines of code (or bugs) did you write per hour?

Some new definitions

Assume we're running a program on a machine X

Performance:

$$\text{Performance}(X) = 1 / \text{Execution time } (x)$$

It's the inverse of the runtime.

“X is n times faster than Y”

$$\text{Performance}(X) / \text{Performance}(Y) \Rightarrow \text{speedup factor of } n$$

We evaluate a computer's performance using benchmarks

- (Real) Programs
 - In the form of collection of programs
 - E.g., SPEC, Winstone, SYSMARK, 3D Winbench, EEMBC
- Kernels
 - Small key pieces of real programs
 - E.g., Livermore Fortran Loops Kernels (LFK), Linpack
- Modified (or scripted)
 - To focus on particular aspects (e.g. remove I/O, focus on CPU)
- “Toy” Benchmarks
 - Produce expected results
 - Why are these so important?
- Synthetic Benchmarks:
 - Representative instruction mix
 - E.g., Dhrystone, Whetstone

Important for:
Understanding architectural and
microarchitectural design trade-off

Competitive analysis of
real products

Check out MLPerf!

On averages..

Arithmetic mean: Use when using the same units

$$\frac{1}{n} \sum_{i=1}^n Time_i \quad \text{or} \quad \frac{1}{n} \sum_{i=1}^n Weight_i * Time_i$$

Harmonic mean: Use when using data values are ratios of variables with different rates

$$\frac{n}{\sum_{i=1}^n \frac{1}{Rate_i}} \quad \text{or} \quad \frac{n}{\sum_{i=1}^n \frac{Weight_i}{Rate_i}}$$



NYU

TANDON SCHOOL
OF ENGINEERING

Why Harmonic Mean? Driving example

30 mph for the first 10 miles

90 mph for the next 10 miles

Average speed? $(30+90)/2 = 60 \text{ mph}??$

Wrong!

Average speed = total distance / total time

$(10+10)/(10/30 + 10/90) = 45 \text{ mph}$

Make the common case fast

Drive NYC -> Boston

- 130 mph NYC -> Stamford (38 mi)
- 65 mph Stamford -> Boston (176 mi)
- First takes $38/130 = 17.5$ minutes
- Second takes $176/65 = 162.5$ minutes
- Total = 180 min.

Just go 65 the whole way: 197 minutes

$$\text{Speedup} = 197 / 180 = 1.09x$$

Not great.. Probably not worth getting arrested.

Make the common case fast!

3 h 43 min (214 miles)

via I-95 N

3 h 43 min without traffic

⚠ This route has tolls.

Washington Square Park

New York, NY 10012

- Take 3rd Ave, E 23rd St and Fdr Drive Service Rd E to FDR Dr
10 min (2.0 mi)
- Follow FDR Dr, I-278 E and I-95 N to Greyrock Pl in Stamford. Take exit 8 from I-95 N
42 min (36.0 mi)
- Take Tresser Blvd to Atlantic St
3 min (0.5 mi)



Stamford

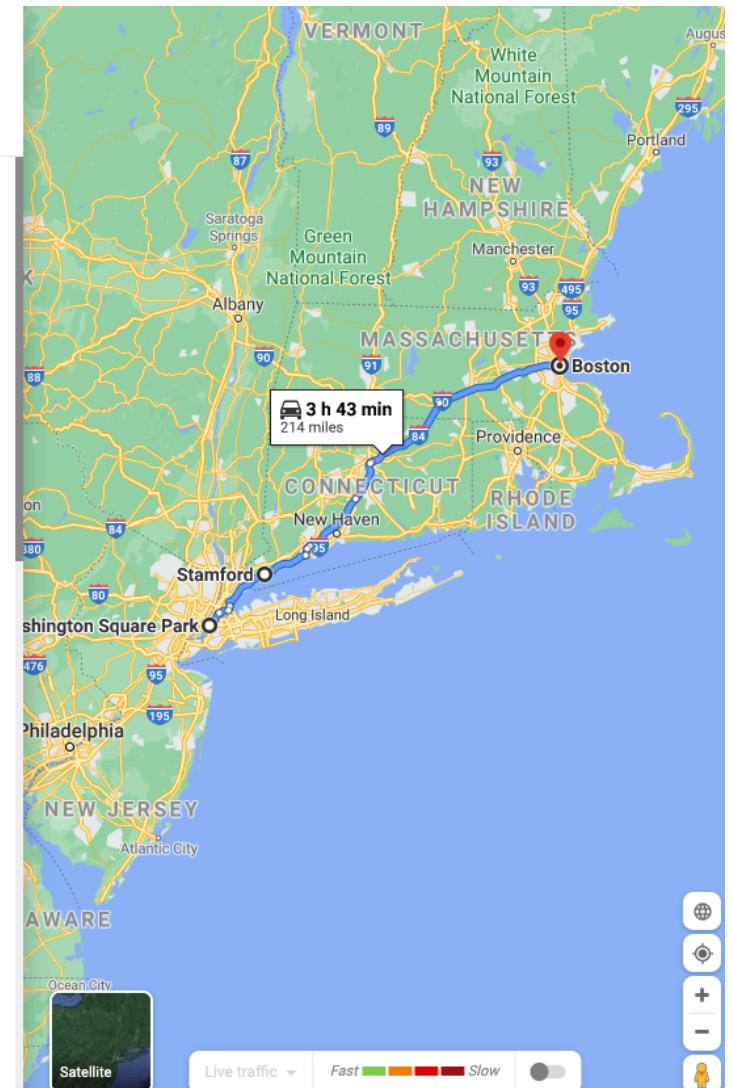
Connecticut

- Get on I-95 N from Broad St and Greyrock Pl
4 min (0.9 mi)
- Continue on I-95 N. Take CT-15 N, I-91 N, I-84 E and I-90 E to North St in Boston. Take exit 23 from I-93 N
2 h 40 min (175 mi)
- Continue on North St. Take Congress St and Court St to Cambridge St
3 min (0.5 mi)



Boston

Massachusetts



NYU

TANDON SCHOOL
OF ENGINEERING

Amdahl's Law (Law of Diminishing Returns)

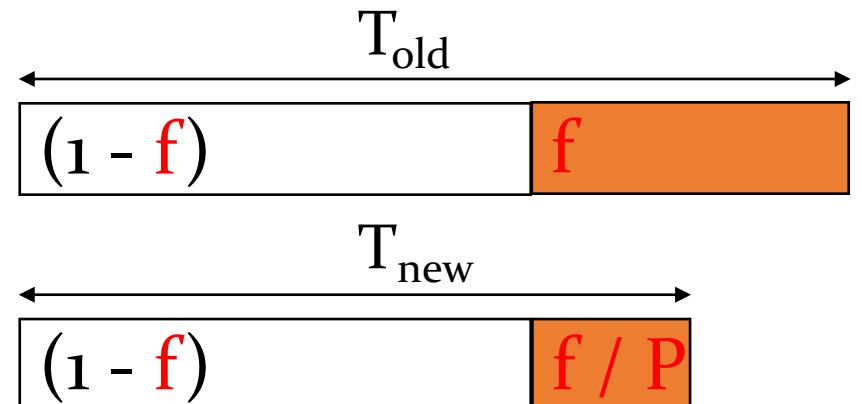
Performance improvement is limited by the fraction of time it can be applied

$$\text{Speedup} = \text{Perf}_{\text{new}} / \text{Perf}_{\text{old}} = T_{\text{old}} / T_{\text{new}} = \frac{1}{(1-f) + \frac{f}{P}}$$

f = fraction of code sped up

P = speedup factor

Q: If one function takes 90% of the time,
what's max speedup? (Hint: $P \rightarrow \infty$)



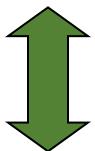
CPU Performance: Iron law of performance

Execution Time = Seconds / Program

$$\frac{\text{Instructions}}{\text{program}} \times \frac{\text{cycles}}{\text{Instruction}} \times \frac{\text{seconds}}{\text{cycle}}$$



- Programmer
- Algorithms
- ISA
- Compilers



- Microarchitecture
- System architecture



- Microarchitecture, pipeline depth
- Circuit design
- Technology



NYU

TANDON SCHOOL
OF ENGINEERING

Example of Performance Evaluation (I)

Operation	Frequency	Clock cycle count
ALU Ops (reg-reg)	43%	1
Loads	21%	2
Stores	12%	2
Branches	24%	2

Assume 25% of the ALU ops directly use a loaded operand that is not used again.
We propose adding ALU instructions that have one src operand in memory.
These new reg-mem instructions take 2 clock cycles.

Also assume that the extended instruction set increases the
branch inst. clock cycle count by 1 but does not impact to cycle time.

Would this change improve performance ?

Example of Performance Evaluation (I)

Operation	Frequency	Clock cycle count
ALU Ops (reg-reg)	43%	1
Loads	21%	2
Stores	12%	2
Branches	24%	2

Assume 25% of the ALU ops directly use a loaded operand that is not used again.
We propose adding ALU instructions that have one src operand in memory.
These new reg-mem instructions take 2 clock cycles.

Also assume that the extended instruction set increases the
branch inst. clock cycle count by 1 but does not impact to cycle time.

Would this change improve performance ?

New inst.

$$Cycles_{old} = 0.43 * 1 + 0.21 * 2 + 0.12 * 2 + 0.24 * 2 = 1.57$$

$$Cycles_{new} = 0.25 * 0.43 * 2 + (0.43 - 0.25 * 0.43) * 1 + (0.21 - 0.25 * 0.43) * 2 + 0.12 * 2 + 0.24 * 3 = 1.703$$

reg-reg
alu

Save load time!

Store and branch stays same.



Example of Performance Evaluation (II)

FP instructions = 25%

Average CPI of FP instructions = 4.0

Average CPI of other instructions = 1.33

FPSQRT = 2% of all instructions, CPI of FPSQRT = 20

Design Option 1: decrease the CPI of FQSQRT to 2

Design Option 2: decease the average CPI of all FP instructions to 2.5

Calculate: original CPI, Des opt 1 CPI, Des. Opt 2 CPI.

$$\text{Original CPI} = 0.25 * 4 + 1.33 * (1 - 0.25) = 2.0$$

$$\text{Option 1 CPI} = 2.0 - 2\% * (20 - 2) = 1.64$$

$$\text{Option 2 CPI} = 0.25 * 2.5 + 1.33 * (1 - 0.25) = 1.625$$

$$\text{Speedup of Option 1} = 2 / 1.64 = 1.2195$$

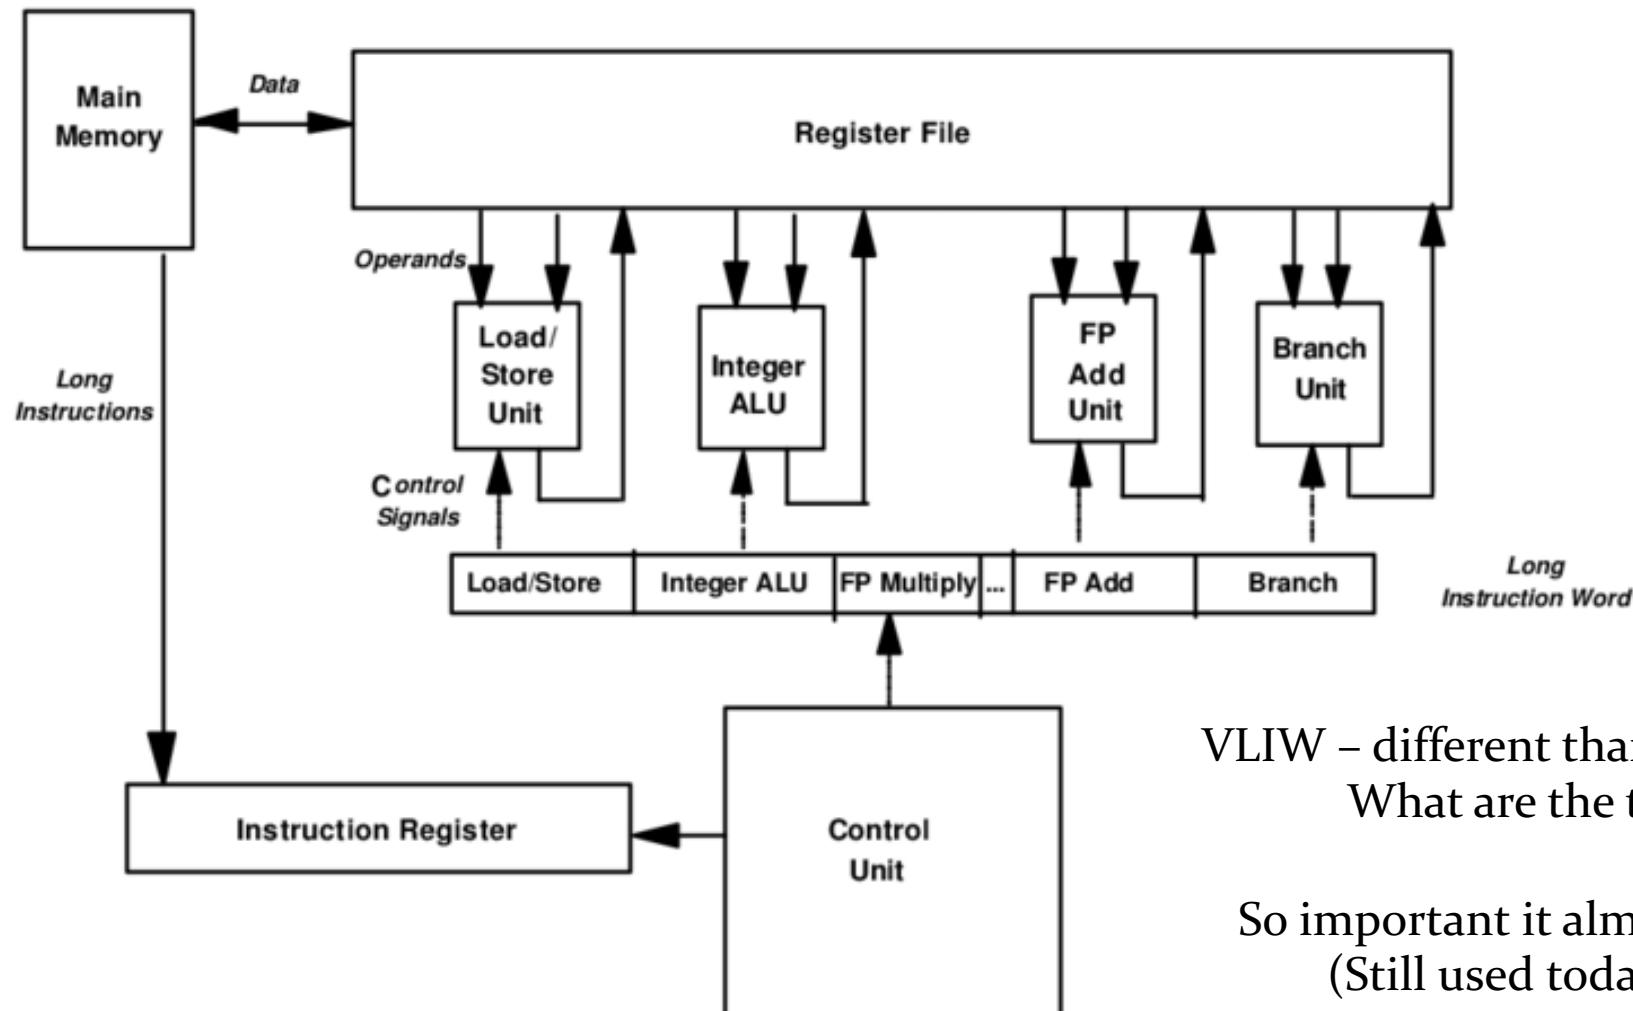
$$\text{Speedup of Option 2} = 2 / 1.625 = 1.2308$$



NYU

TANDON SCHOOL
OF ENGINEERING

Does NYU even do architecture?



VLIW – different than RISC and CISC
What are the tradeoffs?

So important it almost killed Intel
(Still used today in DSPs)

Who invented VLIW?



NYU

TANDON SCHOOL
OF ENGINEERING

Josh Fisher!

Josh Fisher



Josh Fisher in 2014.

Highest award in field.
Given at ISCA each year.

~Second highest award in
field. Given at MICRO
each year.

Born July 22, 1946 (age 74) [\[1\]](#)
Bronx, NY, USA [\[1\]](#)

Alma mater Courant Institute of Mathematical Sciences ([New York University](#))

Known for The Invention of VLIW Architectures, Instruction-level Parallelism, Trace Scheduling, Co-Founder of Multiflow Computer

Awards [Eckert-Mauchly Award](#), (IEEE/ACM 2003)
[B. Ramakrishna Rau Award](#) (IEEE-CS 2012)
Connecticut Entrepreneur of the Year (1987)
Presidential Young Investigator's Award (NSF 1984)

Scientific career

Fields Computer Architecture, Compiling, Embedded Systems

Institutions Yale University, Multiflow Computer, Hewlett-Packard Laboratories (retired)

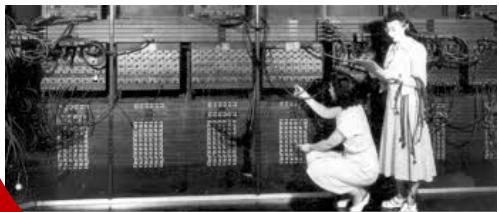


NYU

TANDON SCHOOL
OF ENGINEERING

Computers over time

ENIAC
1945



IBM 360
1964



Pentium Pro
1995
P6 cores!



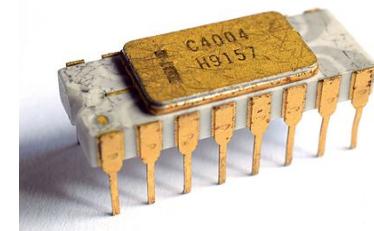
Apple A12 SoC



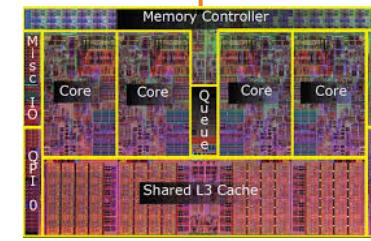
EDVAC
1949



Intel 4004
1971



Intel Nahelem
2008



MICRO 2020

happening soon



Type	Early <i>until 10/10</i>	Late <i>after 10/10</i>
ACM/IEEE Member	EUR 40	EUR 55
ACM/IEEE Lifetime/Retired Member	EUR 25	EUR 35
Non-Member	EUR 60	EUR 80
ACM/IEEE Student Member	EUR 25	EUR 35
Student Non-Member	EUR 40	EUR 55



NYU

TANDON SCHOOL
OF ENGINEERING