

Virtual Memory and Dynamic Scheduling

Computer Architecture
ECE 6913

Brandon Reagen



NYU

TANDON SCHOOL
OF ENGINEERING

Announcements

1) Exam

- 1) Taking a while to grade the branch prediction question

2) Lab2

- 1) < Take pulse >
- 2) Please check-in with TAs with any questions!

3) Lab3

- 1) Will assign this weekend

4) Today: Finish VM, intro OoO

- 1) Take it easy, been a crazy week! (And we're still ahead)



NYU

TANDON SCHOOL
OF ENGINEERING

Class “poll” results

You were all relatively positive..

If you have complaints, please let me know!

Major feedback:

- Too fast
- Hard to hear, can't see when I walk to board
- Want list of papers/readings to accompany lecture
- Grading unfair: we're willing to work with you

My goal is to have the grading as fair as possible.



NYU

TANDON SCHOOL
OF ENGINEERING

CSAW (FREE)!

Learn about security and privacy

CSAW'20

**GLOBAL CONFERENCE
AGENDA**

<https://www.csaw.io/agenda>



NYU

TANDON SCHOOL
OF ENGINEERING

Virtual Memory



NYU

**TANDON SCHOOL
OF ENGINEERING**

Virtual memory: Terminology

Page: A single unit of virtual memory. Like a line/block in caches

Frame: A single unit of physical memory (the page of PM)

Page fault: “Miss” in the virtual memory. No mapping exists between virtual and physical addresses, bring in from disk

VM address: Address produced by the program/processor by the program

PM address: Address used to access main memory

Address translation: Process of mapping virtual addresses to physical ones

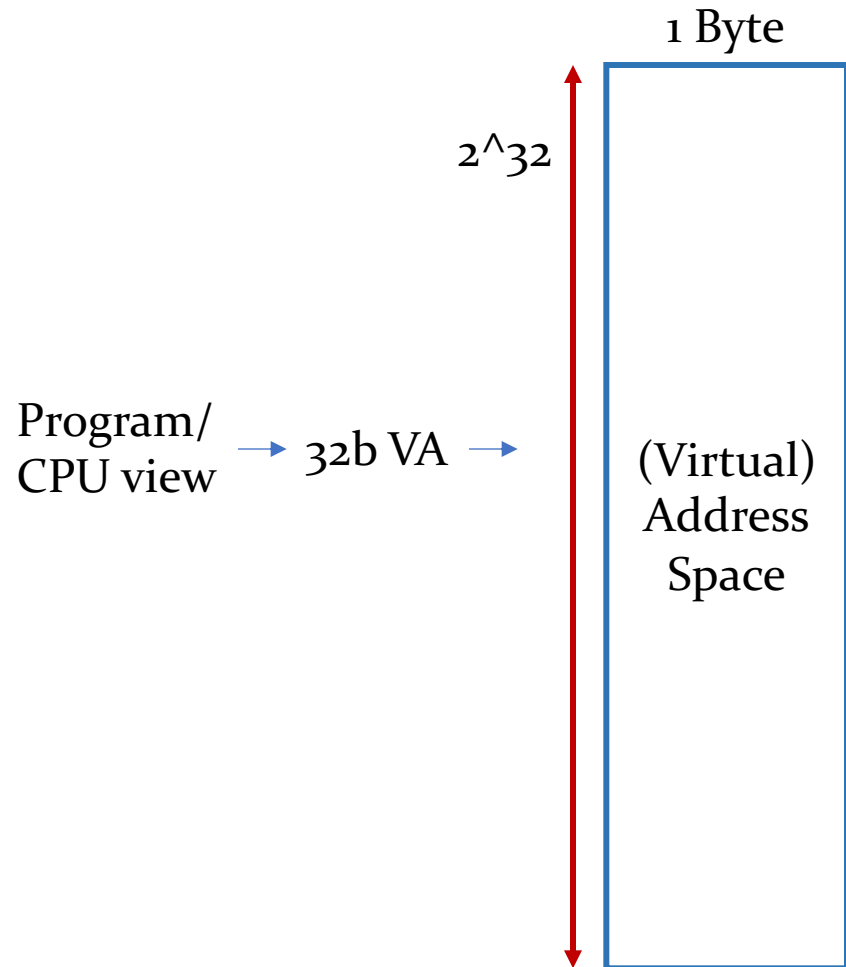
Relocation: Process of mapping physical addresses to virtual ones



NYU

**TANDON SCHOOL
OF ENGINEERING**

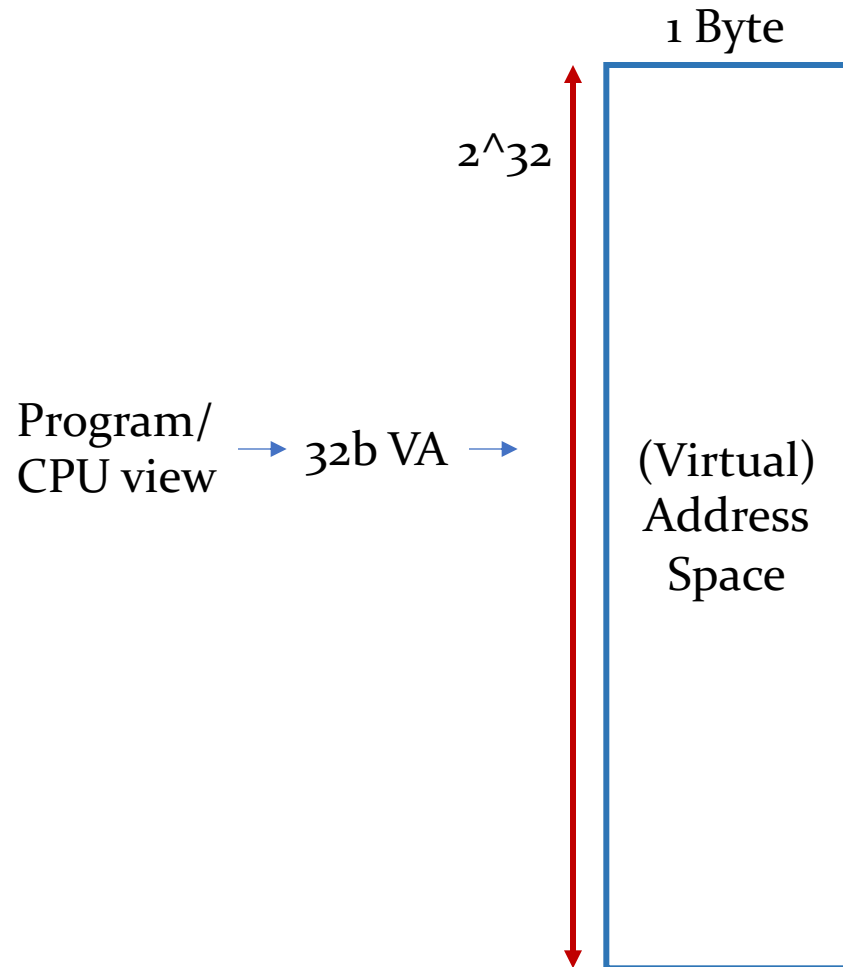
A programmer's perspective (from the clouds..)



NYU

TANDON SCHOOL
OF ENGINEERING

A programmer's perspective (from the clouds..)



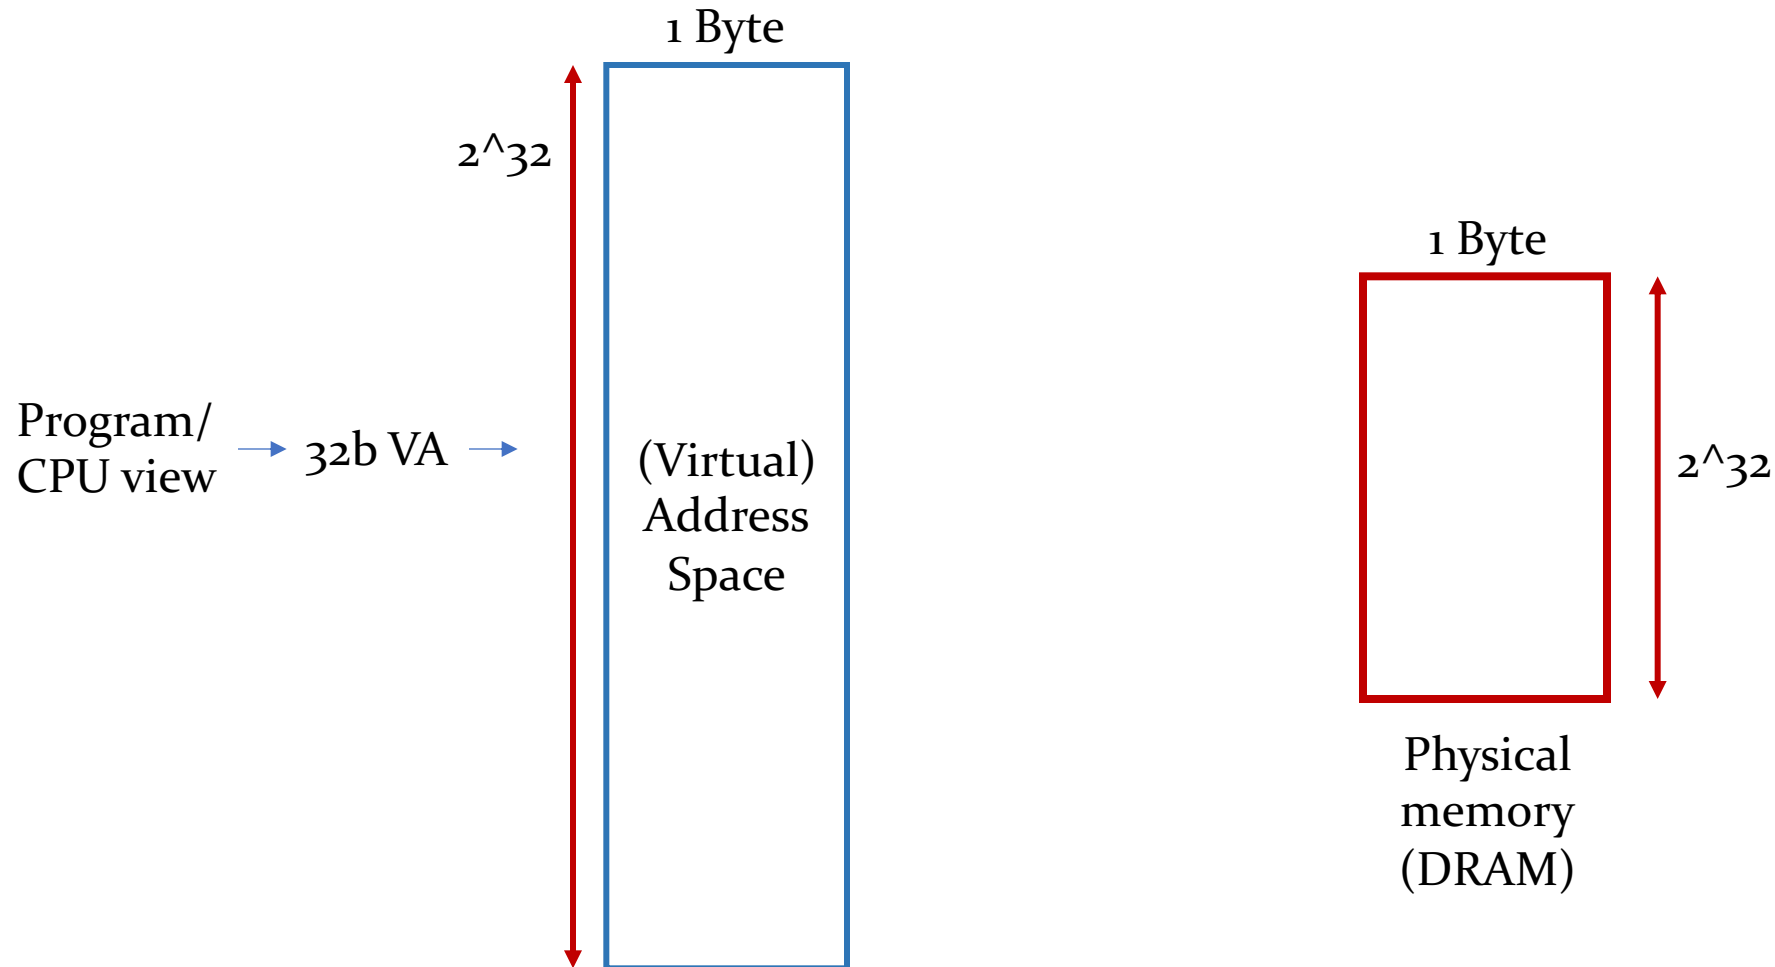
Once again the complexities
are dumped on the hardware and OS



NYU

TANDON SCHOOL
OF ENGINEERING

Problem 1: How does everything fit?



Each program needs
4GiB in a 32b machine

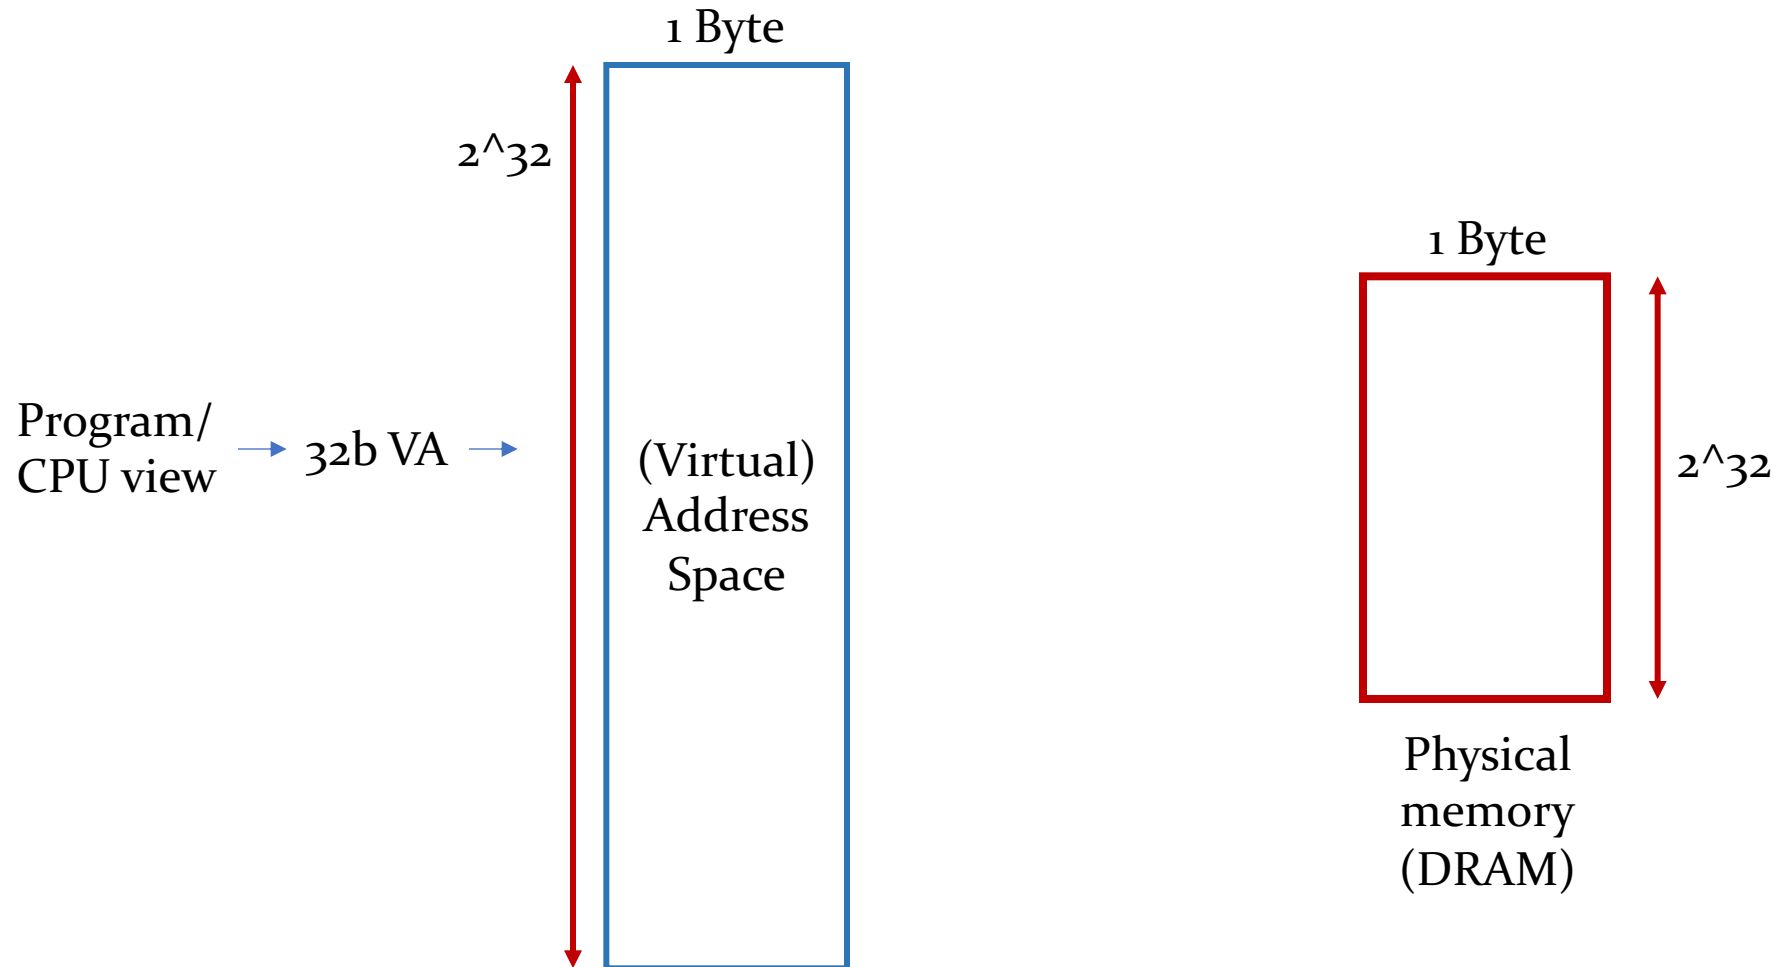
You have tens-hundreds
of programs running right
now, **how does it fit?**



NYU

TANDON SCHOOL
OF ENGINEERING

Problem 1: How does everything fit?



Each program needs
4GiB in a 32b machine

You have tens-hundreds
of programs running right
now, **how does it fit?**

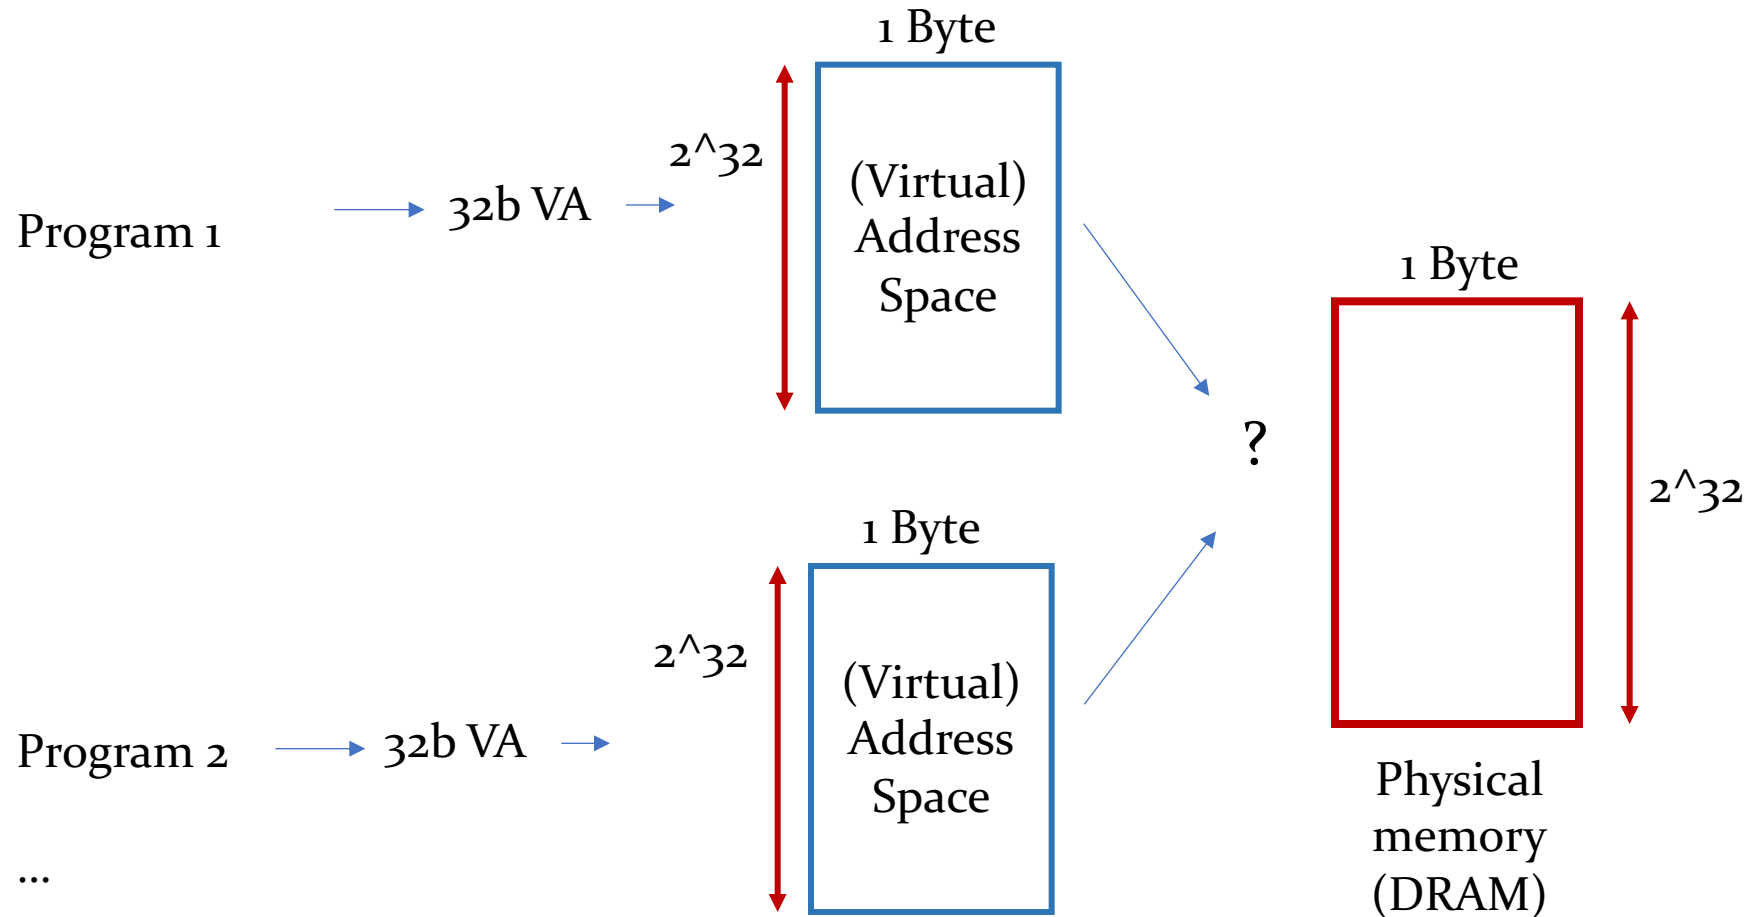
It doesn't..



NYU

TANDON SCHOOL
OF ENGINEERING

Problem 2: Where does everything go?



Each program has a VA space larger than the machine's physical memory.

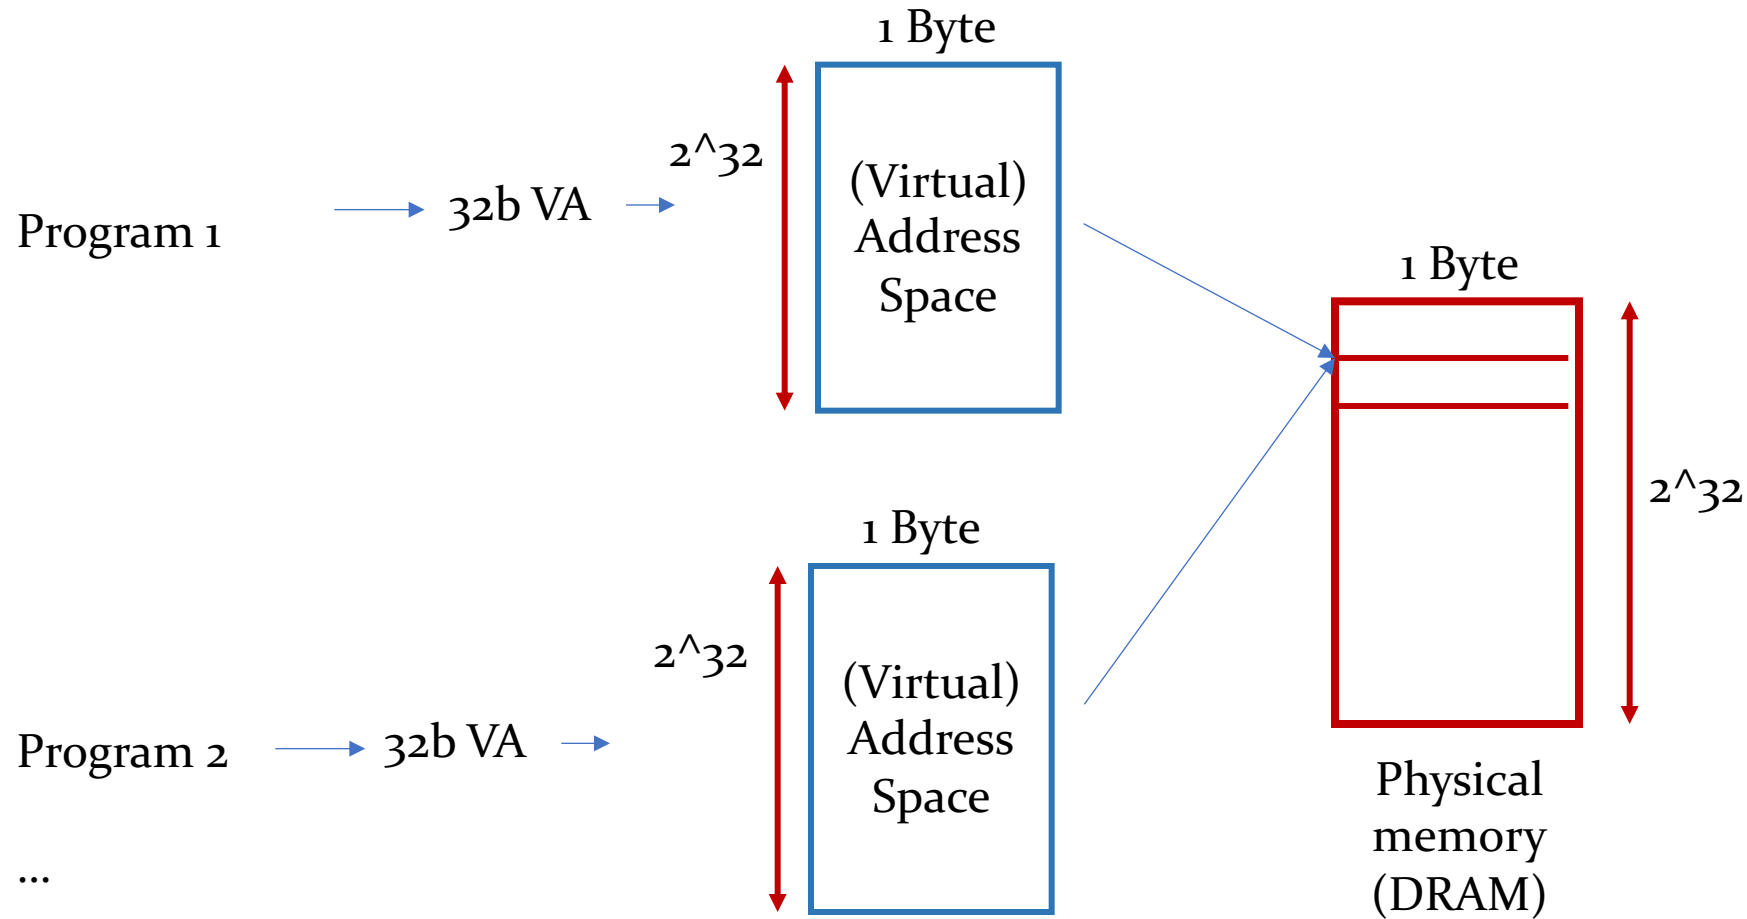
How should we divide up the limited physical memory resource?



NYU

TANDON SCHOOL
OF ENGINEERING

Problem 3: Memory protection



Protect against the benign and malicious case.

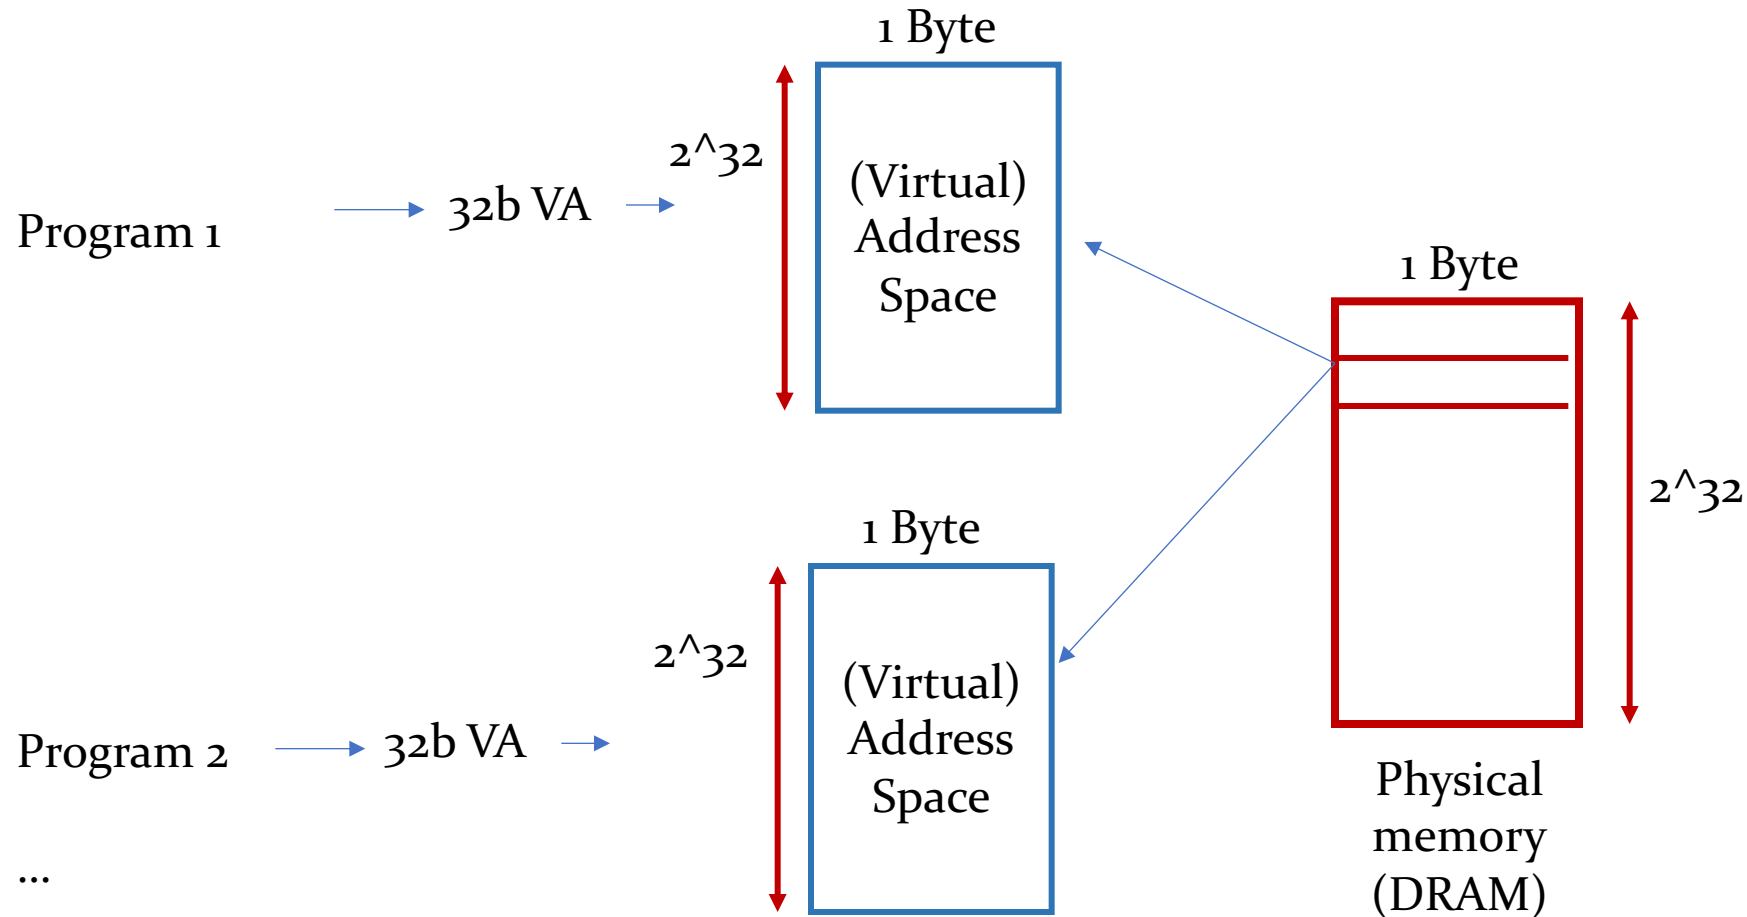
What's that mean?



NYU

TANDON SCHOOL
OF ENGINEERING

Problem 4: What about sharing?



Why would we want to share?

What could we share?



NYU

TANDON SCHOOL
OF ENGINEERING

How do we solve all these problems?

Hint: famous saying in computer science.

Indirection!

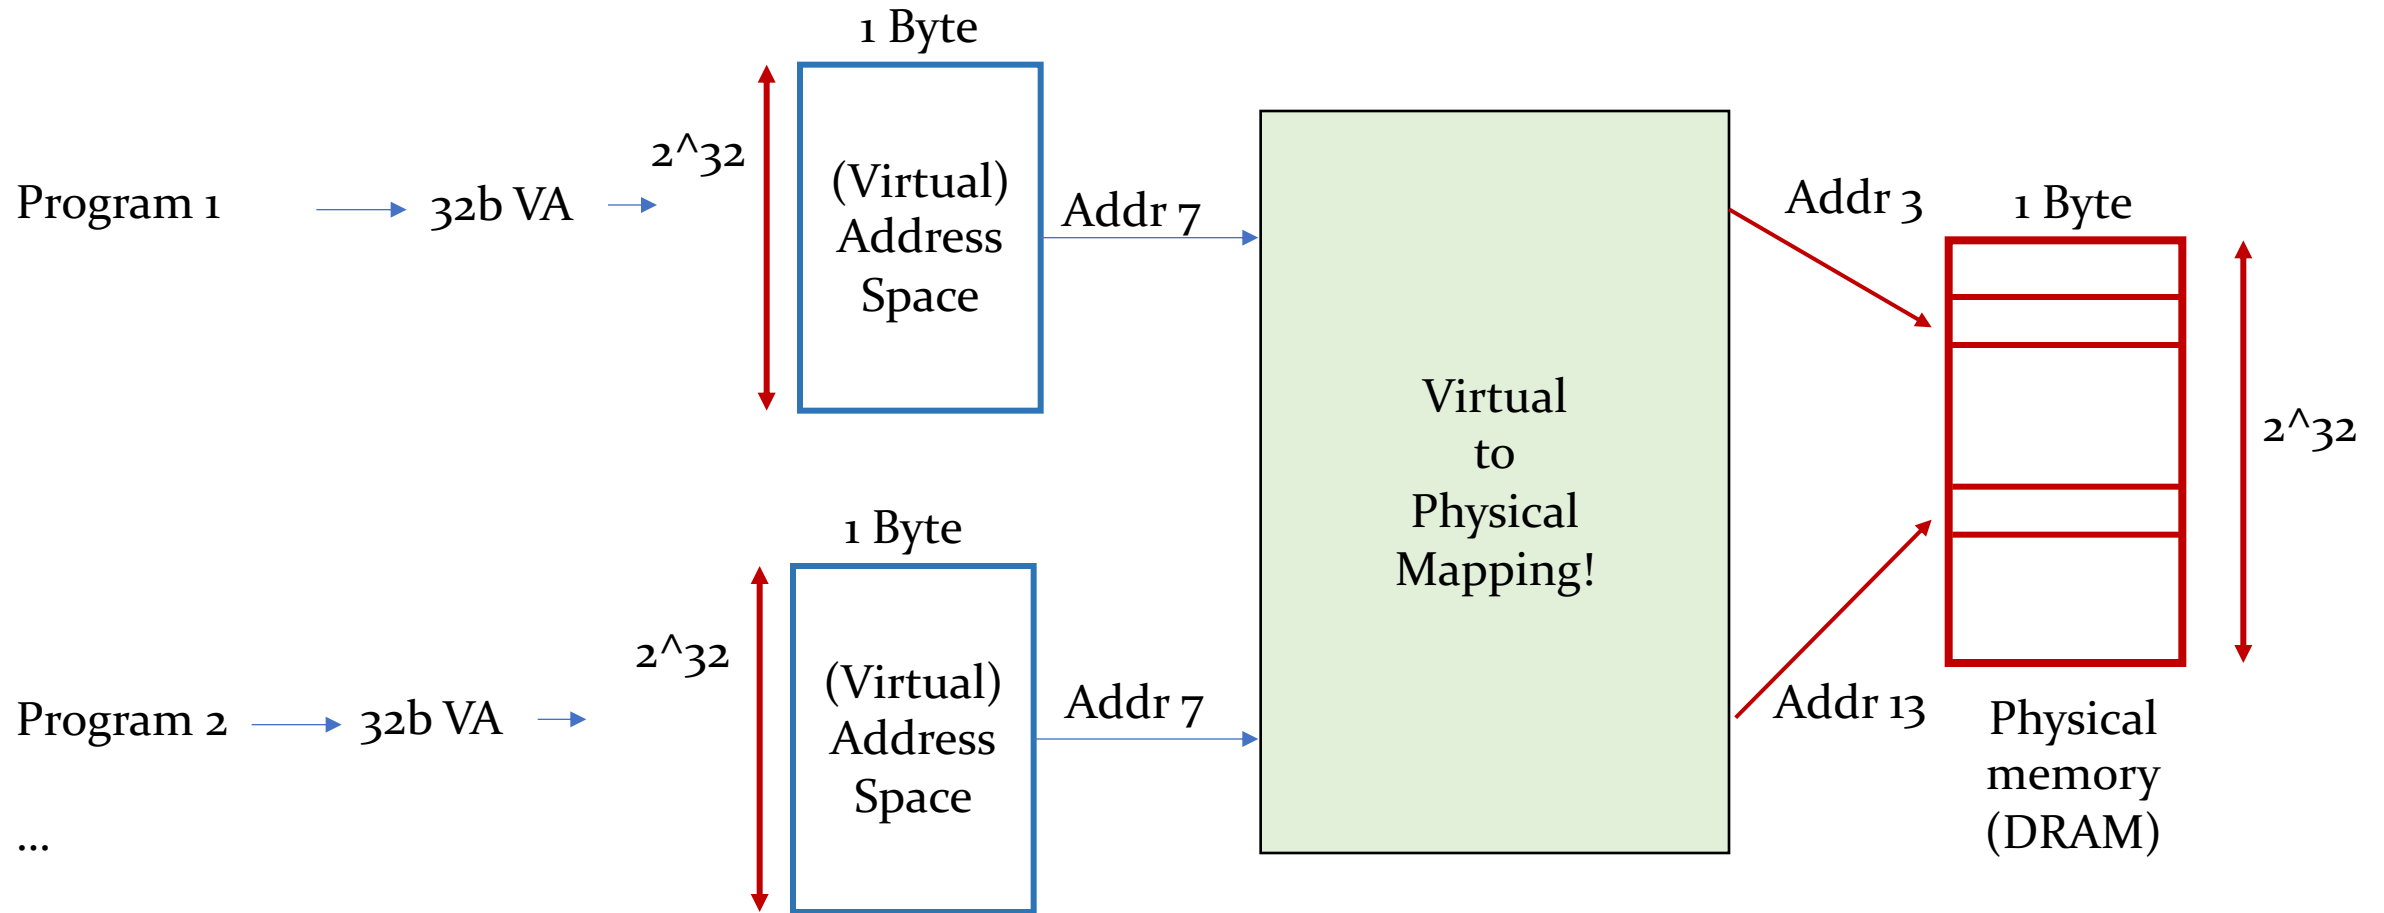
“All problems in CS can be solved by another level of indirection”
- David Wheeler



NYU

TANDON SCHOOL
OF ENGINEERING

Add a level of indirection: Virt -> Phys mapping



The hardware that handles address mapping, or **Translation**, is called the **MMU**: Memory Management Unit



NYU

TANDON SCHOOL
OF ENGINEERING

What does this VM (indirection) fix?

1) Isolation of program address spaces

- 1) Programs can't access each other's memory (regardless of intent)
- 2) Users can't access privileged memory (OS code)

2) Significantly simplifies programmer's life!

No need to worry about any of this:

- 1) Write code like have all address space
- 2) Don't worry about overwriting
- 3) Code portable to machine with any size memory!

3) Provides means to use physical memory (DRAM) as a cache

- 1) Memory space too big, so only store “hot” pages in DRAM
- 2) Everything must live somewhere, so there we must rely on disk.

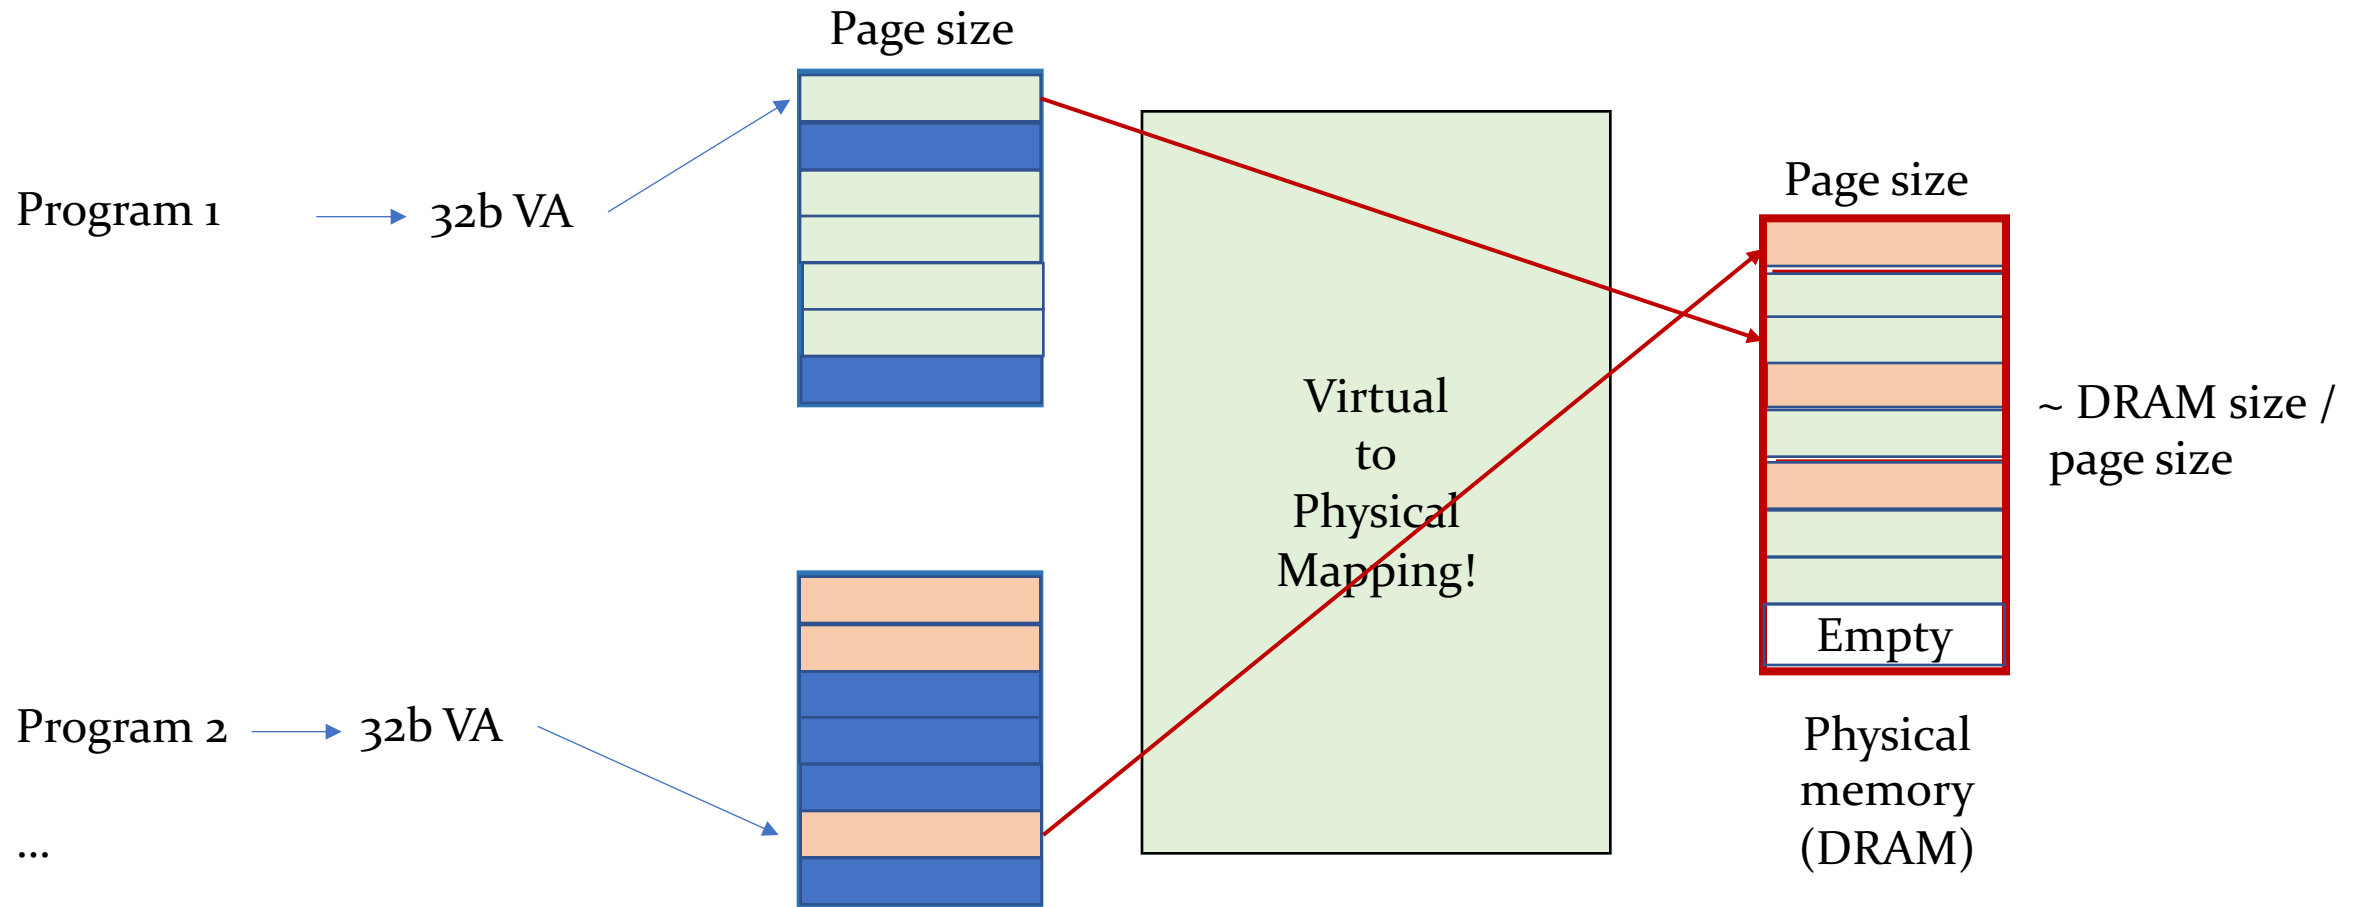
What's happening
here?



NYU

TANDON SCHOOL
OF ENGINEERING

What it looks like

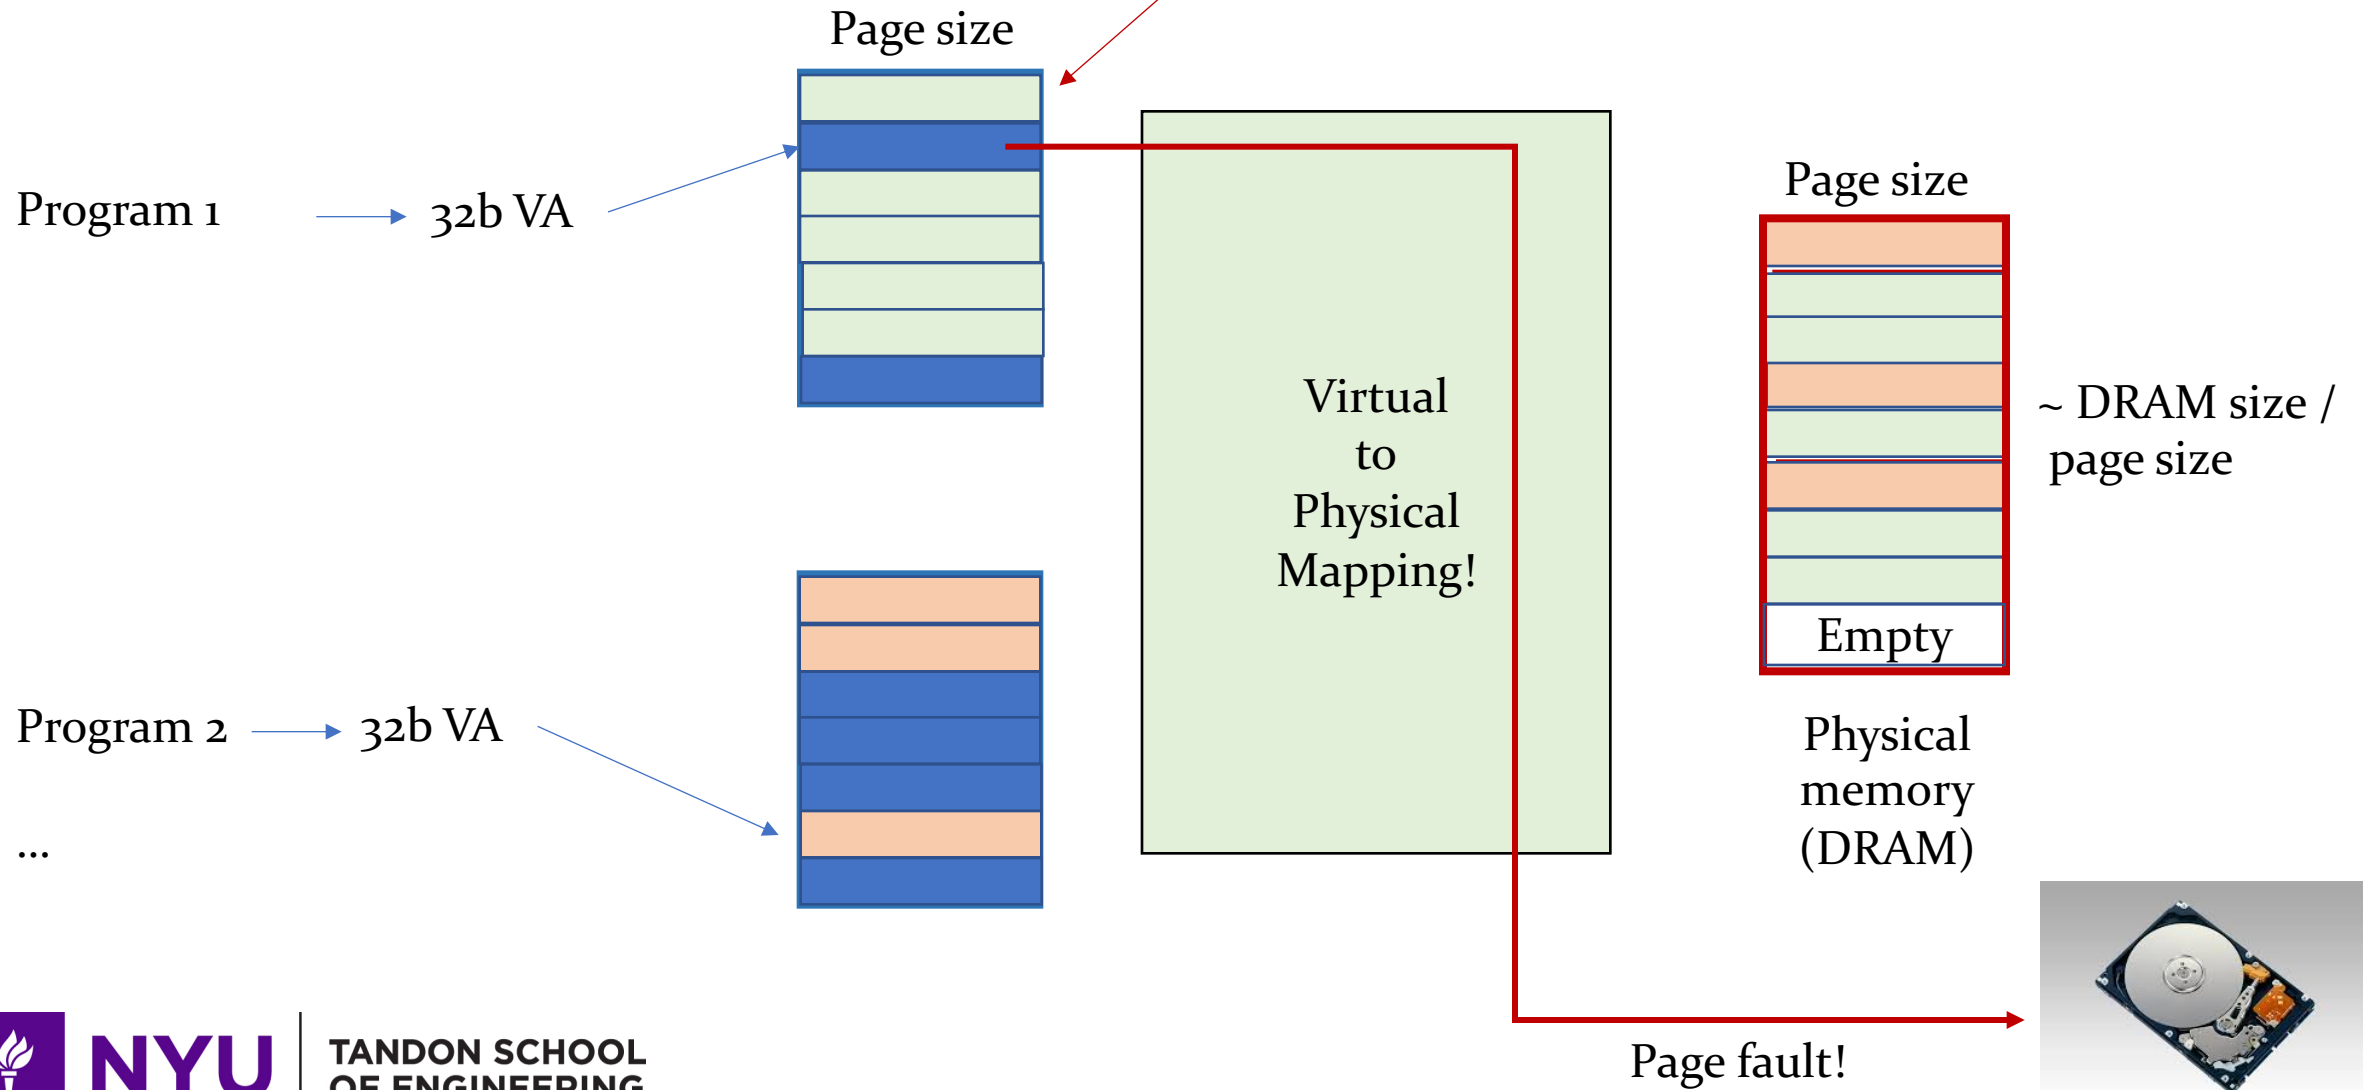


NYU

TANDON SCHOOL
OF ENGINEERING

What it looks like

Remember, this isn't real!
It's just an illusion.



Addressing memory

Virtual address has two parts: **virtual page number** and **page offset**

The size of a page is determined by the page offset

This means we don't have to translate it!

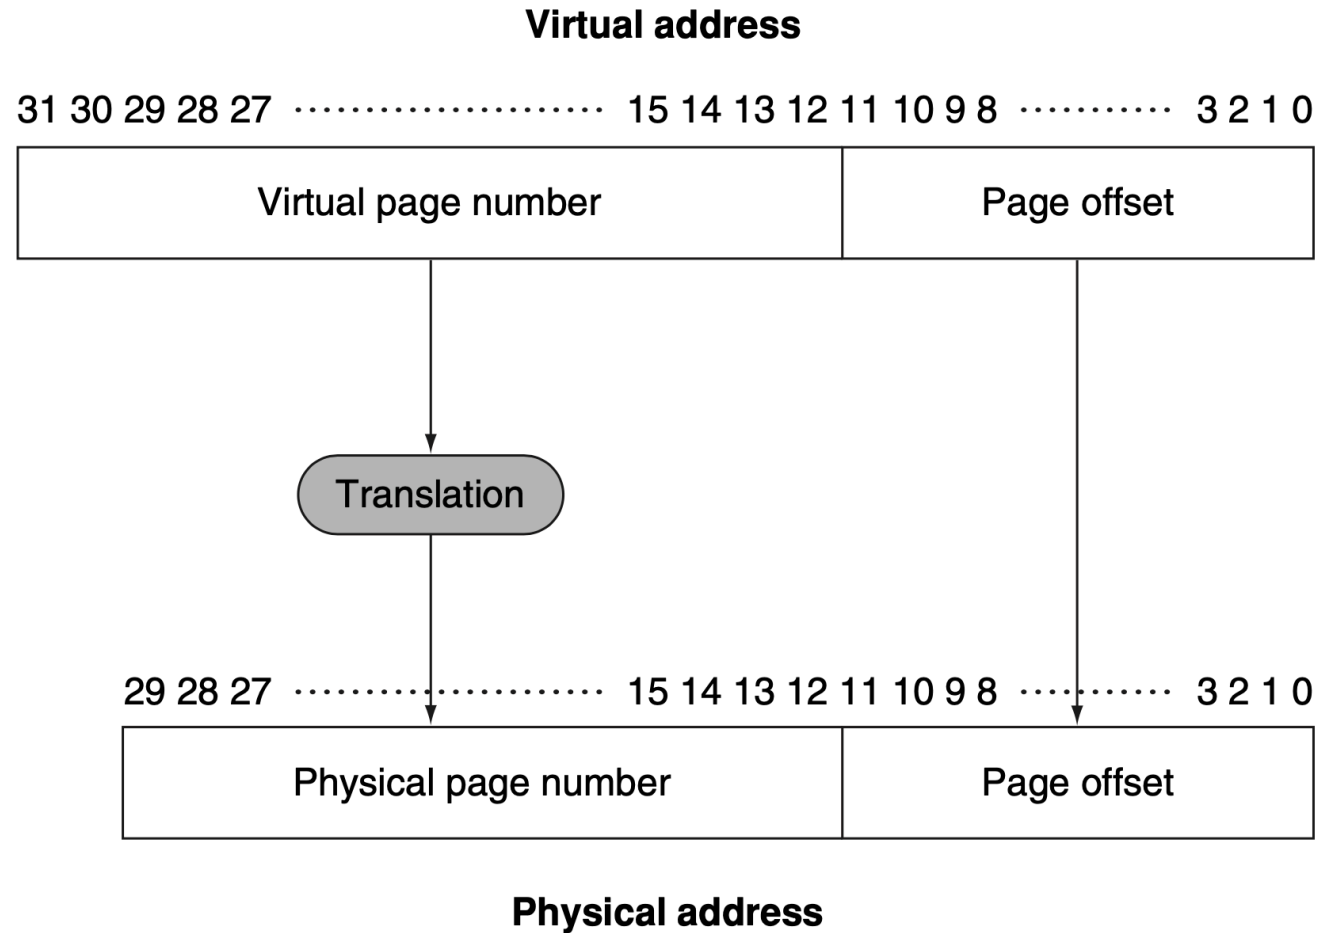
Translation: Convert a program's virtual page number to the machine's physical page (frame) number



NYU

TANDON SCHOOL
OF ENGINEERING

Address translation - simplified



Placing and finding pages

Placement:

OS can map any virtual page to any physical page

Using full associativity minimizes conflicts for limited resource

What was the problem with fully-associative caches?

Finding is hard.

Takes long time, or requires expensive parallel access

VM uses a **Page Table**:

a special, auxiliary data structure to track VA \rightarrow PA mappings



NYU

TANDON SCHOOL
OF ENGINEERING

Page tables

A page table takes virtual addresses as input and outputs physical address

Track **all** virtual to physical mappings

This means they can get quite large

Page tables live in memory

A special register (**page table register**)
points to the start of each processes page table

Quiz: How large is a single program's page table given:

32-bit virtual address, 4 KiB pages, and 4 bytes per page table entry



NYU

TANDON SCHOOL
OF ENGINEERING

Page tables

A page table takes virtual addresses as input and outputs physical address

Track all virtual to physical mappings

This means they can get quite large

Page tables live in memory

A special register (**page table register**)

points to the start of each processes page table

Quiz: How large is a single program's page table given:

32-bit virtual address, 4 KiB pages, and 4 bytes per page table entry

Num entries: $2^{32} / 2^{12} = 2^{20}$.

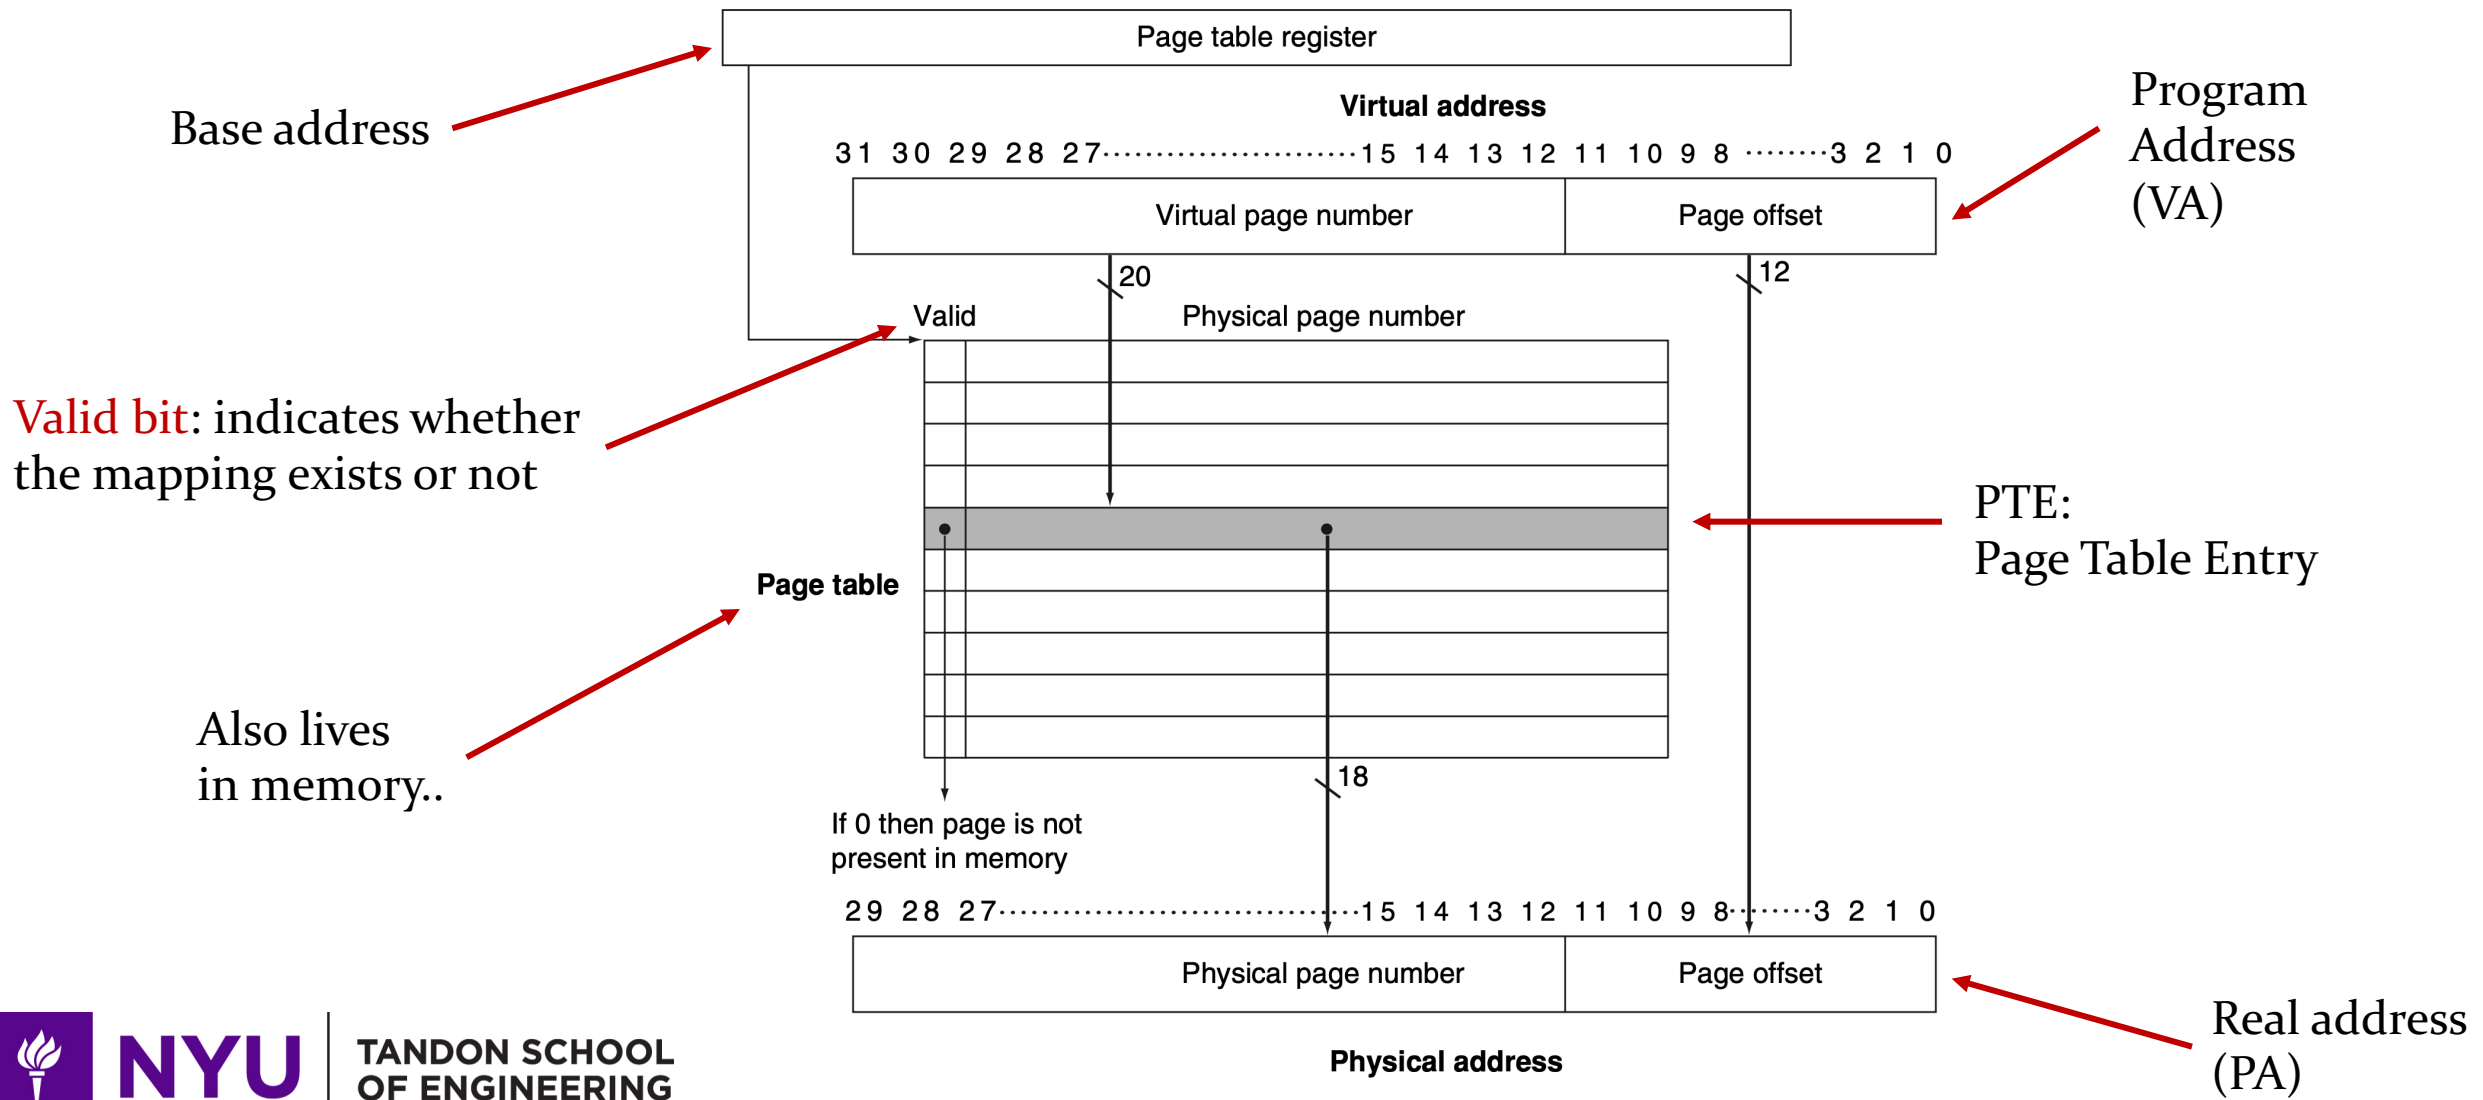
Total size: $2^{20} * 2^2 \text{ <bytes per PTE>} = 4 \text{ MiB}$



NYU

TANDON SCHOOL
OF ENGINEERING

Page table example



Page faults

If we access a PTE and the valid bit is zero we incur a **page fault**

Raise an exception and give control to the OS

OS will go out and find the data stored on disk, bring it into main memory, and create a valid mapping in the page table

When process starts, OS creates “**swap space**” on disk

This is a pre-allocation of all virtual addresses

OS creates data structure for where all program's VM addresses are stored on disk

This can be part of the page table or separate structure

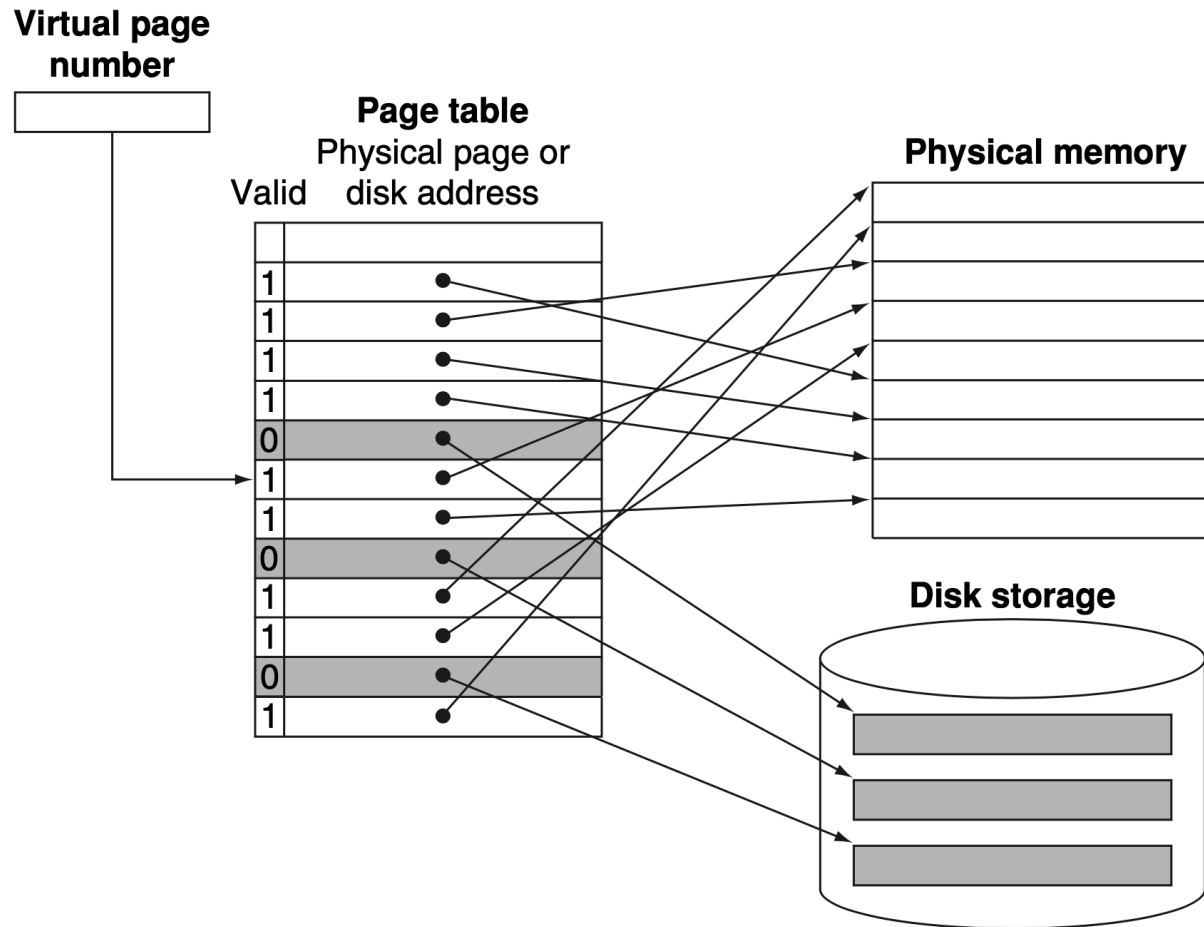
Usually separate structures, enables PT optimizations



NYU

TANDON SCHOOL
OF ENGINEERING

Page faults



Page faults are bad because we must go to disk.
Can't really get around that..

But we also don't like going to memory..
and the page table lives in main memory

What does this imply?

How can we fix this?



NYU

TANDON SCHOOL
OF ENGINEERING

Translation Lookaside Buffer (TLB)

Each time we need to access memory, need to make 2 trips!

One for translation, one for access event

We can **cache** virtual to physical address translations

This is what the TLB does:

Leverage locality to store recently used translations close to processor

TLB is why professors spend so long thinking about paper names and titles

Terrible name. Translation Cache much better.

Does this make sense?

- a) We have caches
- b) VM let's DRAM cache pages from disk
- c) TLB caches translations to avoid trip to memory



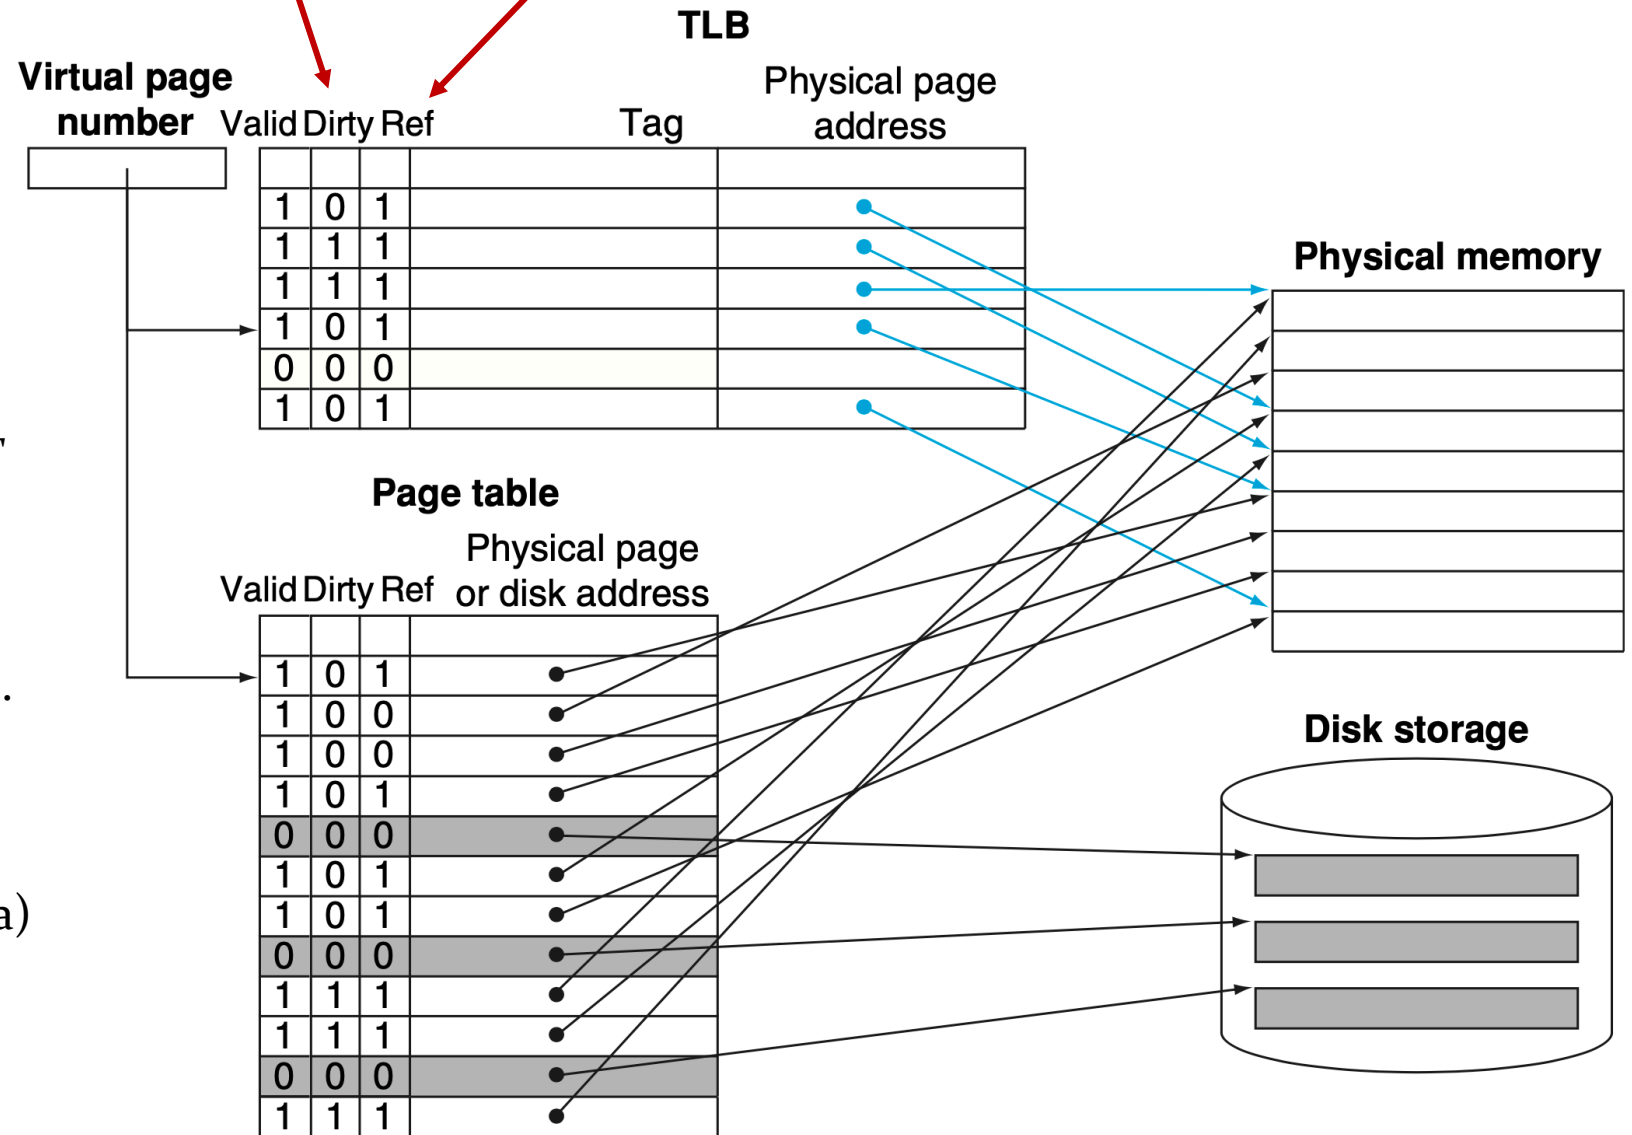
NYU

TANDON SCHOOL
OF ENGINEERING

TLB example

Pages can be dirty
(recall we use WB)

Ref bits used for
replacement



TLB holds a subset of all VA->PA translations

Hits in the TLB avoid going to the PT
Avoid 1 memory access!

Use high associativity,
we really don't want to go memory..

Note:
TLB is a cache (tag and data)
Page Table is **not a cache** (just data)



NYU

**TANDON SCHOOL
OF ENGINEERING**

TLB access process

On every memory event, access the TLB:

If TLB hit: the physical page number is used to make the physical address,
Ref bit set high. If write event: turn on dirty bit

If TLB miss: Determine if simple TLB miss or page fault

If **PTE valid**, processor looks up VA in PT,
loads PTE info in TLB, replays access

If **page fault**, invoke OS, trigger page fault event, go to disk,
create mapping in PT, load mapping into TLB, replay access

On eviction, only must write back reference and dirty bit
Mappings cannot change in the TLB!



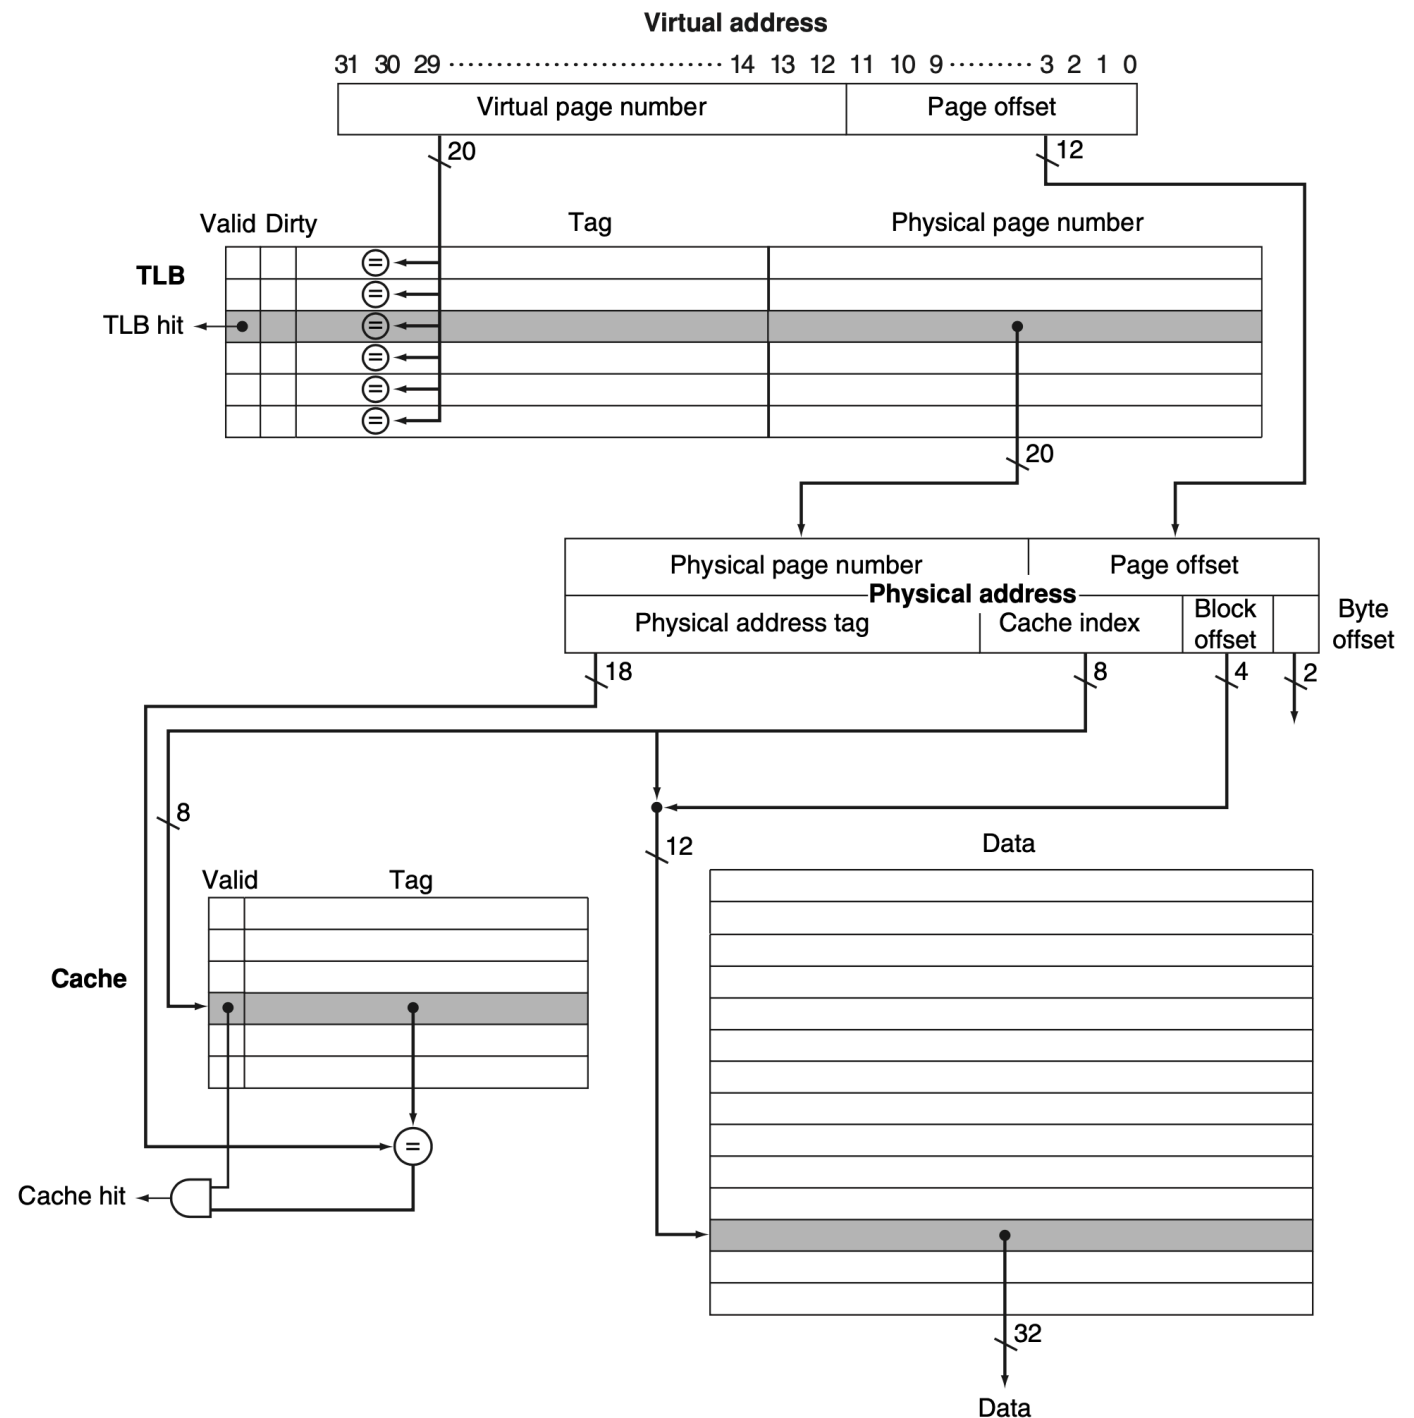
NYU

TANDON SCHOOL
OF ENGINEERING

Whole picture

Wait.. What's going on with this cache??
The data has more entries than tag array

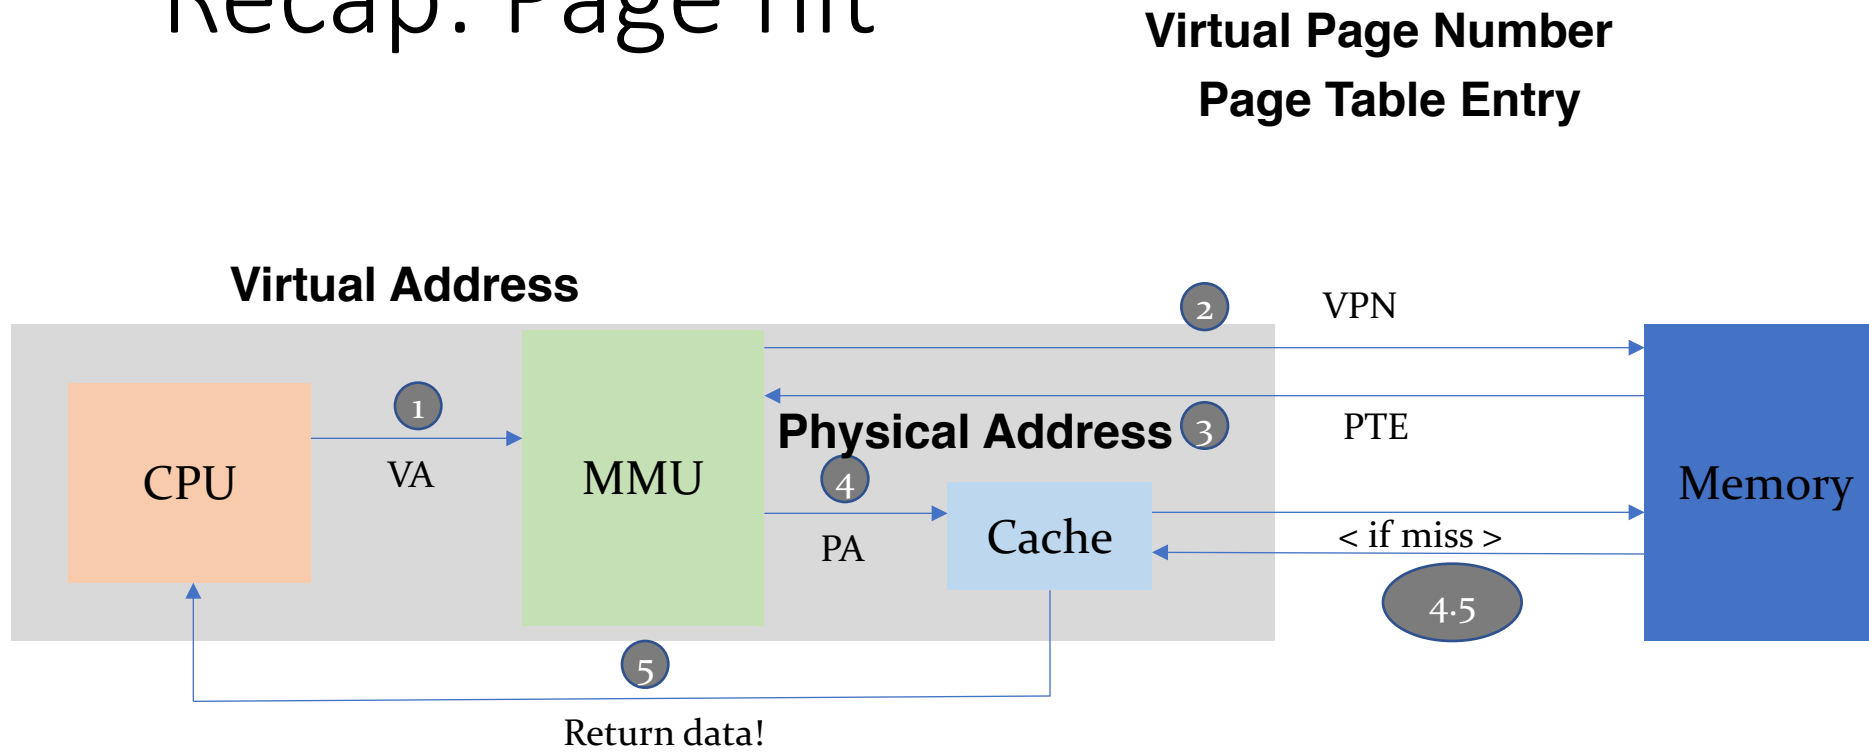
It's a DM cache that uses offset to index.
This saves a mux



NYU

TANDON SCHOOL
OF ENGINEERING

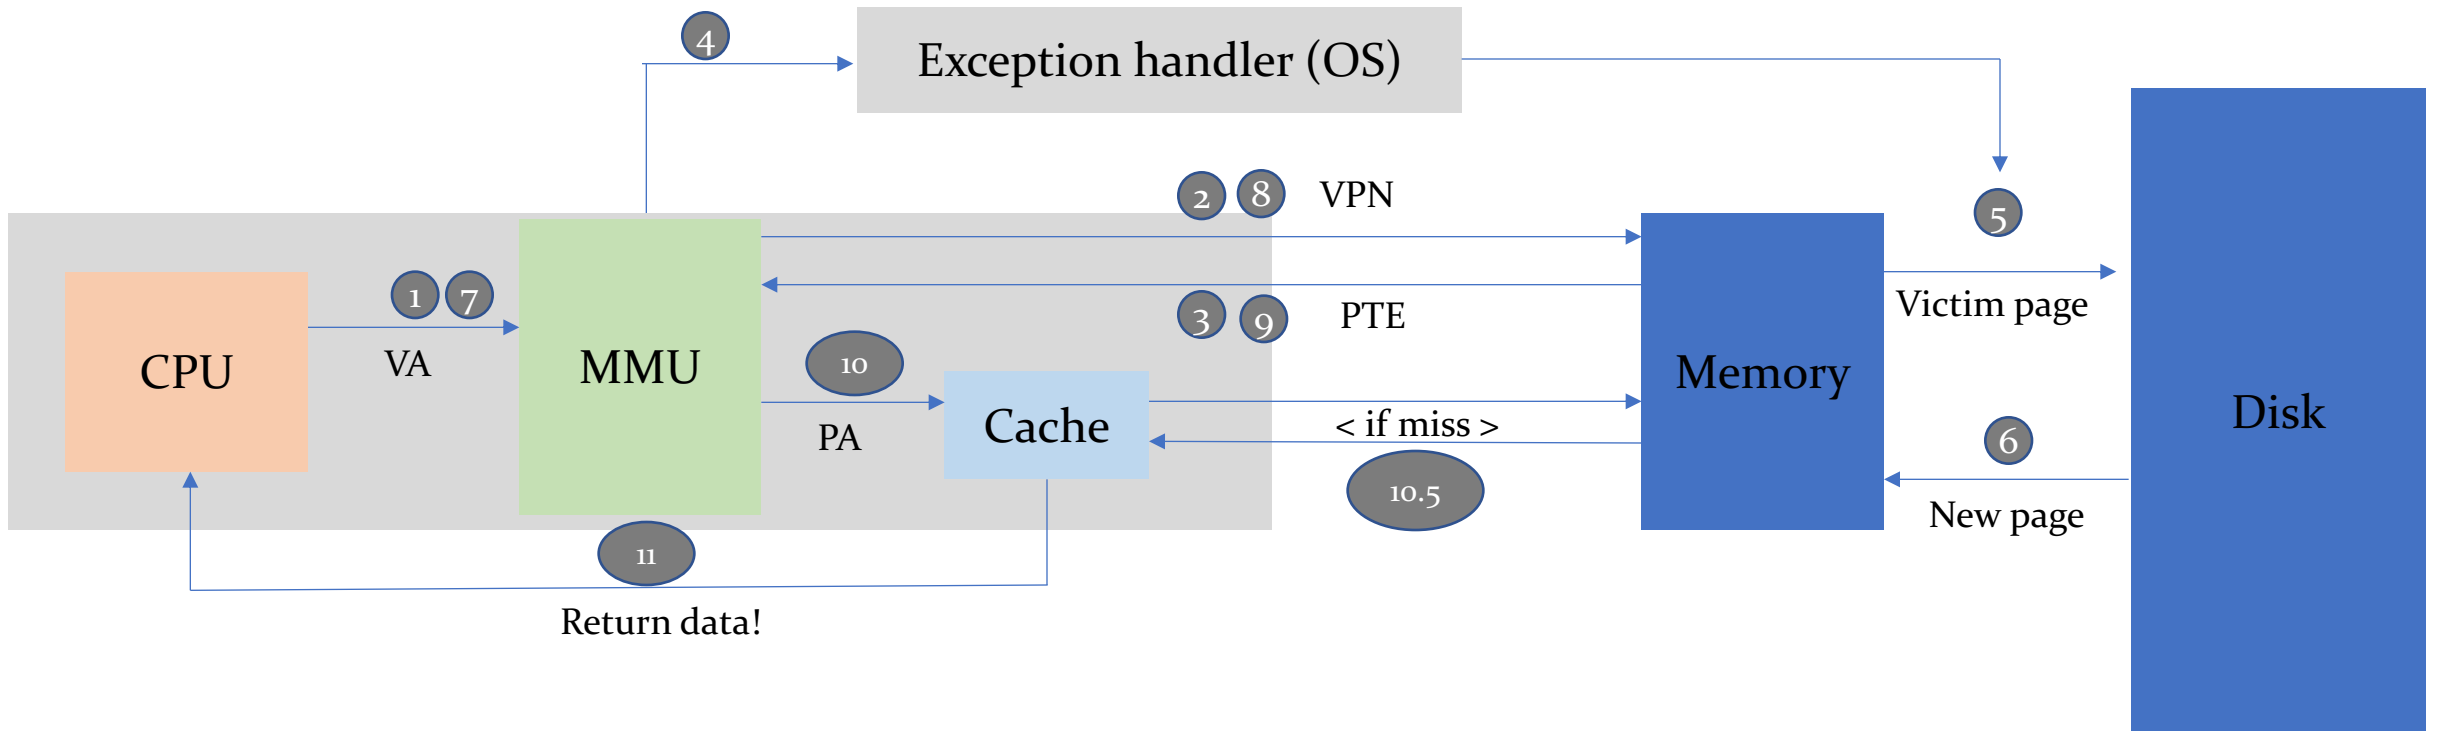
Recap: Page hit



NYU

**TANDON SCHOOL
OF ENGINEERING**

Recap: Page fault



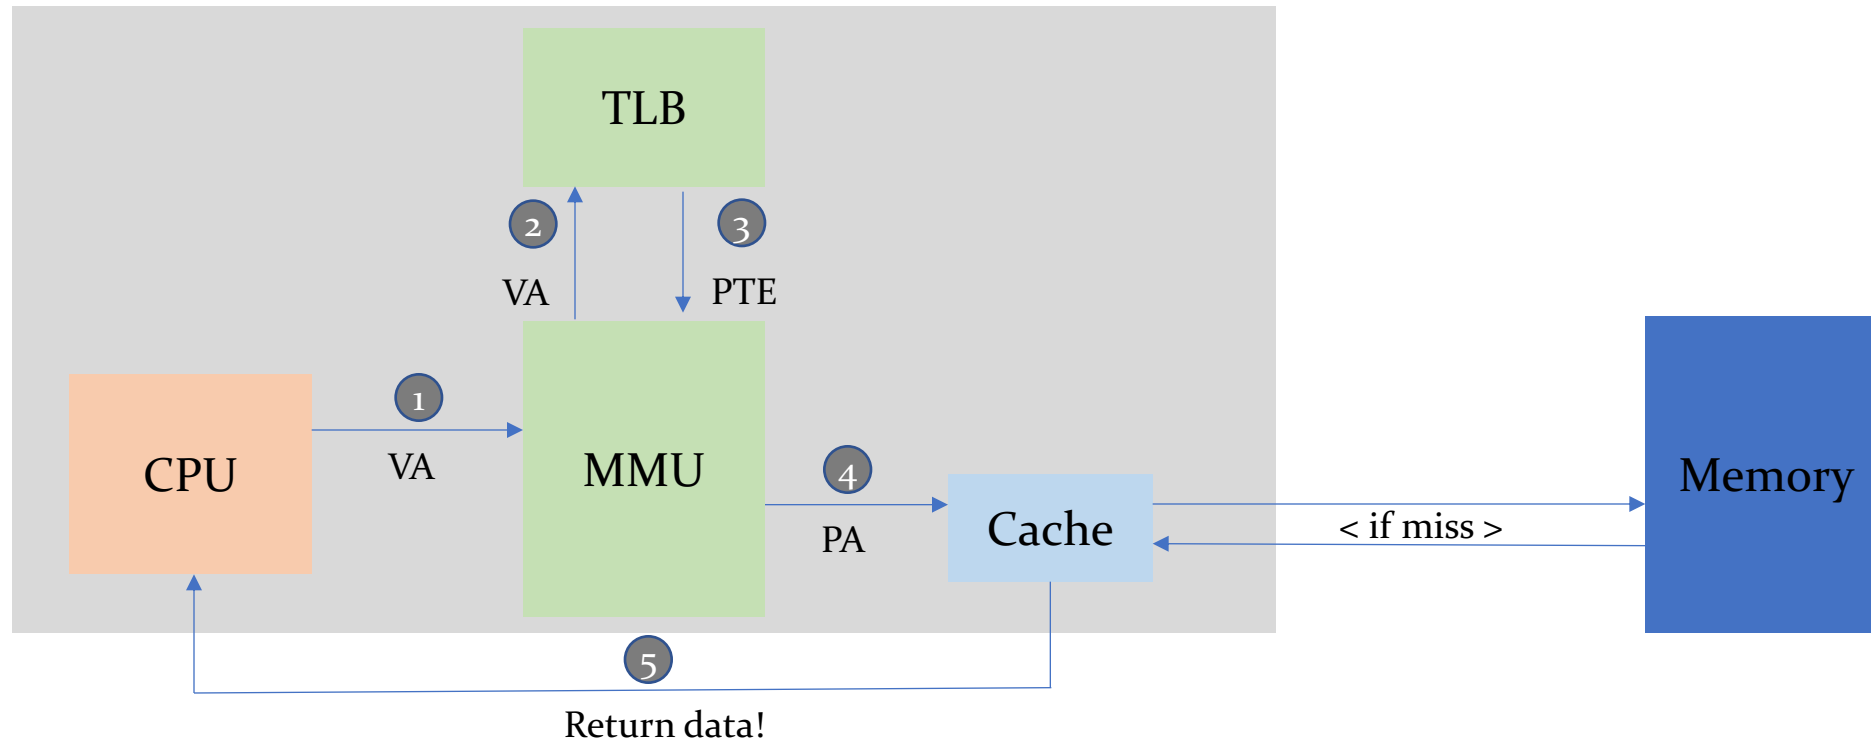
A lot of work..



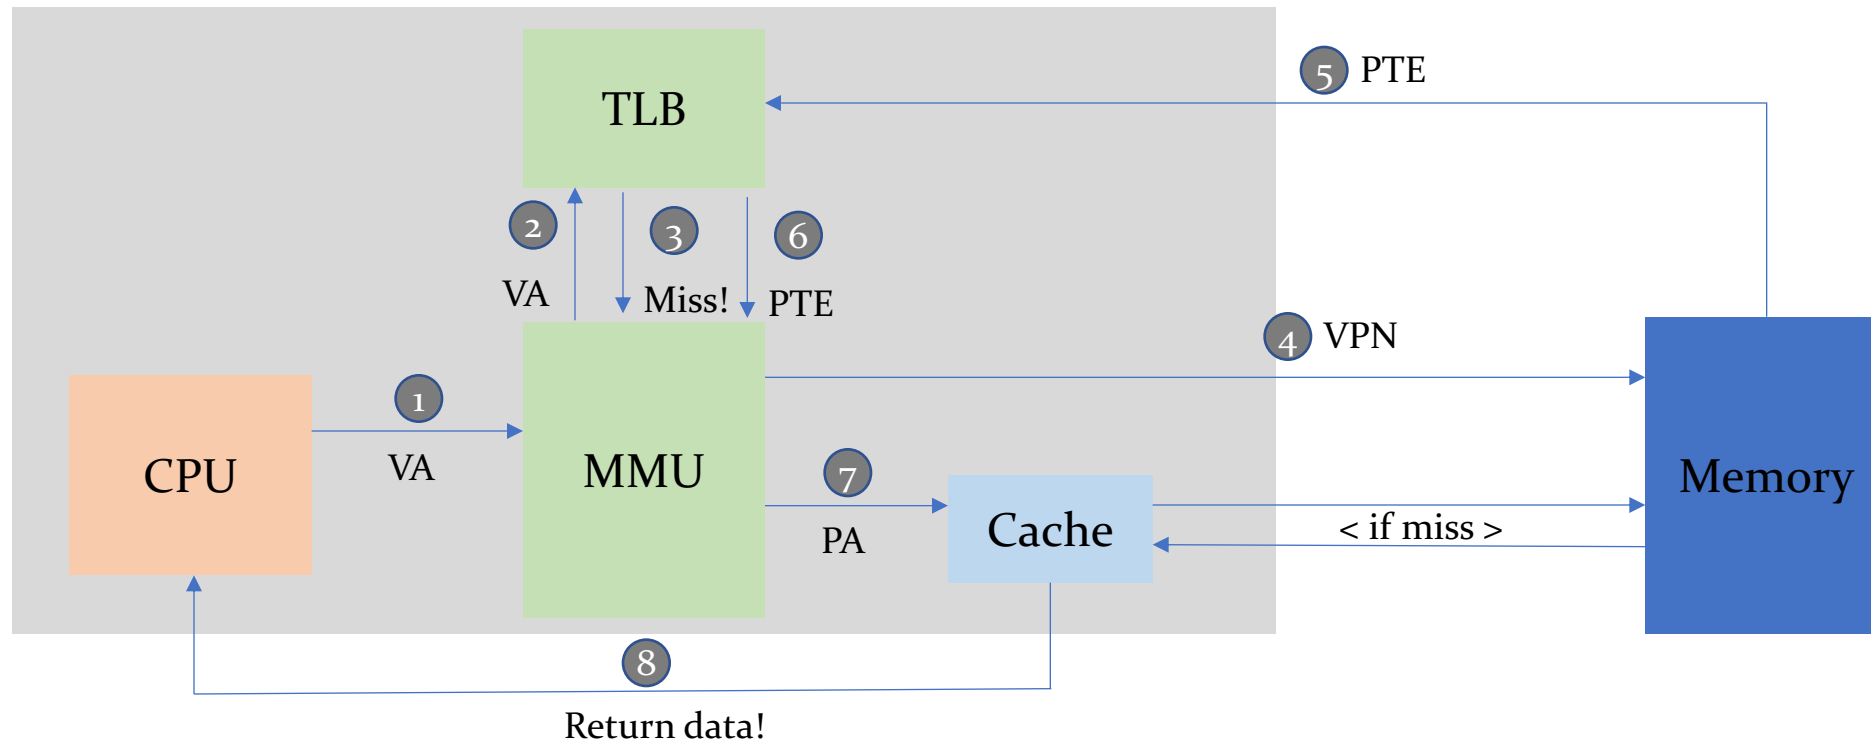
NYU

TANDON SCHOOL
OF ENGINEERING

Recap: TLB hit



Recap: TLB miss



Common interview question:
"How does a load work?"



NYU

TANDON SCHOOL
OF ENGINEERING

Examples!



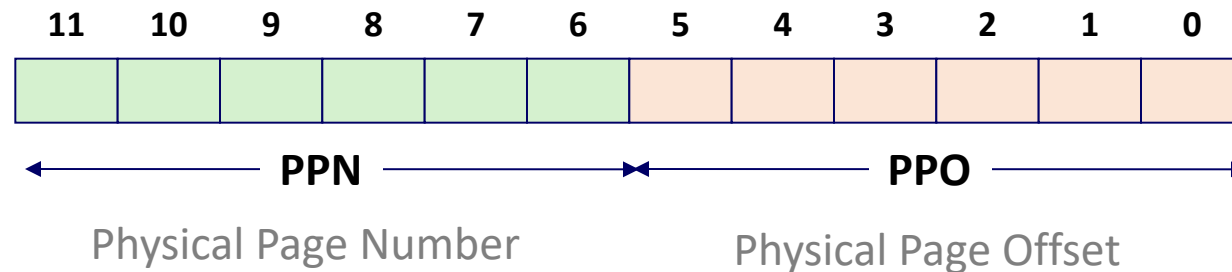
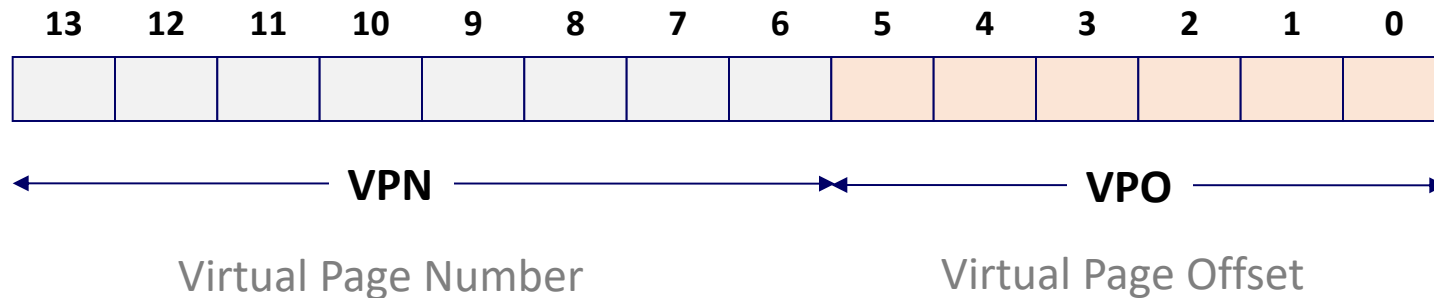
NYU

**TANDON SCHOOL
OF ENGINEERING**

Simple Memory System Example

Addressing

- 14-bit virtual addresses
- 12-bit physical address
- Page size = 64 bytes



Simple Memory System Page Table

Only show first 16 entries (out of 256)

<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
00	28	1
01	–	0
02	33	1
03	02	1
04	–	0
05	16	1
06	–	0
07	–	0

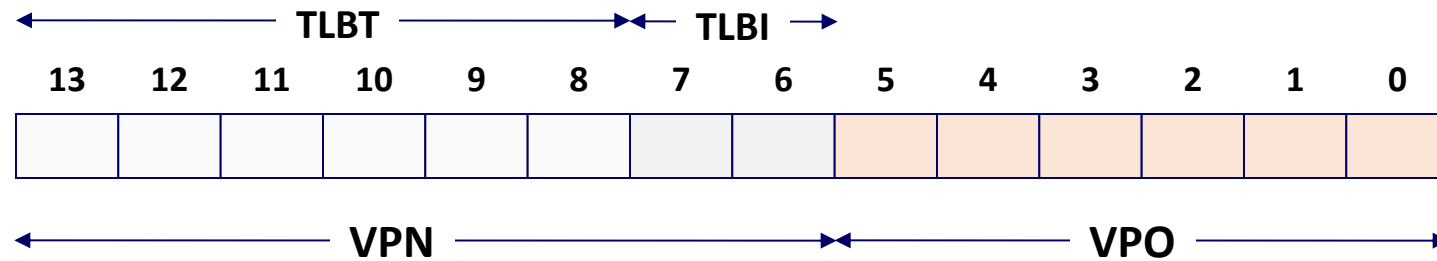
<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
08	13	1
09	17	1
0A	09	1
0B	–	0
0C	–	0
0D	2D	1
0E	11	1
0F	0D	1



Simple Memory System TLB

16 entries

4-way associative



<i>Set</i>	<i>Tag</i>	<i>PPN</i>	<i>Valid</i>	<i>Tag</i>	<i>PPN</i>	<i>Valid</i>	<i>Tag</i>	<i>PPN</i>	<i>Valid</i>	<i>Tag</i>	<i>PPN</i>	<i>Valid</i>
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0



NYU

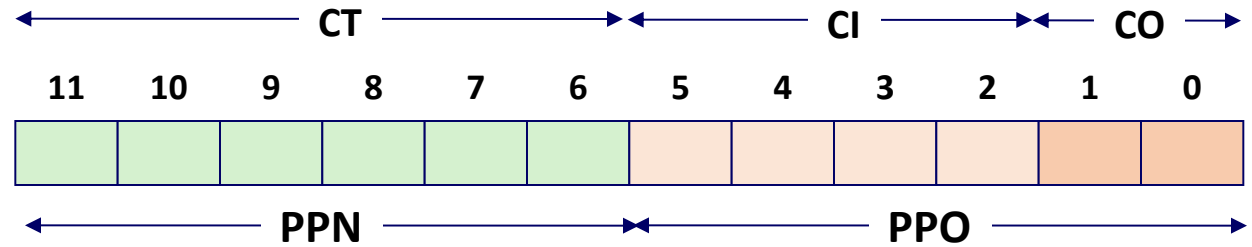
TANDON SCHOOL
OF ENGINEERING

Simple Cache

16 lines, 4-byte block size

Physically addressed

Direct mapped



<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03

<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–



NYU

TANDON SCHOOL
OF ENGINEERING

Current state of caches/tables

TLB

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0

VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	08	13	1
01	–	0	09	17	1
02	33	1	0A	09	1
03	02	1	0B	–	0
04	–	0	0C	–	0
05	16	1	0D	2D	1
06	–	0	0E	11	1
07	–	0	0F	0D	1

Cache

Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03

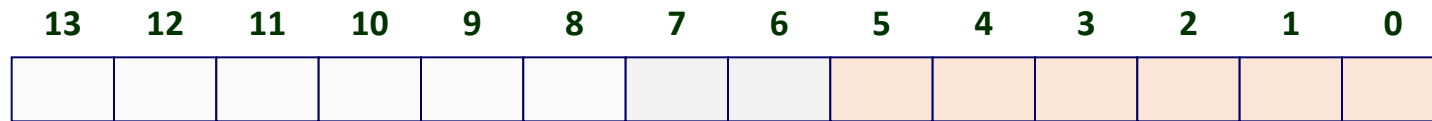
Page table

Idx	Tag	Valid	B0	B1	B2	B3
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–



Address Translation Example #1

Virtual Address: 0x03D4



TLB Index

VPN ____

TLBI ____

TLBT ____

TLB Hit? ____

Page Fault? ____

PPN: ____

Virtual Page Number

TLB Tag

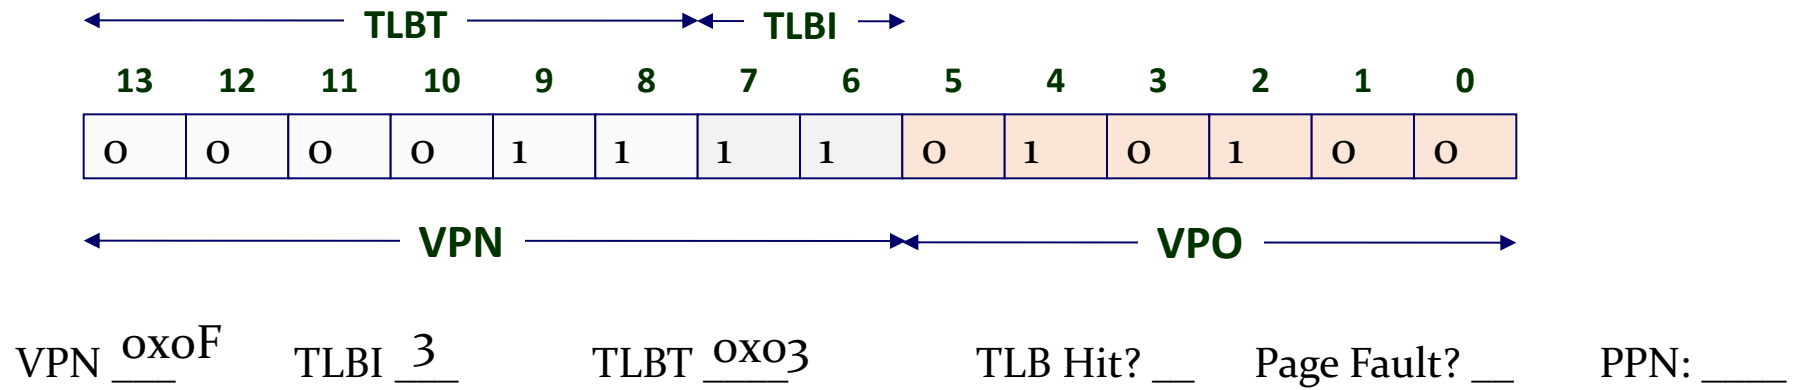


NYU

**TANDON SCHOOL
OF ENGINEERING**

Address Translation Example #1

Virtual Address: 0x03D4



NYU

TANDON SCHOOL
OF ENGINEERING

Current state of caches/tables

TLB

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0

VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	08	13	1
01	–	0	09	17	1
02	33	1	0A	09	1
03	02	1	0B	–	0
04	–	0	0C	–	0
05	16	1	0D	2D	1
06	–	0	0E	11	1
07	–	0	0F	0D	1

Cache

Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03

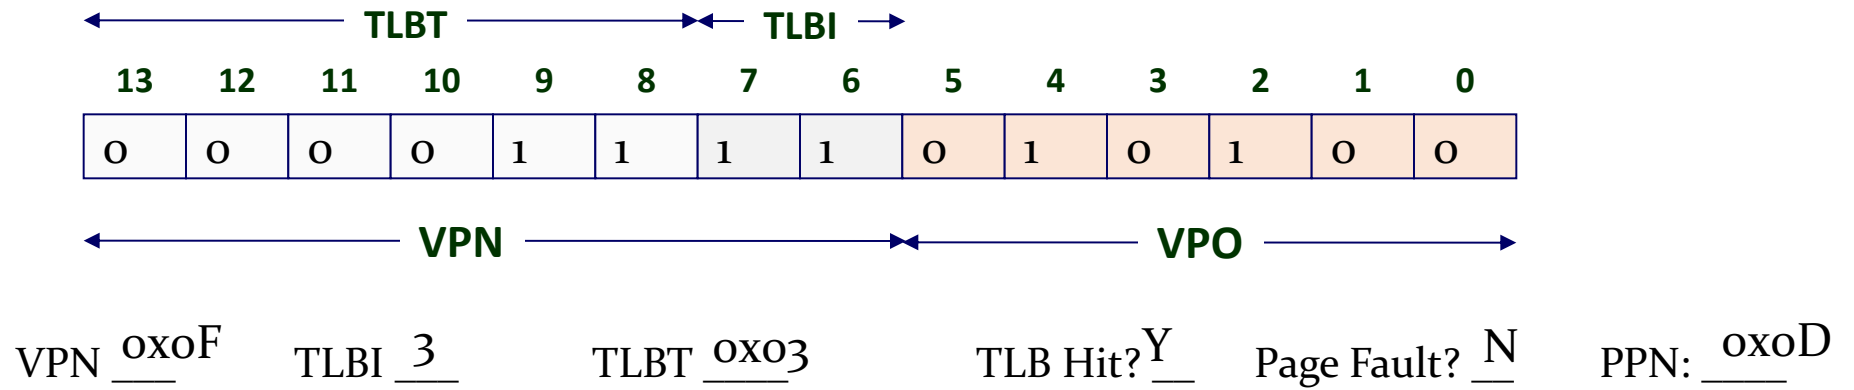
Page table

Idx	Tag	Valid	B0	B1	B2	B3
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–



Address Translation Example #1

Virtual Address: 0x03D4

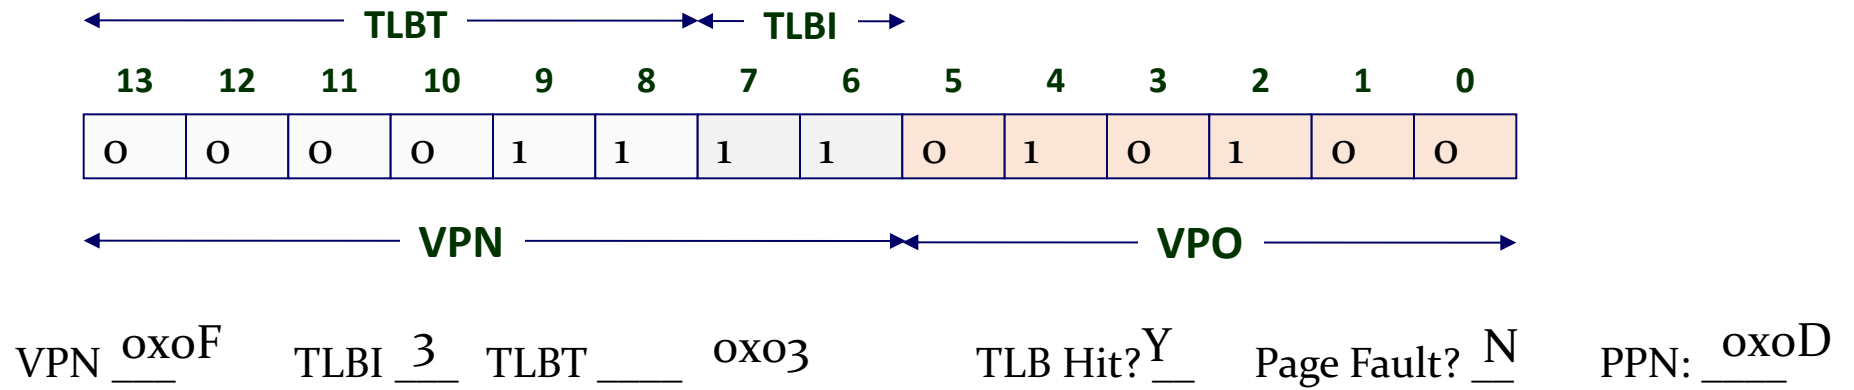


NYU

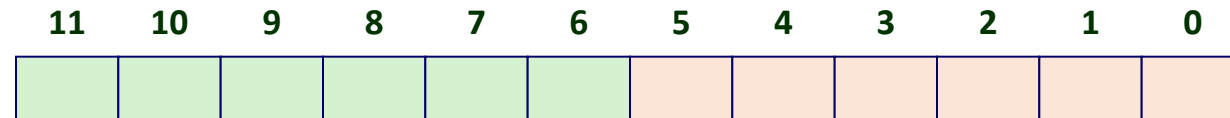
TANDON SCHOOL
OF ENGINEERING

Address Translation Example #1

Virtual Address: 0x03D4



Physical Address



16 entry
DM cache

CO

CI

CT

Hit?

Byte

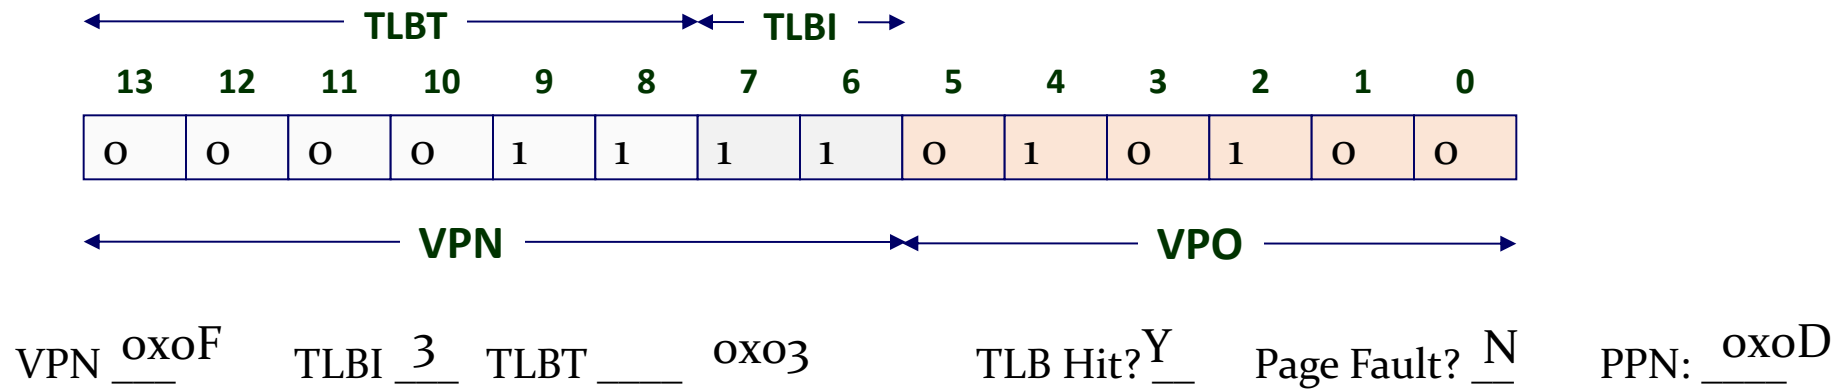


NYU


TANDON SCHOOL
OF ENGINEERING

Address Translation Example #1

Virtual Address: 0x03D4



Physical Address



16 entry
DM cache

Cache offset Cache index

CO

CI

CT

Hit?

Byte

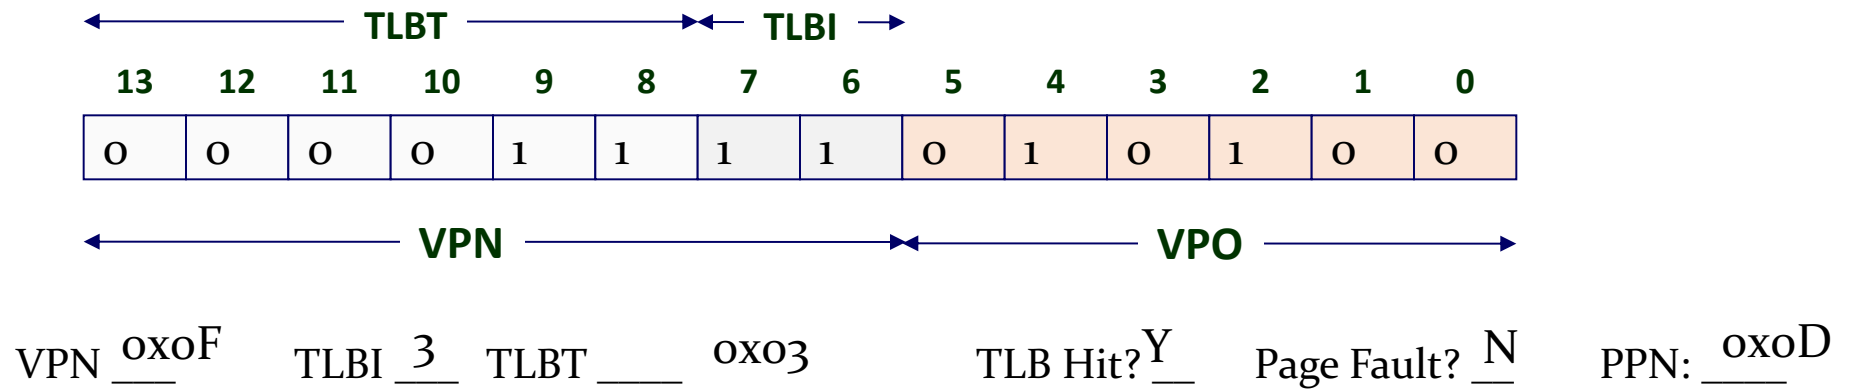


NYU

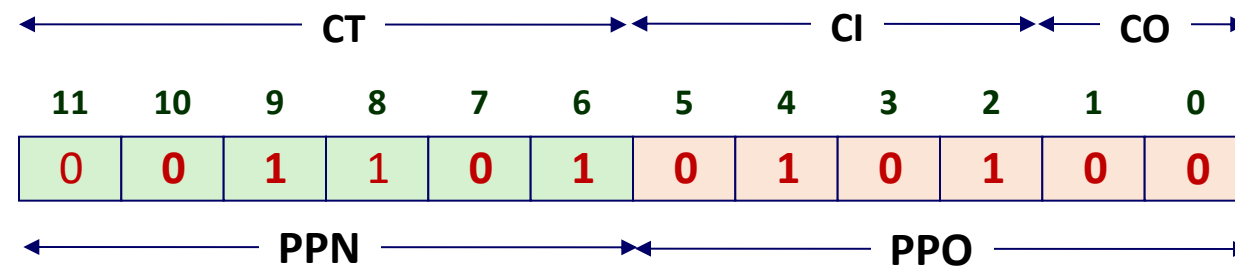
TANDON SCHOOL
OF ENGINEERING

Address Translation Example #1

Virtual Address: 0x03D4



Physical Address



16 entry
DM cache

CO: 0 CI: 0x5 CT: 0x0D Hit? Y Byte



NYU

TANDON SCHOOL
OF ENGINEERING

Current state of caches/tables

TLB

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0

VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	08	13	1
01	–	0	09	17	1
02	33	1	0A	09	1
03	02	1	0B	–	0
04	–	0	0C	–	0
05	16	1	0D	2D	1
06	–	0	0E	11	1
07	–	0	0F	0D	1

Cache

Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03

Page table

Idx	Tag	Valid	B0	B1	B2	B3
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–

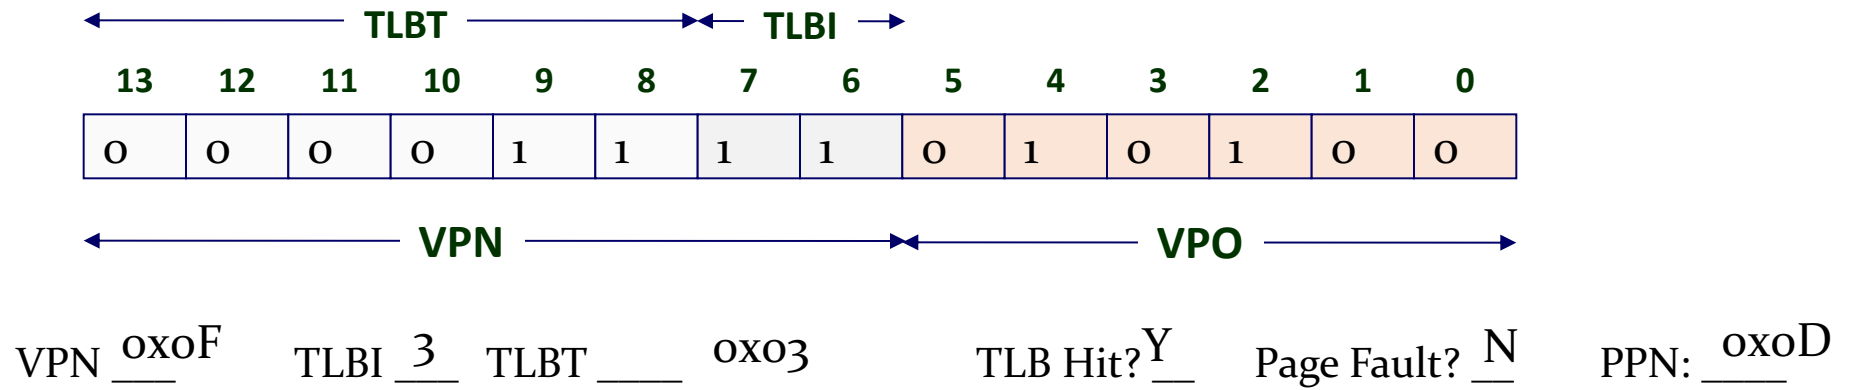


NYU

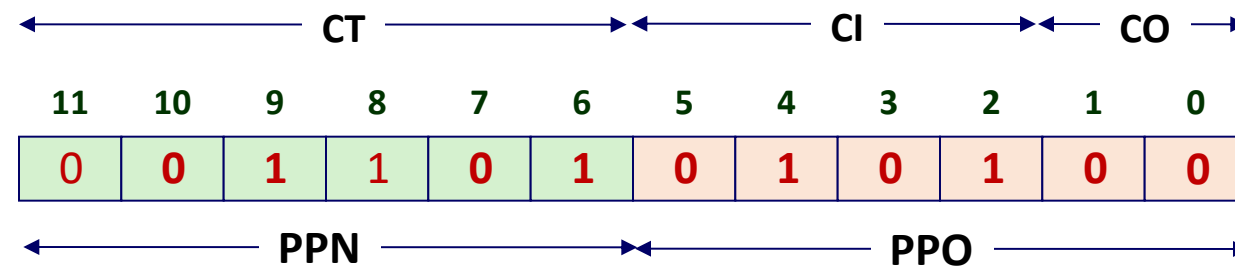
TANDON SCHOOL
OF ENGINEERING

Address Translation Example #1

Virtual Address: 0x03D4



Physical Address



CO: 0, CI: 0x5, CT: 0x0D, Hit? Y, Byte: 0x36

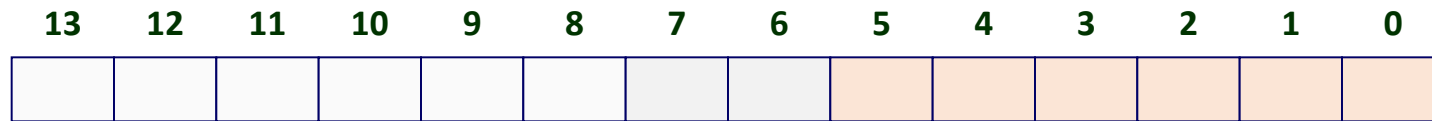


NYU

TANDON SCHOOL
OF ENGINEERING

Address Translation Example #2

Virtual Address: 0x0B8F



VPN ____

TLBI ____

TLBT ____

TLB Hit? ____

Page Fault? ____

PPN: ____

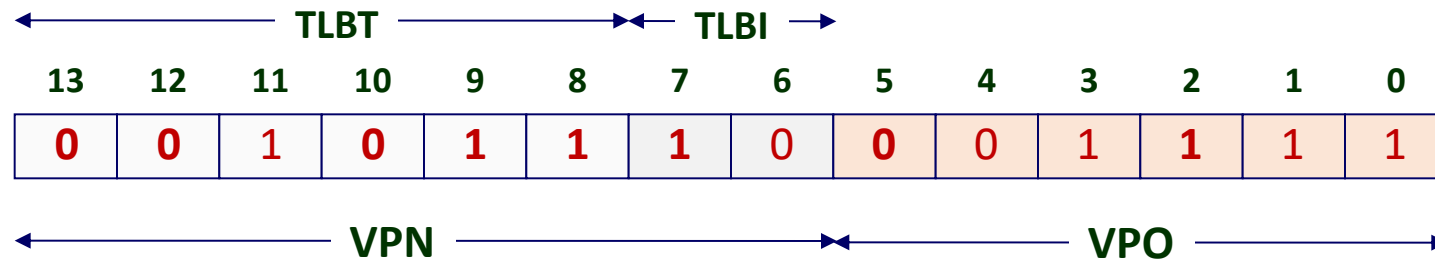


NYU

TANDON SCHOOL
OF ENGINEERING

Address Translation Example #2

Virtual Address: 0x0B8F



VPN 0x2E TLBI 2 TLBT 0x0B TLB Hit? Page Fault? PPN:



NYU

TANDON SCHOOL
OF ENGINEERING

Current state of caches/tables

TLB

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0

VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	08	13	1
01	–	0	09	17	1
02	33	1	0A	09	1
03	02	1	0B	–	0
04	–	0	0C	–	0
05	16	1	0D	2D	1
06	–	0	0E	11	1
07	–	0	0F	0D	1

Cache

Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03

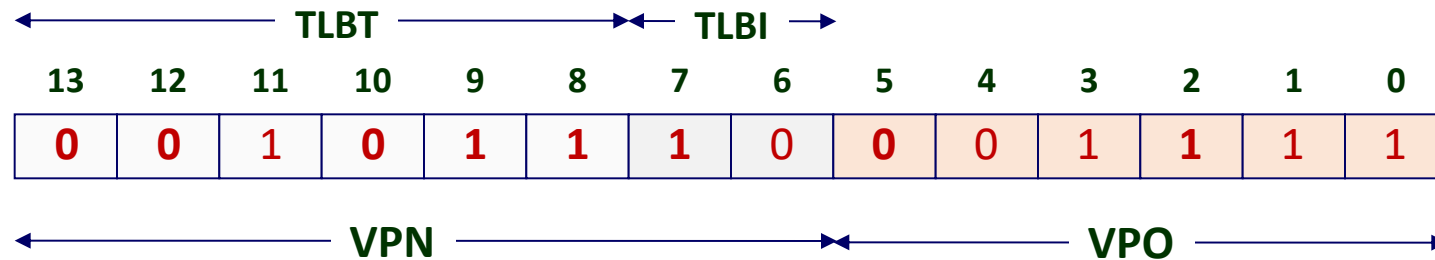
Page table

Idx	Tag	Valid	B0	B1	B2	B3
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–



Address Translation Example #2

Virtual Address: 0x0B8F



VPN 0x2E

TLBI 2

TLBT 0x0B

TLB Hit? N

Page Fault?

PPN:



NYU

TANDON SCHOOL
OF ENGINEERING

Current state of caches/tables

TLB

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0

VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	08	13	1
01	–	0	09	17	1
02	33	1	0A	09	1
03	02	1	0B	–	0
04	–	0	0C	–	0
05	16	1	0D	2D	1
06	–	0	0E	11	1
07	–	0	0F	0D	1



Cache

Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03

Page table

Idx	Tag	Valid	B0	B1	B2	B3
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–

We'll assume
Valid bit is zero

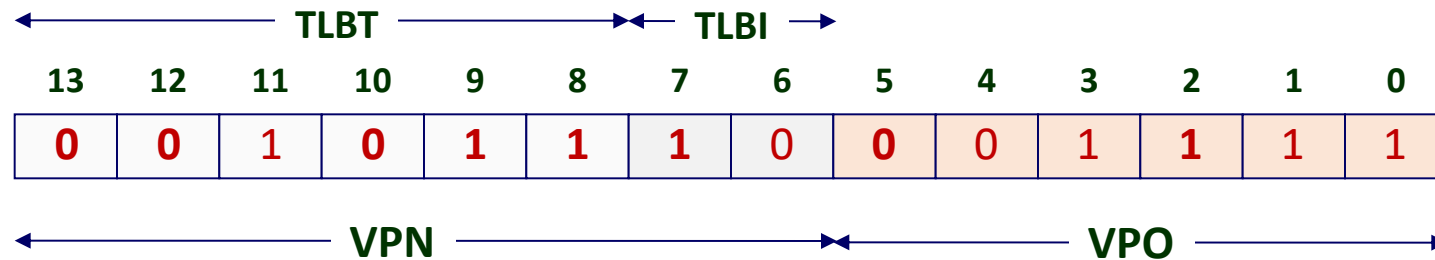


NYU

TANDON SCHOOL
OF ENGINEERING

Address Translation Example #2

Virtual Address: 0x0B8F



VPN 0x2E TLBI 2 TLBT 0x0B

TLB Hit? N Page Fault? Y PPN: TBD

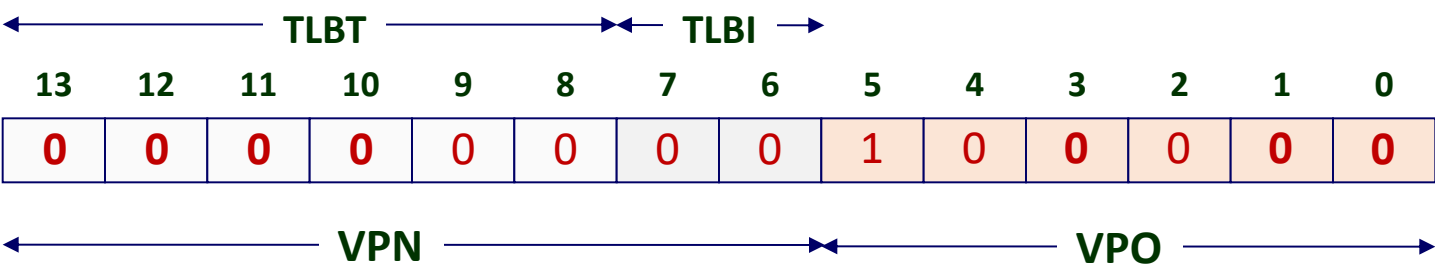


NYU

TANDON SCHOOL
OF ENGINEERING

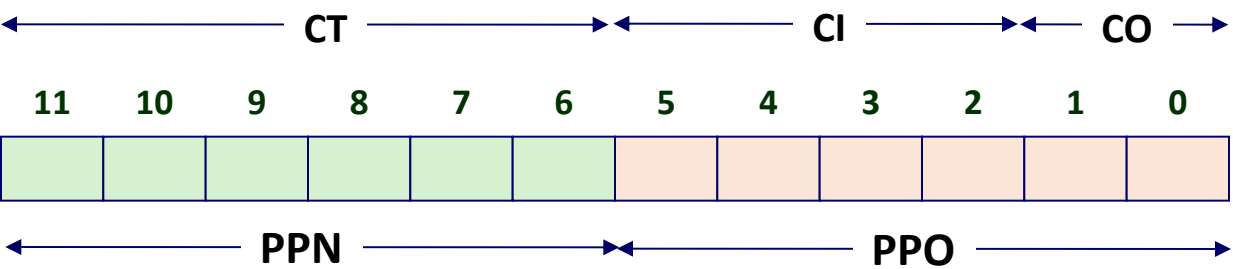
Address Translation Example #3

Virtual Address: 0x0020



VPN 0x00 TLBI 0 TLBT 0x00 TLB Hit? Page Fault? PPN:

Physical Address



CO CI CT Hit? Byte:

Current state of caches/tables

TLB

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0

VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	08	13	1
01	–	0	09	17	1
02	33	1	0A	09	1
03	02	1	0B	–	0
04	–	0	0C	–	0
05	16	1	0D	2D	1
06	–	0	0E	11	1
07	–	0	0F	0D	1

Cache

Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03

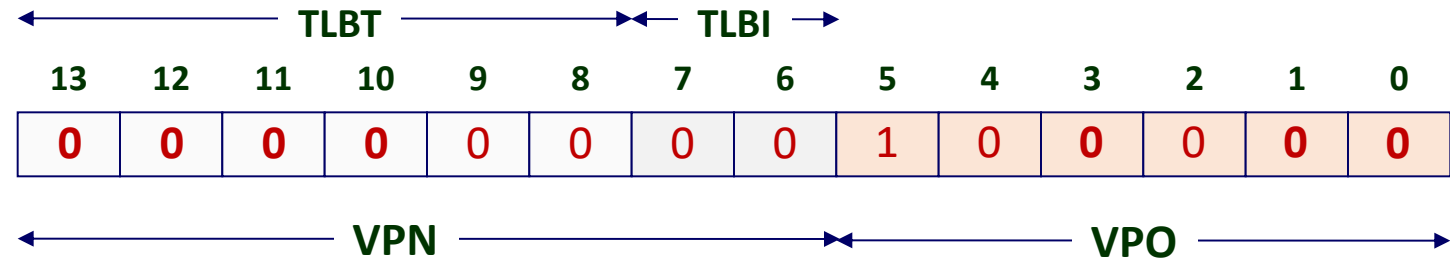
Page table

Idx	Tag	Valid	B0	B1	B2	B3
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–



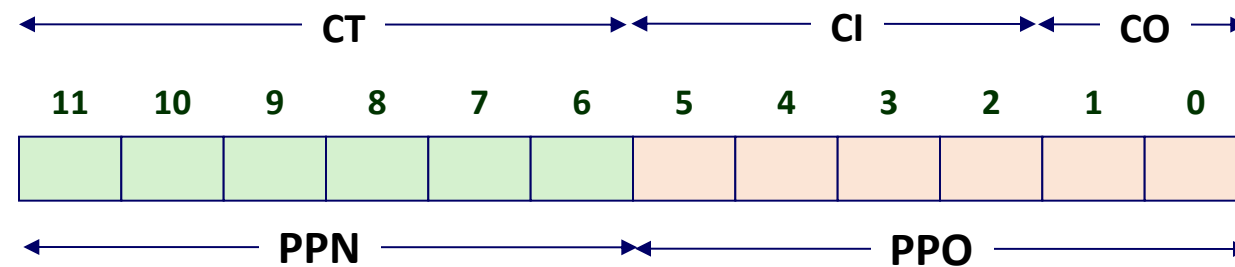
Address Translation Example #3

Virtual Address: 0x0020



VPN 0x00 TLBI 0 TLBT 0x00 TLB Hit? N Page Fault? PPN:

Physical Address



CO CI CT Hit? Byte:



NYU

TANDON SCHOOL
OF ENGINEERING

Current state of caches/tables

TLB

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0



VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	08	13	1
01	–	0	09	17	1
02	33	1	0A	09	1
03	02	1	0B	–	0
04	–	0	0C	–	0
05	16	1	0D	2D	1
06	–	0	0E	11	1
07	–	0	0F	0D	1

Cache

Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03


Page table

Idx	Tag	Valid	B0	B1	B2	B3
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–




Address Translation Example #3

Virtual Address: 0x0020



VPN 0x00 TLBI 0 TLBT 0x00 TLB Hit? N Page Fault? N PPN: 0x28

Physical Address



CO 0 CI 0x8 CT 0x28 Hit? Byte:

Current state of caches/tables

TLB

Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0

VPN	PPN	Valid	VPN	PPN	Valid
00	28	1	08	13	1
01	–	0	09	17	1
02	33	1	0A	09	1
03	02	1	0B	–	0
04	–	0	0C	–	0
05	16	1	0D	2D	1
06	–	0	0E	11	1
07	–	0	0F	0D	1

Cache

Idx	Tag	Valid	B0	B1	B2	B3
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03


Page table

Idx	Tag	Valid	B0	B1	B2	B3
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–




Address Translation Example #3

Virtual Address: 0x0020



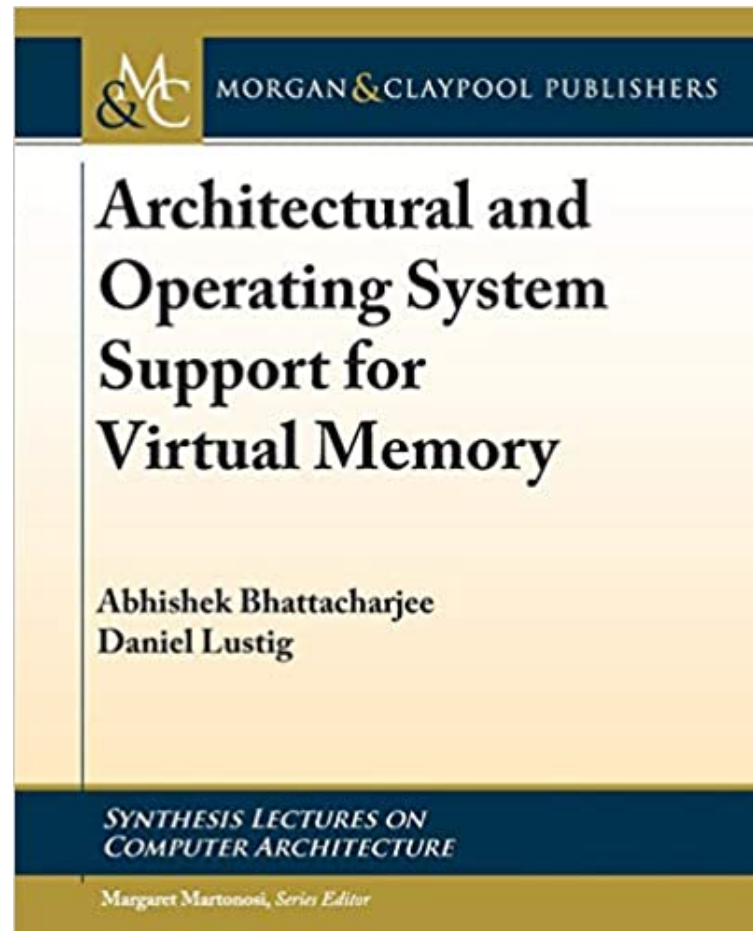
VPN 0x00 TLBI 0 TLBT 0x00 TLB Hit? N Page Fault? N PPN: 0x28

Physical Address



CO 0 CI 0x8 CT 0x28 Hit? N Byte: < from mem >

If you want to know more check out



NYU

**TANDON SCHOOL
OF ENGINEERING**

Dynamic scheduling

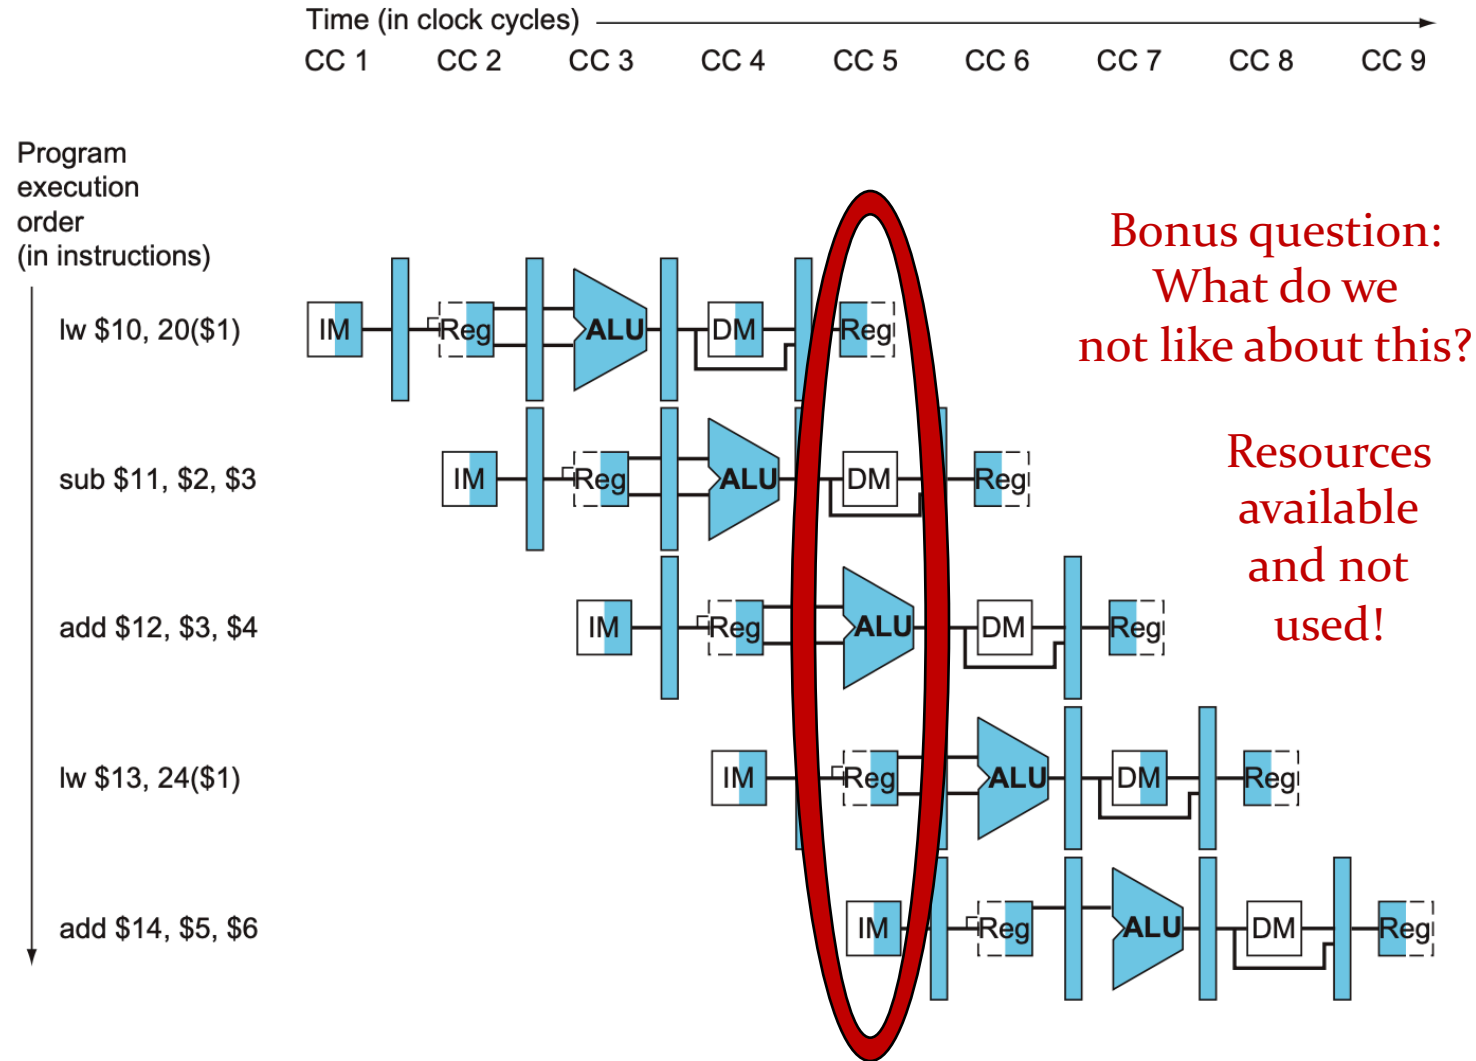


NYU

**TANDON SCHOOL
OF ENGINEERING**

Problem 1

```
lw    $10, 20($1)
sub    $11, $2, $3
add    $12, $3, $4
lw    $13, 24($1)
add    $14, $5, $6
```



NYU

TANDON SCHOOL
OF ENGINEERING

Problem 2: Unnecessary stalls (causes underutilization)

Add r3, r2, r1
Lw r2, o(ro)
Sub r4, r2, r1
Add r6, r7, r8

CPI?

$$5/4 = 1.25$$

Add r3, r2, r1
Lw r2, o(ro)
Add r6, r7, r8
Sub r4, r2, r1

CPI?

$$4/4 = 1.0!$$

Same functionality, better performance!



NYU

TANDON SCHOOL
OF ENGINEERING

Solution: Data flow!

$a := x + y$
 $b := a \times a$
 $c := 4 - a$

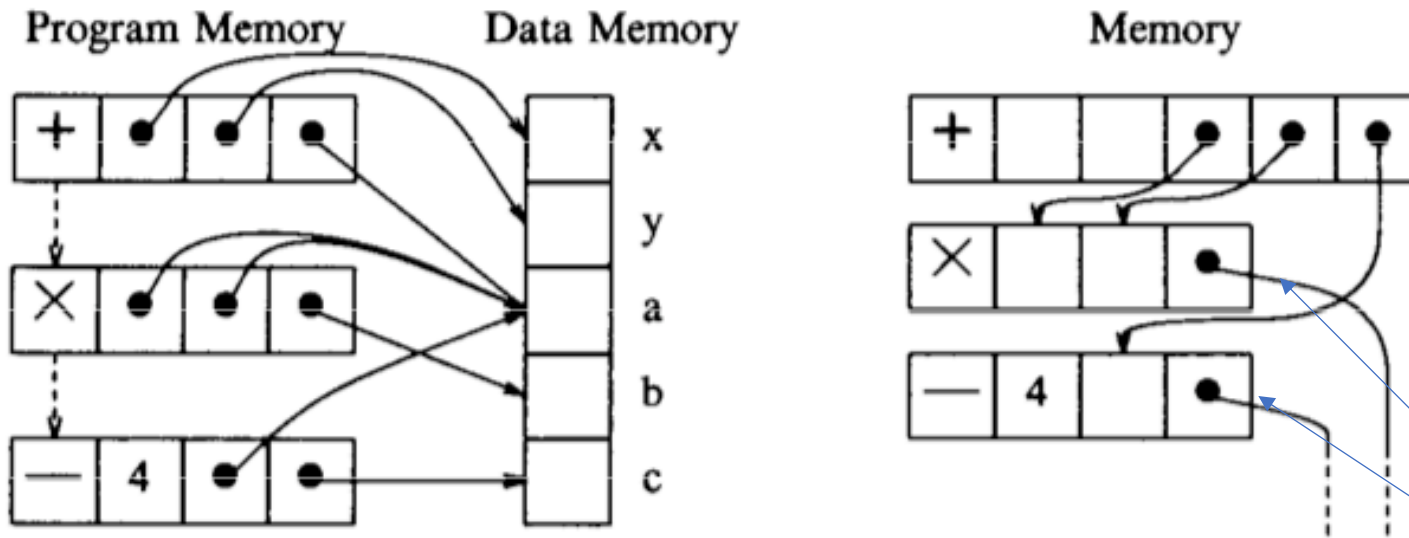
How are we going to do the conversion?

uArch tricks!

Instructions dependent, even if results aren't.

This is how you write programs

It's very restrictive



Only tracks real dependencies.
“Which instructions need what I’m computing?”

This is what OoO machines do!

What can we say about these two?

Figure 2. A comparison of control flow and dataflow programs. On the left a control flow program for a computer with memory-to-memory instructions. The arcs point to the locations of data that are to be used or created. Control flow arcs are indicated with dashed arrows; usually most of them are implicit. In the equivalent dataflow program on the right only one memory is involved. Each instruction contains pointers to all instructions that consume its results.

Parallel instructions

“Dataflow Machine Architecture” Heen, 1986.



NYU

TAN
OF ENGINEERING

Instruction-Level Parallelism (ILP)

Fine-grained parallelism

A measure of **inter-instruction dependency** in an app

- ILP assumes a unit-cycle operation, infinite resources, prefect frontend
 - Theoretically, how many instructions could I process per cycle?
- ILP != IPC
- $IPC = \# \text{ instructions} / \# \text{ cycles}$
- ILP is the **upper bound** of attainable IPC

Enabled and improved by RISC

- More ILP of RISC over CISC does not imply a better overall performance
 - CISC can be implemented like RISC

Limited by **dependencies**



NYU

TANDON SCHOOL
OF ENGINEERING

Back to hazards...

Control: Where do I go next?

- Caused by branch instructions

Data: There's more!

- RAW: Read-After-Write
- WAW: Write-After-Write
- WAR: Write-After-Read

Why am I just telling you this now?

When we dynamically schedule, the program order is broken.


Must guarantee there are no artifacts!

Hazards limit ILP.

WAW:

```
Add r2, r2, r2
sub r2, r2, r2
```

WAR:



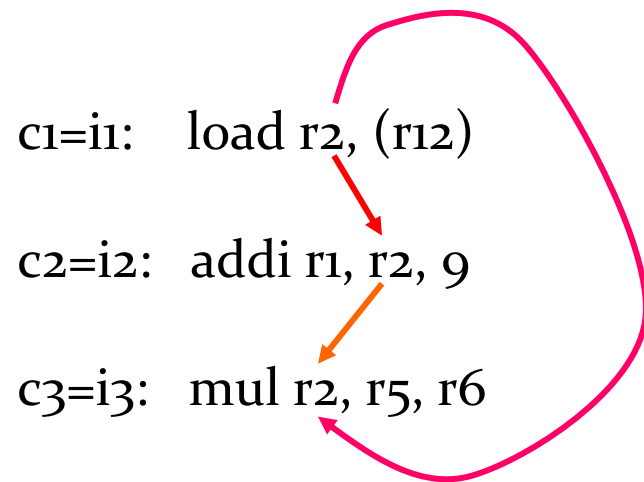
NYU

TANDON SCHOOL
OF ENGINEERING

True/False dependencies and ILP

True dependency forces “sequentiality”

$$\text{ILP} = 3/3 = 1$$



False dependency removed

$$\text{ILP} = 3/2 = 1.5$$

Diagram illustrating false dependency removal:

- i1: load r2, (r12)
- i2: addi r1, r2, 9
- i3: mul r8, r5, r6

A red arrow points from i1 to i2. A large yellow arrow points down from i3, indicating that the false dependency between i2 and i3 is removed.

c1: load r2, (r12)

c2: addi r1, r2, 9

mul r8, r5, r6



NYU

TANDON SCHOOL
OF ENGINEERING

Exploiting ILP

- Control speculation

Branch prediction!

- Dynamic scheduling

Reschedule program ordering of instructions

today



- Register renaming

Break false dependencies

Use uArch to provide illusion of “more” registers..

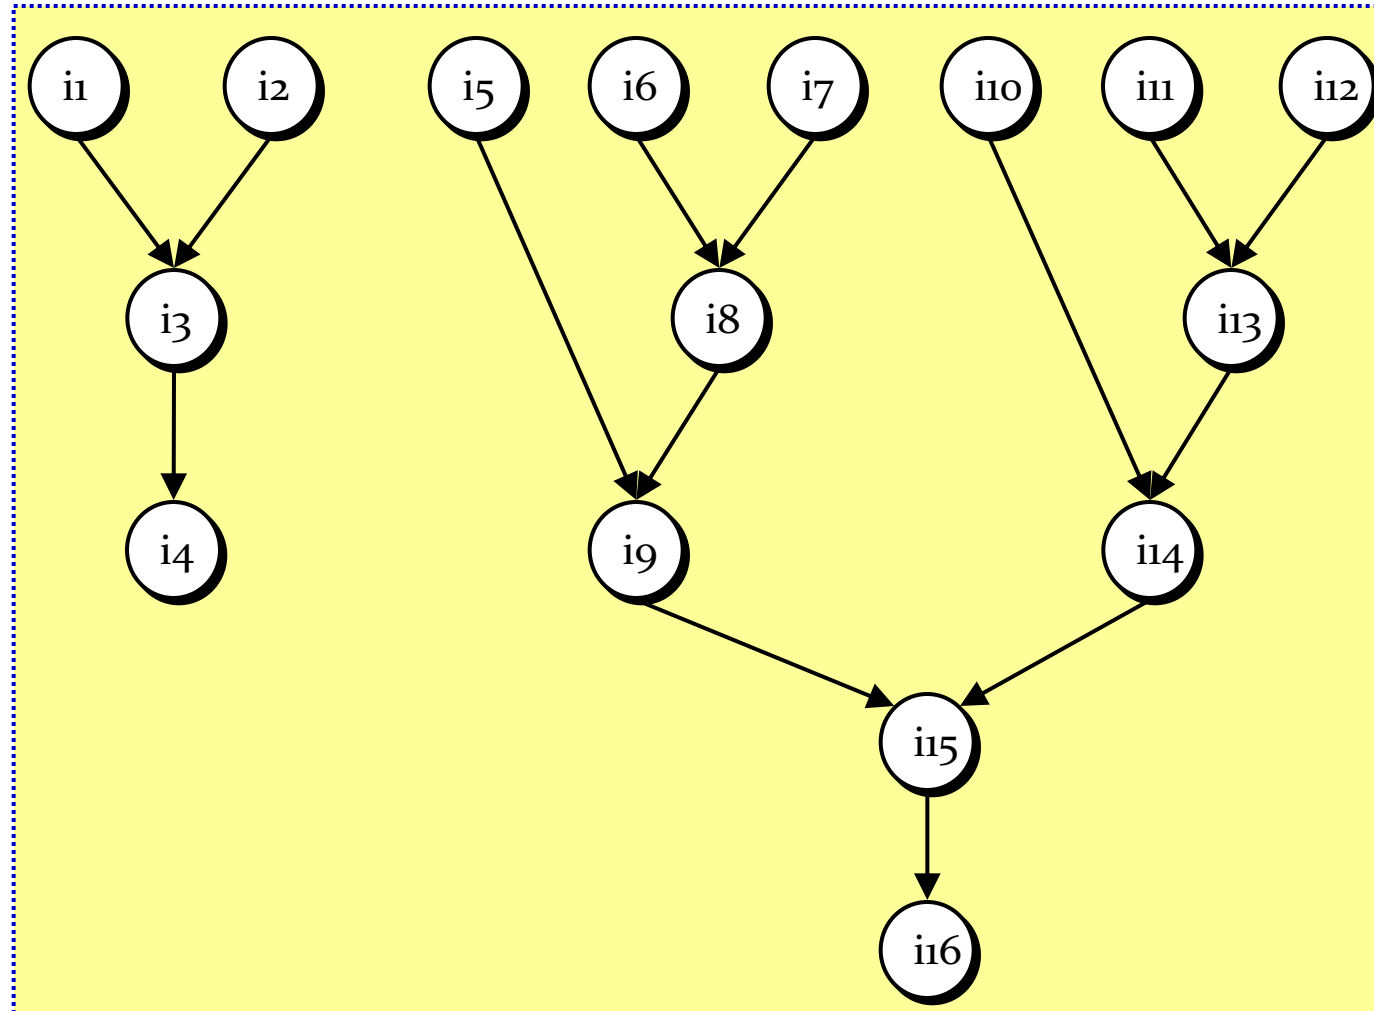
- Dynamic memory disambiguation

Discuss later



Step 1: Construct data flow graph

i1: $r2 = 4(r22)$
i2: $r10 = 4(r25)$
i3: $r10 = r2 + r10$
i4: $4(r26) = r10$
i5: $r14 = 8(r27)$
i6: $r6 = (r22)$
i7: $r5 = (r23)$
i8: $r5 = r6 - r5$
i9: $r4 = r14 * r5$
i10: $r15 = 12(r27)$
i11: $r7 = 4(r22)$
i12: $r8 = 4(r23)$
i13: $r8 = r7 - r8$
i14: $r8 = r15 * r8$
i15: $r8 = r4 - r8$
i16: $(r28) = r8$



This is ideal.
Why not
possible?

Data Flow Graph (or Data Dependency Graph, DDG)

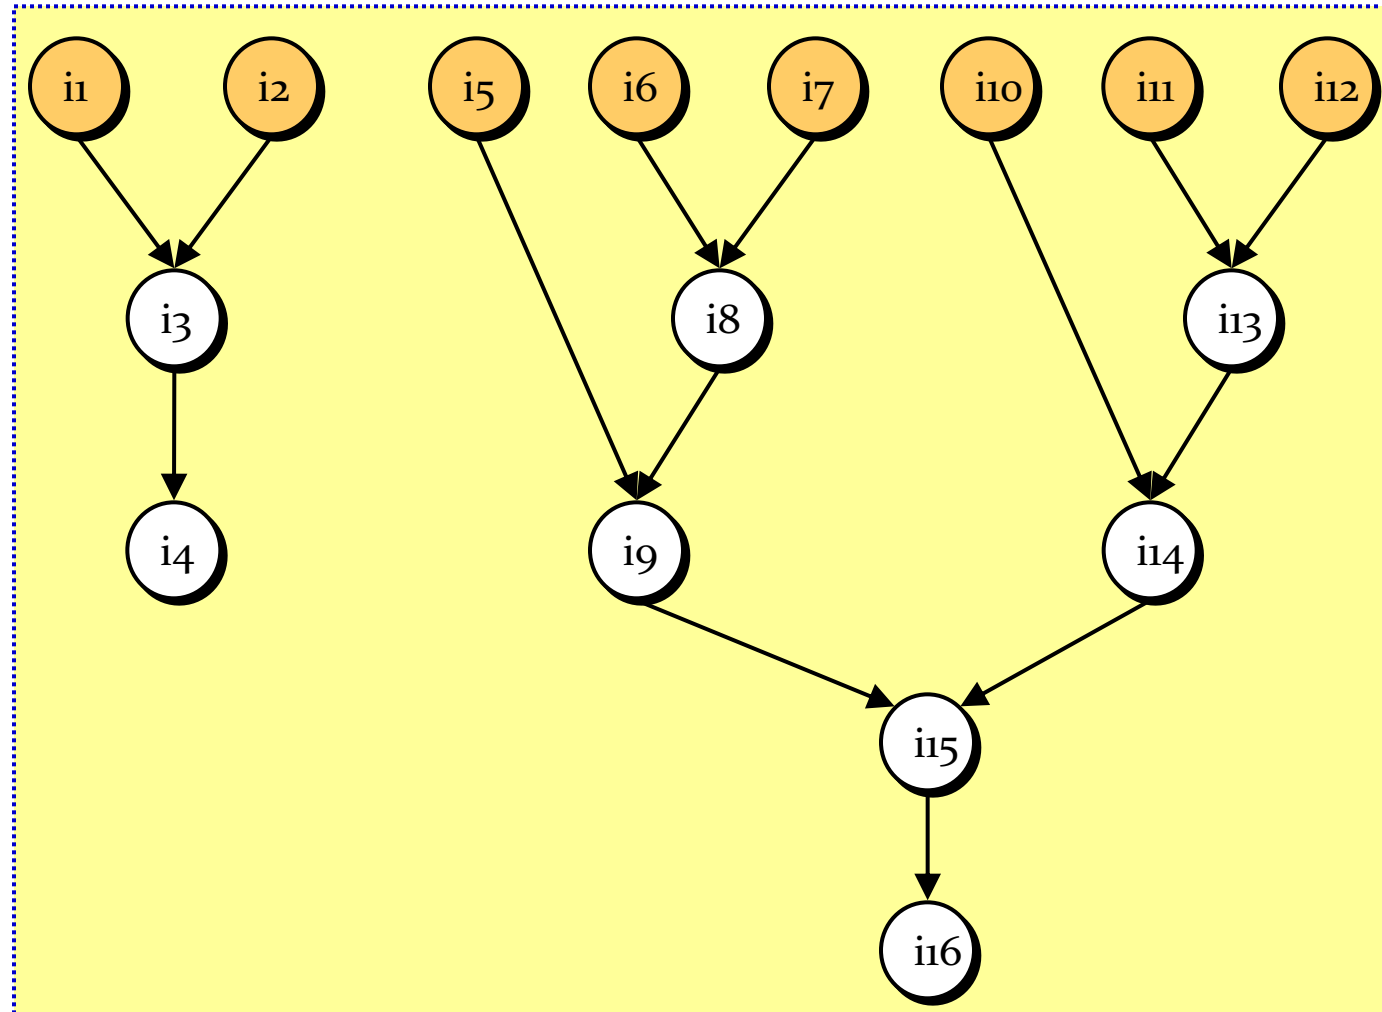


NYU

TANDON SCHOOL
OF ENGINEERING

Step 1: Construct data flow graph

i1: $r2 = 4(r22)$
i2: $r10 = 4(r25)$
i3: $r10 = r2 + r10$
i4: $4(r26) = r10$
i5: $r14 = 8(r27)$
i6: $r6 = (r22)$
i7: $r5 = (r23)$
i8: $r5 = r6 - r5$
i9: $r4 = r14 * r5$
i10: $r15 = 12(r27)$
i11: $r7 = 4(r22)$
i12: $r8 = 4(r23)$
i13: $r8 = r7 - r8$
i14: $r8 = r15 * r8$
i15: $r8 = r4 - r8$
i16: $(r28) = r8$



This is ideal.
Why not
possible?



NYU

TANDON SCHOOL
OF ENGINEERING

ILP “Windows”

ILP = 1

$$R_5 = 8(R_6)$$

$$R_7 = R_5 - R_4$$

$$R_9 = R_7 * R_7$$

ILP = 1.5

$$R_{15} = 16(R_6)$$

$$R_{17} = R_{15} - R_{14}$$

$$R_{19} = R_{15} * R_{15}$$

ILP = ?



NYU

TANDON SCHOOL
OF ENGINEERING

Bigger window

$$R_5 = 8(R_6)$$

$$R_7 = R_5 - R_4$$

$$R_9 = R_7 * R_7$$

$$R_{15} = 16(R_6)$$

$$R_{17} = R_{15} - R_{14}$$

$$R_{19} = R_{15} * R_{15}$$



Bigger => better ILP!

C₁: $R_5 = 8(R_6)$

$R_{15} = 16(R_6)$

C₂: $R_7 = R_5 - R_4$

$R_{17} = R_{15} - R_{14}$

$R_{19} = R_{15} * R_{15}$

$R_9 = R_7 * R_7$

ILP = $6/3 = 2$ better than 1 and 1.5

Larger window gives more opportunities

But what limits the window?



Impact of registers on ILP

$$R_1 = 8(R_0)$$

$$R_3 = R_1 - 5$$

$$R_2 = R_1 * R_3$$

$$24(R_0) = R_2$$

$$R_1 = 16(R_0)$$

$$R_3 = R_1 - 5$$

$$R_2 = R_1 * R_3$$

$$32(R_0) = R_2$$

When only 4 registers available

ILP =



NYU

TANDON SCHOOL
OF ENGINEERING

Impact of registers on ILP

$$R_1 = 8(R_0)$$

$$R_3 = R_1 - 5$$

$$R_2 = R_1 * R_3$$

$$24(R_0) = R_2$$

$$R_1 = 16(R_0)$$

$$R_3 = R_1 - 5$$

$$R_2 = R_1 * R_3$$

$$32(R_0) = R_2$$

When only 4 registers available

$$ILP = \#inst / \#cycles = 8 / 8 = 1$$



Impact of registers on ILP

$$R_1 = 8(R_0)$$

$$R_3 = R_1 - 5$$

$$R_2 = R_1 * R_3$$

$$24(R_0) = R_2$$

$$R_1 = 16(R_0)$$

$$R_3 = R_1 - 5$$

$$R_2 = R_1 * R_3$$

$$32(R_0) = R_2$$

When more (**8**) registers available,
rewrite code to improve ILP

ILP =



NYU

TANDON SCHOOL
OF ENGINEERING

Impact of registers on ILP

$$R_1 = 8(R_0)$$

$$R_3 = R_1 - 5$$

$$R_2 = R_1 * R_3$$

$$24(R_0) = R_2$$

$$R_5 = 16(R_0)$$

$$R_6 = R_5 - 5$$

$$R_7 = R_5 * R_6$$

$$32(R_0) = R_7$$

$$R_1 = 8(R_0)$$

$$R_3 = R_1 - 5$$

$$R_2 = R_1 * R_3$$

$$24(R_0) = R_2$$

$$R_5 = 16(R_0)$$

$$R_6 = R_5 - 5$$

$$R_7 = R_5 * R_6$$

$$32(R_0) = R_7$$

When more (8)
registers available,
rewrite code to
improve ILP

$$ILP = \#inst / \# cycles = 8/4 = 2!$$



NYU

TANDON SCHOOL
OF ENGINEERING

Step 2: Execute nodes!

When can an instruction execute?

- All source operands are ready
- Execution unit available
- Destination is ready (to be written)

This is dynamic scheduling,
we no longer follow the program order of instructions,
instead use data flow to execute computation

Dynamic scheduling enables Out-of-Order execution

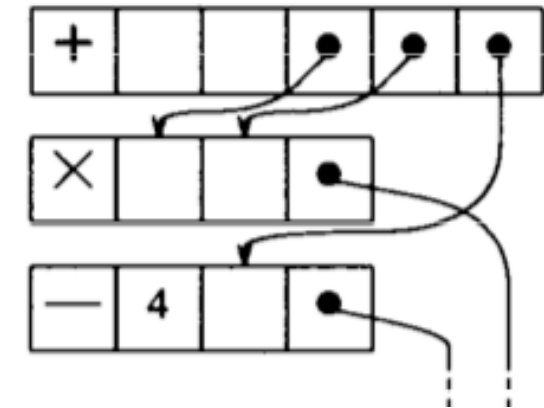
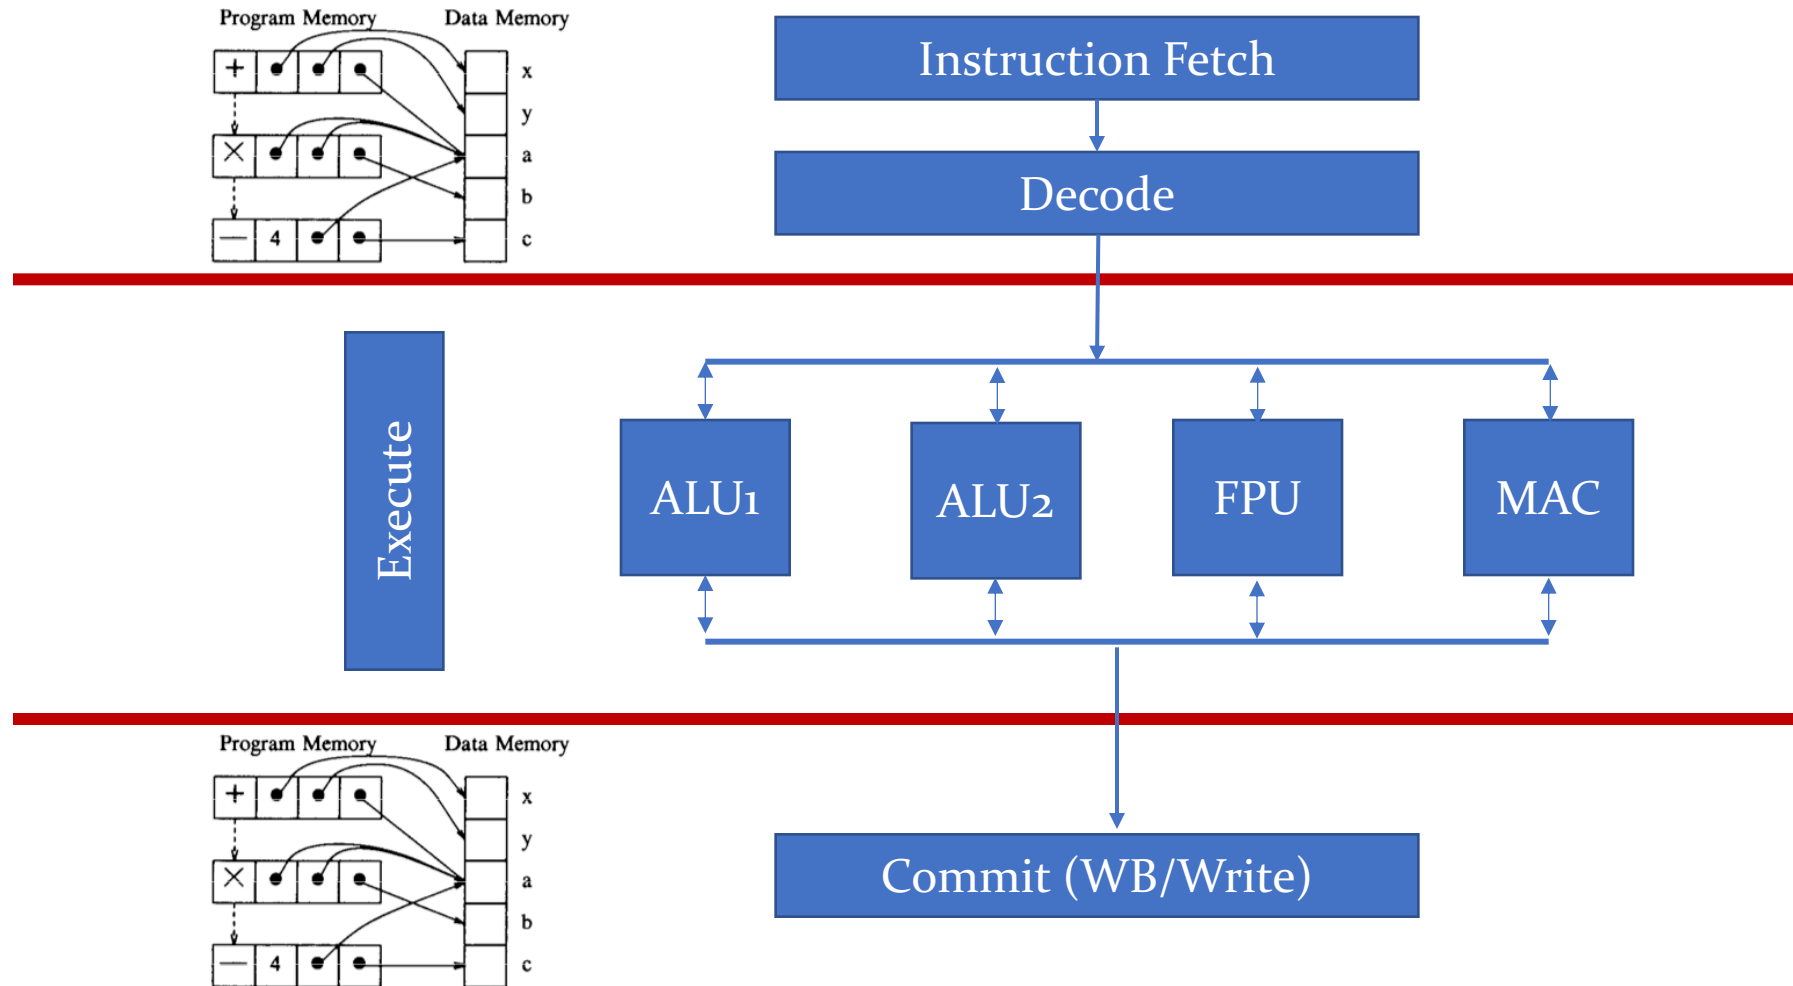


NYU

**TANDON SCHOOL
OF ENGINEERING**

OoO General Scheme

Call this:
“Restricted dataflow”



Next time:
Tomasulo!



NYU

TANDON SCHOOL
OF ENGINEERING

OoO (or “O3” or OOO) execution

OoO execution \equiv out-of-order computation

OoO execution \neq out-of-order retirement (commit)

Commit \Rightarrow when instruction updates the architectural state

No (speculative) instruction allowed to commit
until it is confirmed on the right path

Fetch and decode (i.e., front-end)
are still done in the program order

Why?



NYU

TANDON SCHOOL
OF ENGINEERING